

Predictive Maintenance with Deep Learning and Flink

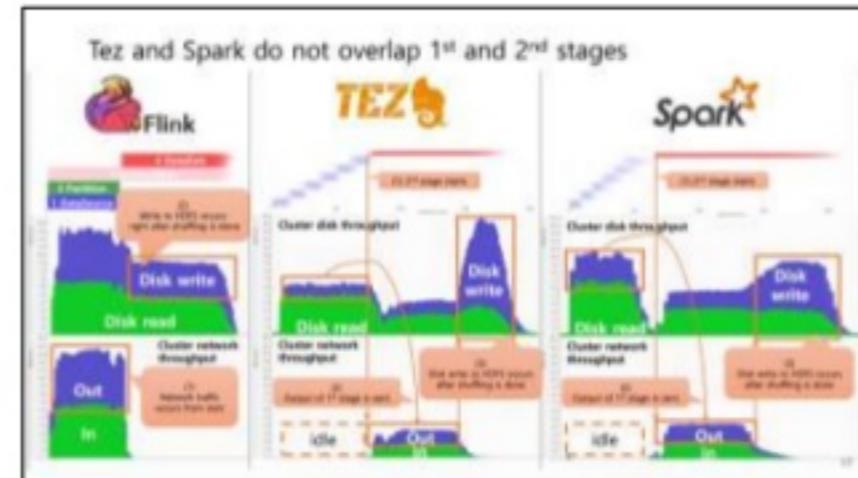
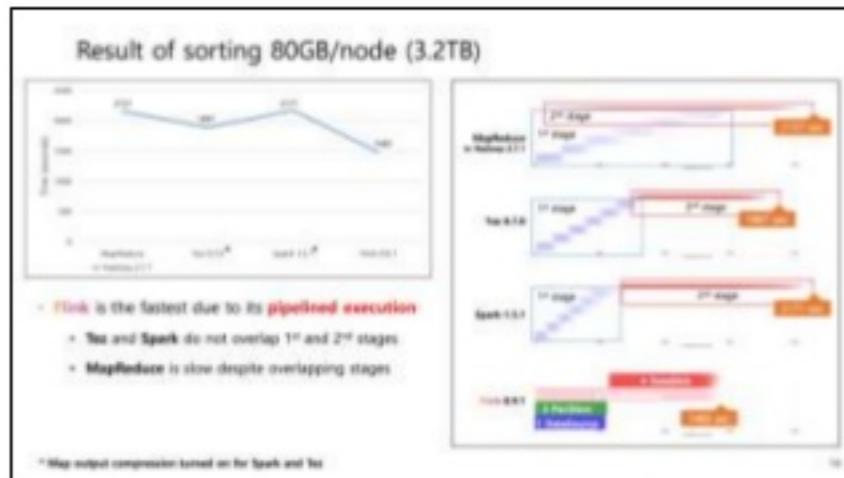


Dongwon Kim, PhD
Solution R&D center
SK Telecom

About me (@eastcirclek)

- Big Data processing engines
 - MapReduce, Tez, **Spark**, **Flink**, Hive, Presto
- Recent interest
 - Deep learning model serving (**TensorFlow serving**)
 - Containerization (**Docker**, **Kubernetes**)
 - Time-series data (**InfluxDB**, **Prometheus**, **Grafana**)
- Flink Forward 2015 Berlin
 - A Comparative Performance Evaluation of **Flink**

Covered
in this talk

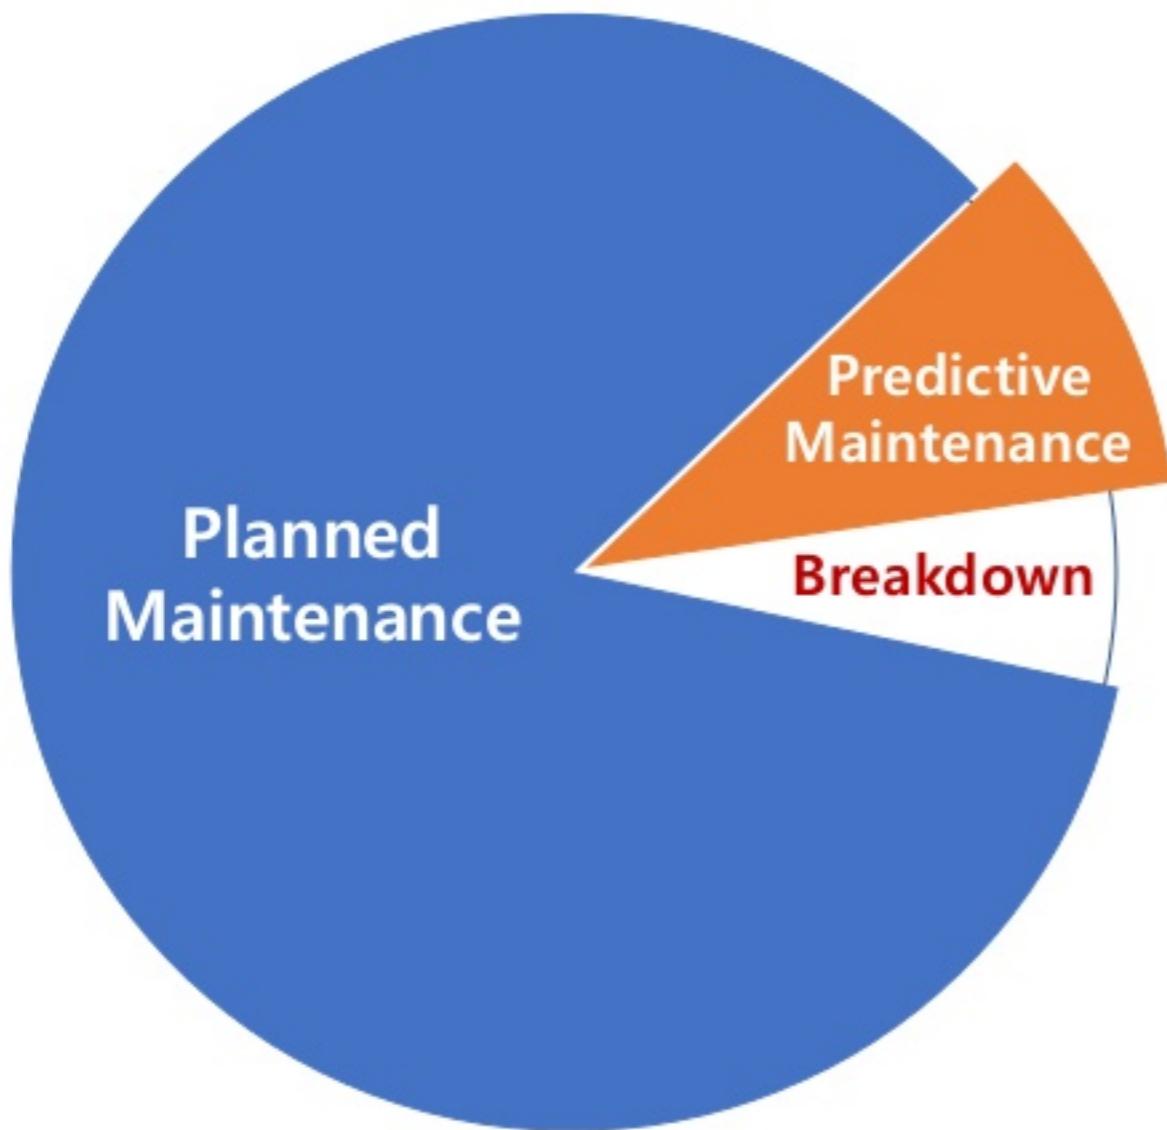


Refinery and semiconductor companies depend on equipment



Breakdown of equipment largely affects company profit

Equipment maintenance to minimize **breakdown**



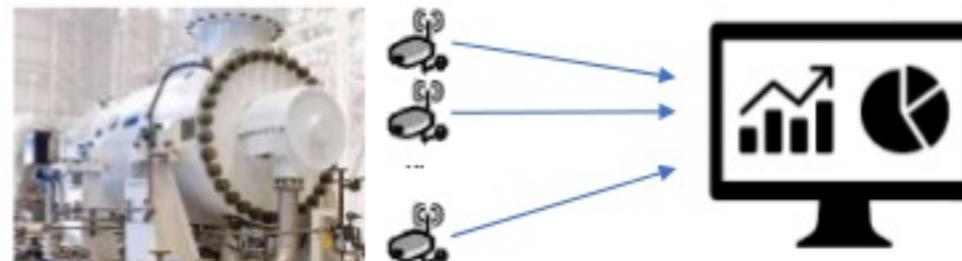
Planned Maintenance

Shutdown equipment on a regular basis for parts replacement, cleaning, adjustments, etc.



Predictive Maintenance (PdM)

Unplanned maintenance based on
prediction of equipment condition

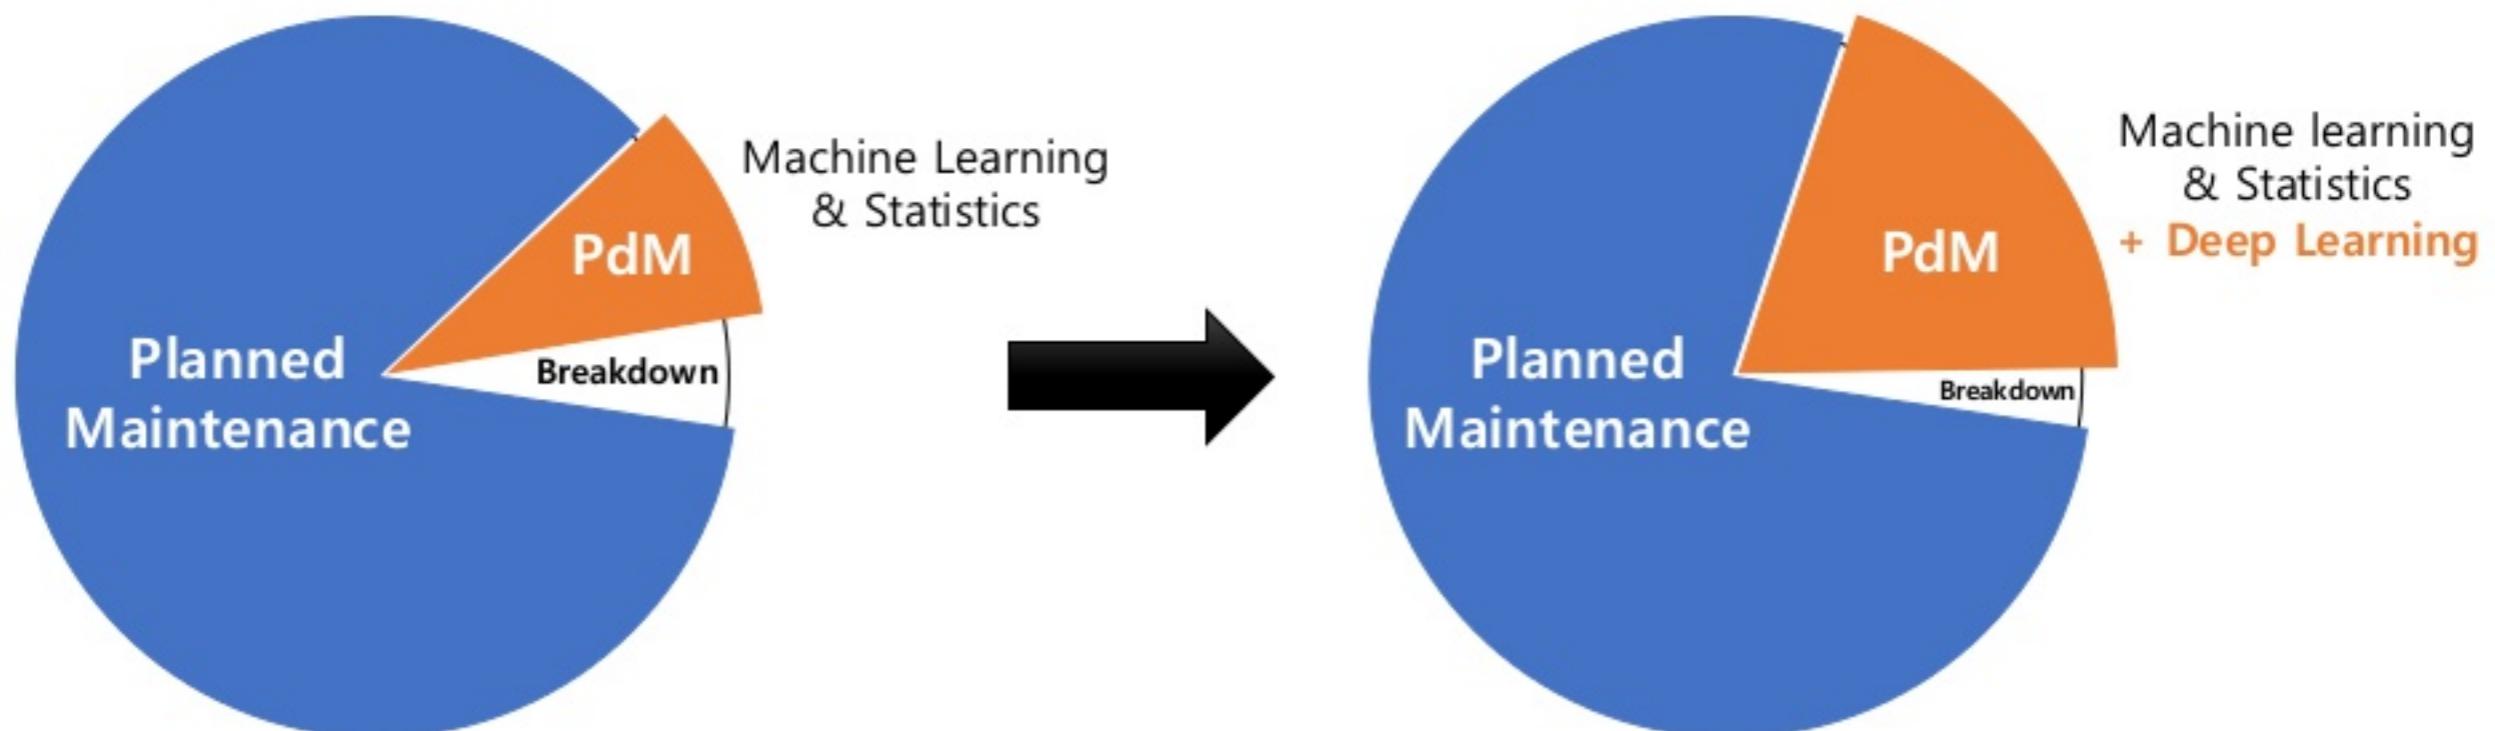


<equipment sensors>

<predictive maintenance system>

Our approach to **Predictive Maintenance**

**Better prediction of equipment condition
using Deep Learning**

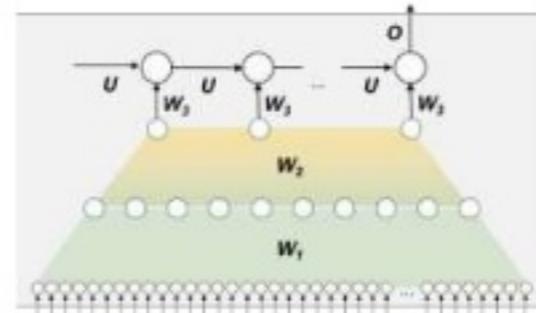


Contents

1

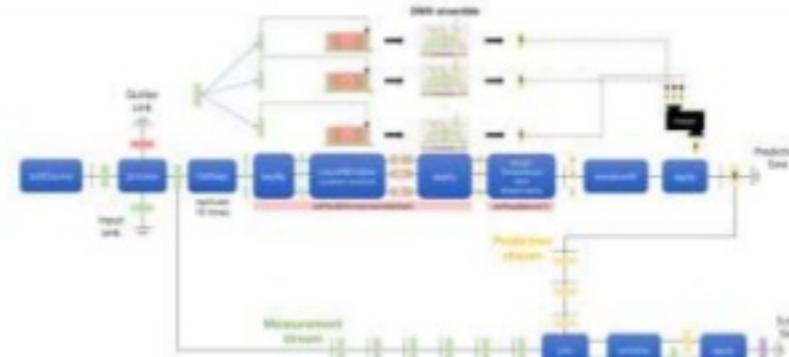
Why we use Flink
for our time-series prediction model

Time-series prediction model



2

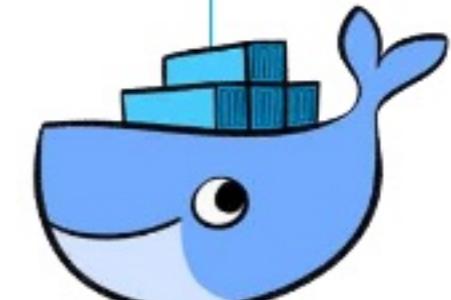
Flink pipeline design for
rendezvous and DNN ensemble



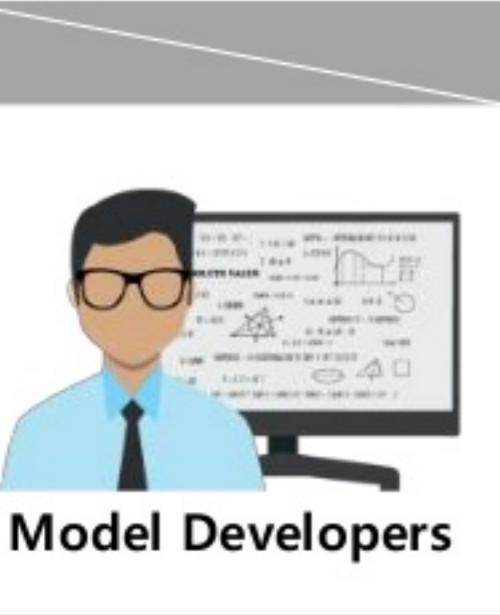
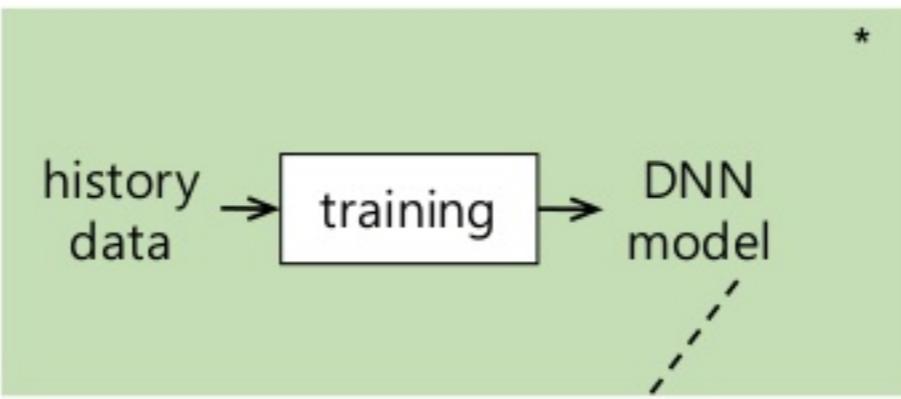
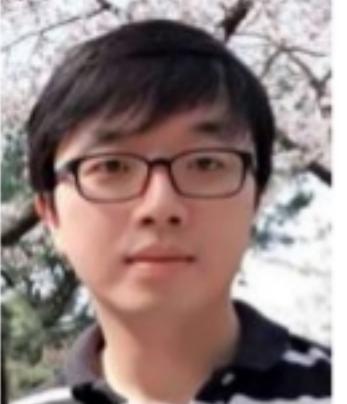
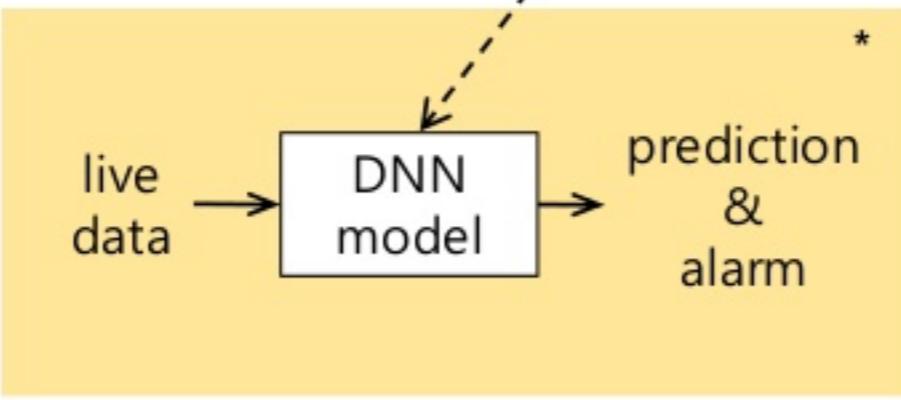
K Keras

3

Solution packaging and monitoring
with Docker and Prometheus



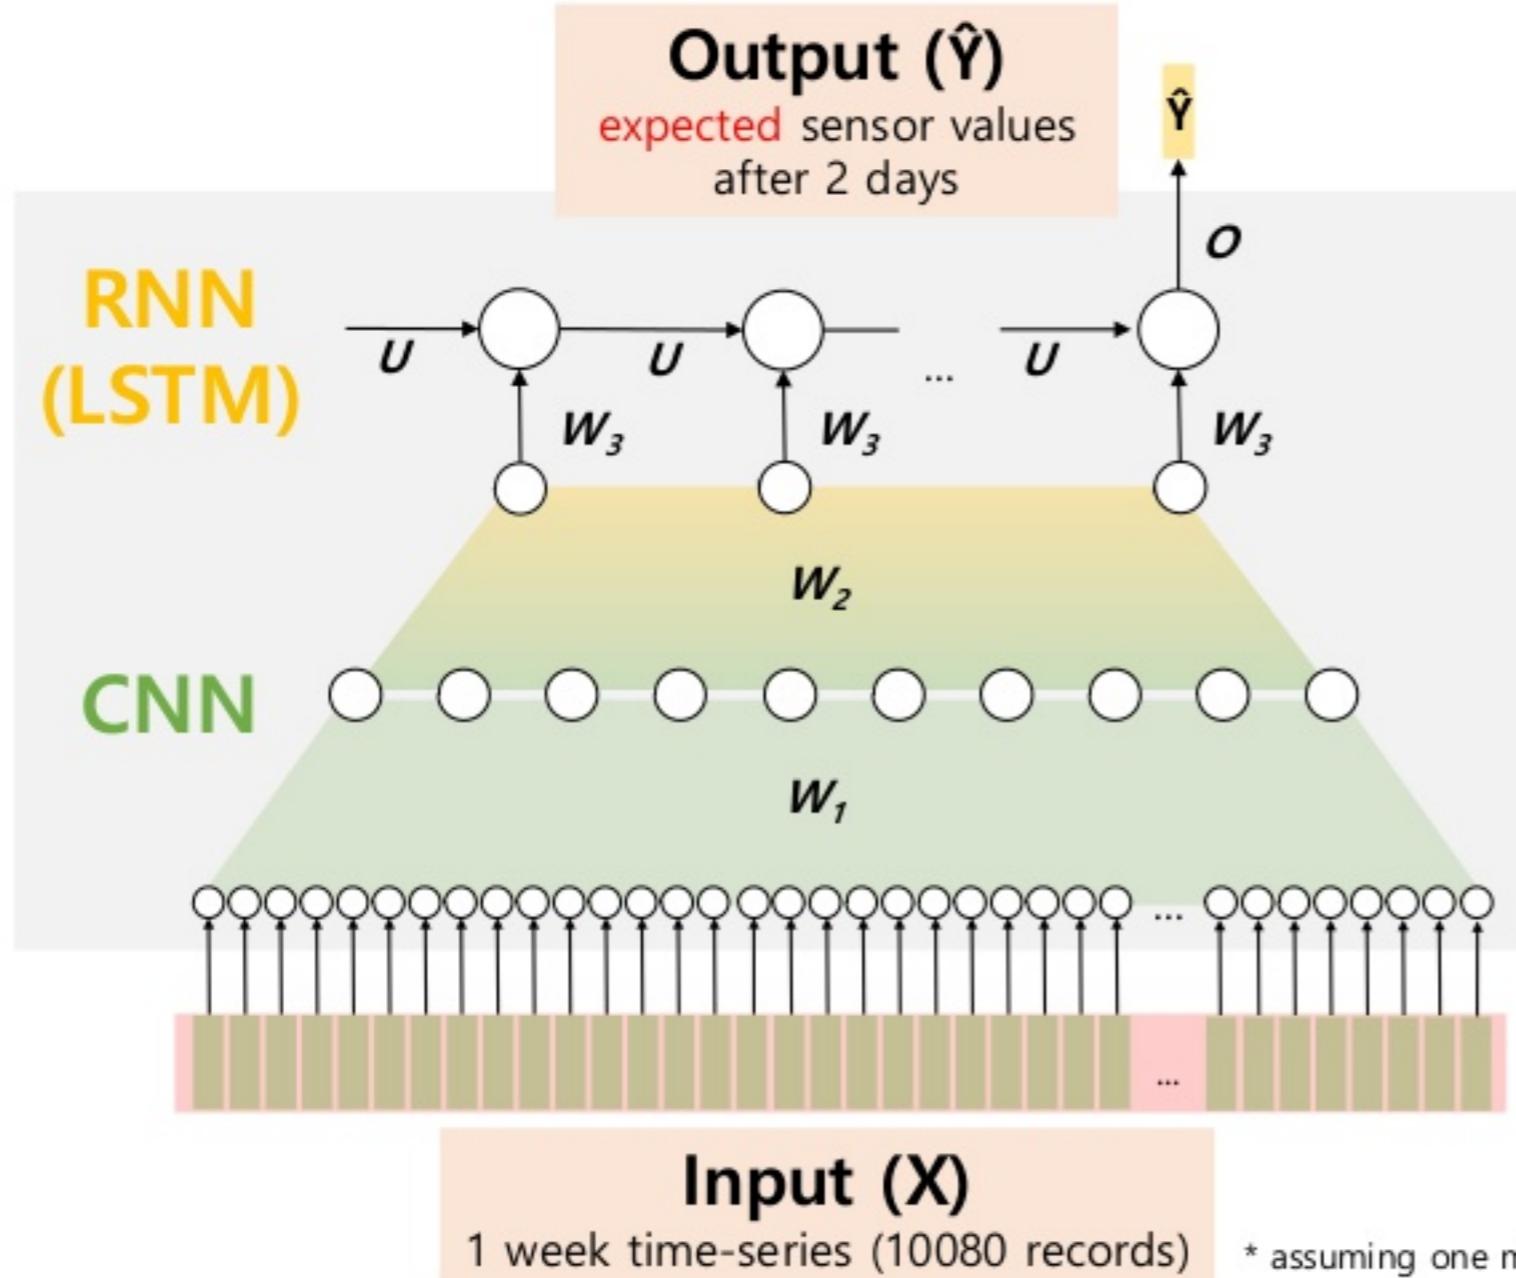
My team consists of two groups

	Role	Toolbox
 Model Developers	 <pre>graph LR; A[history data] --> B[training]; B --> C[DNN model]; C -.-> D[DNN model]; D --> E[prediction & alarm]</pre>	 Keras  TensorFlow
 Data Engineers	 <pre>graph LR; A[live data] --> B[DNN model]; B --> C[prediction & alarm]; C -.-> D[training]</pre>	 python  Apache Spark™ 

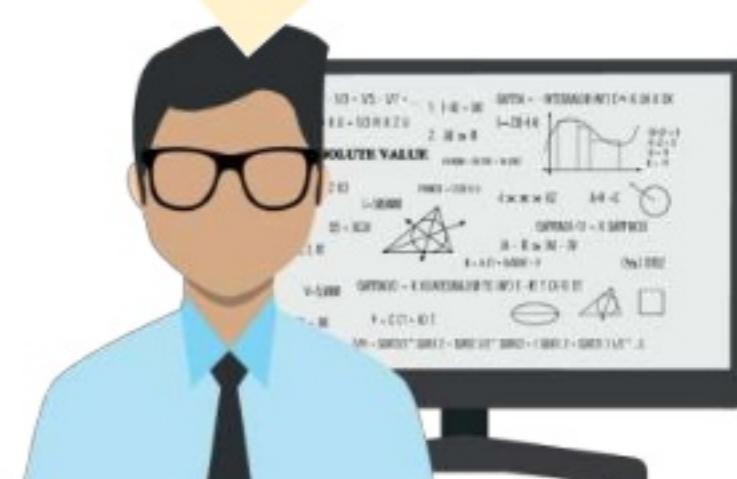
* The diagram is based on Ted Dunning's slides (Flink Forward 2017 SF)



Model developers give **Convolutional LSTM** to engineers



It **does not return** whether the equipment breaks down after 2 days

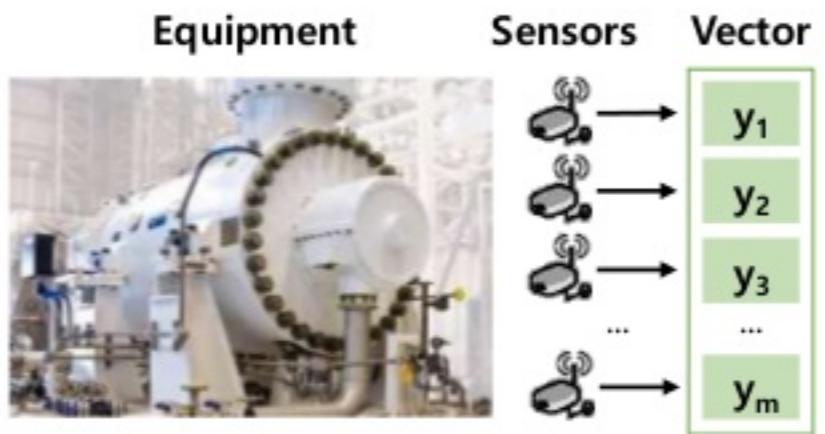


Model developer

* assuming one minute sampling rate



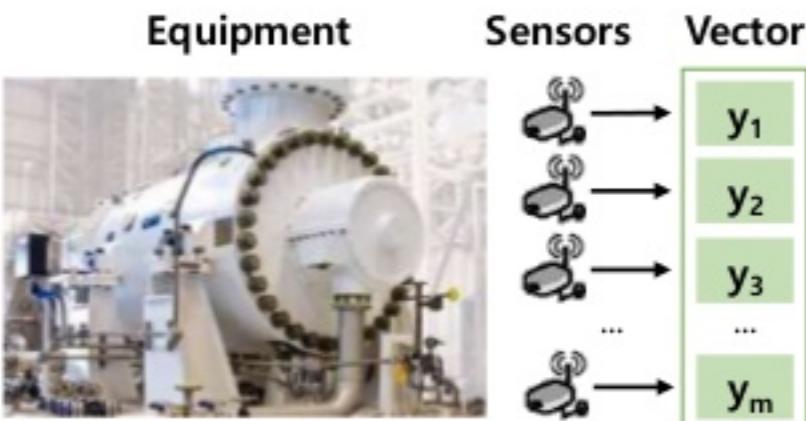
Data engineers apply **Convolutional LSTM** to live sensor data



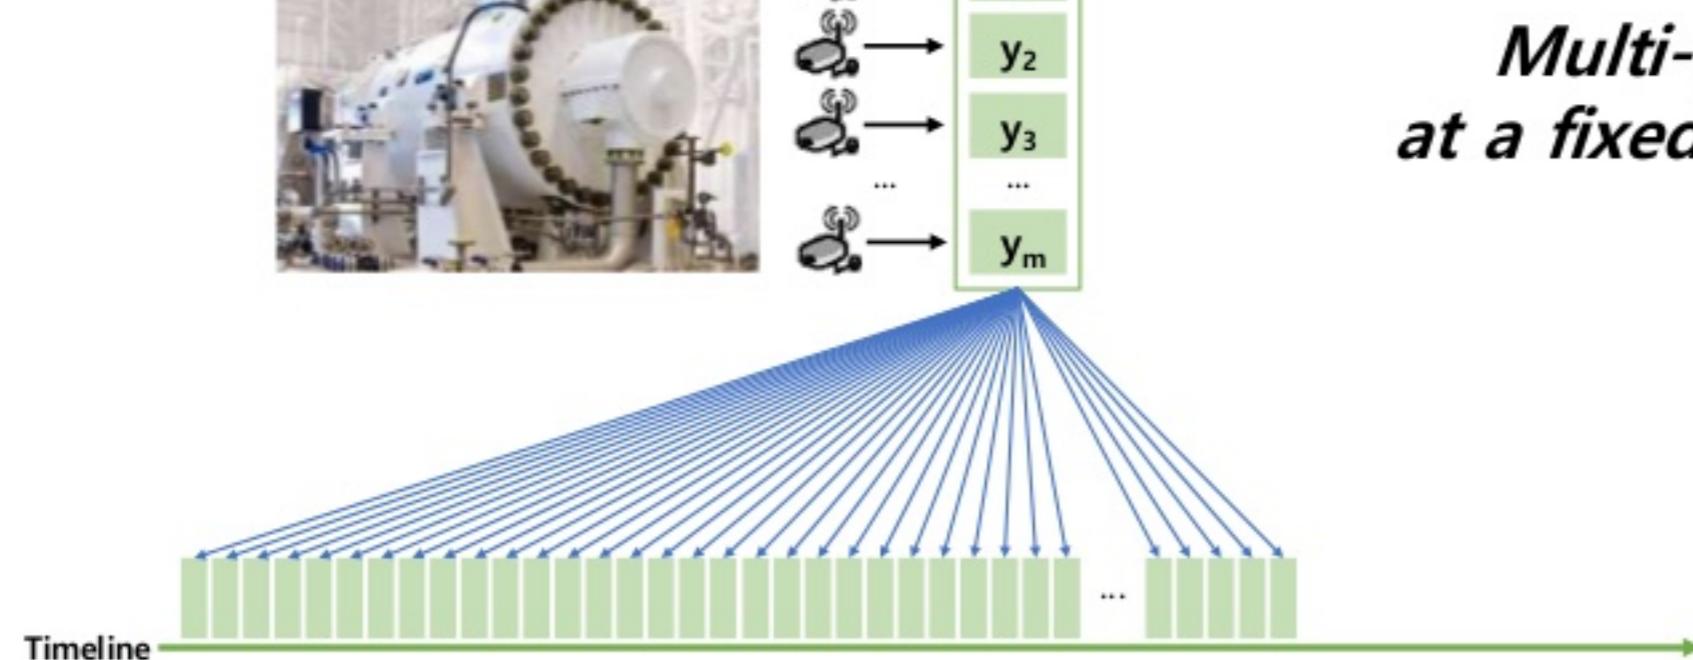
We have multi-sensor data



Data engineers apply **Convolutional LSTM** to live sensor data

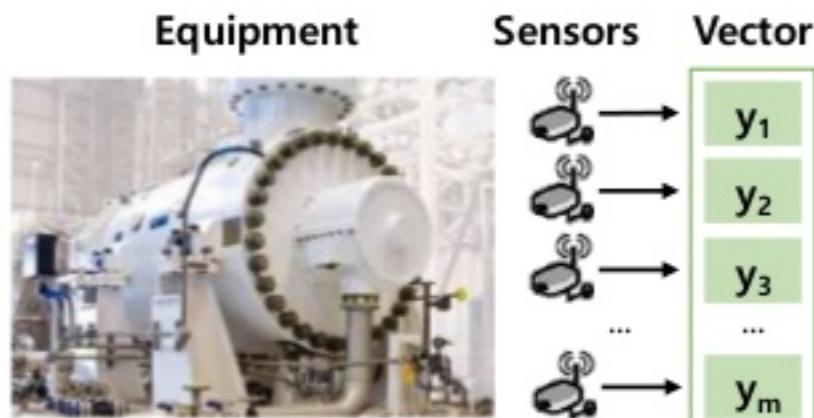


*Multi-sensor data arrive
at a fixed interval of 1 minute*





Data engineers apply **Convolutional LSTM** to live sensor data



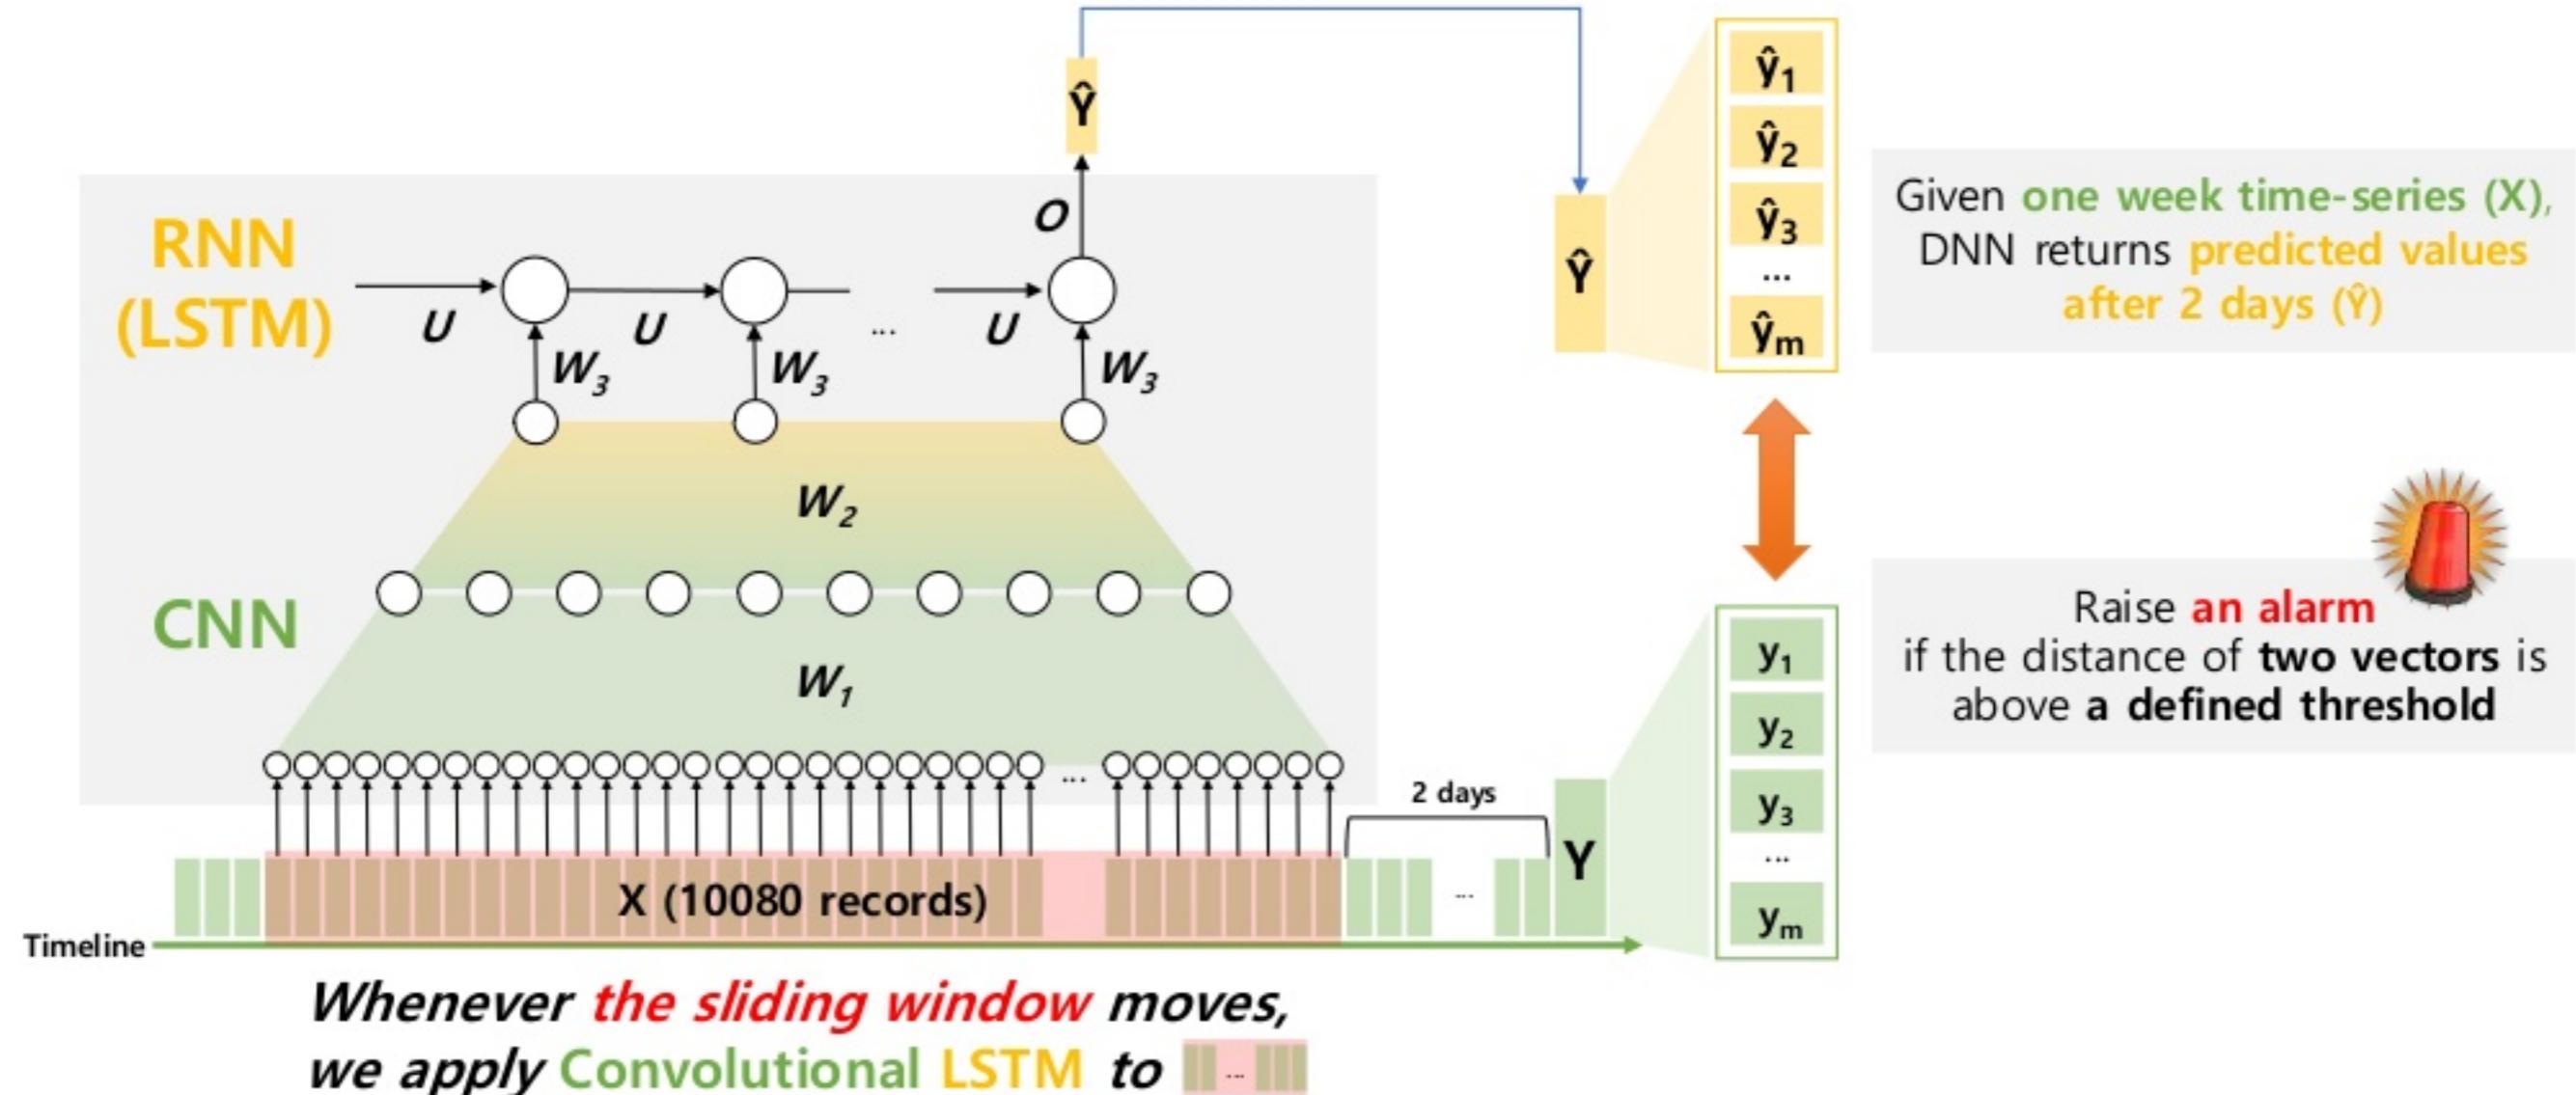
*We maintain **a count window** over **the latest 10080 records***

*It slides as a new record arrives
(a sliding count window)*

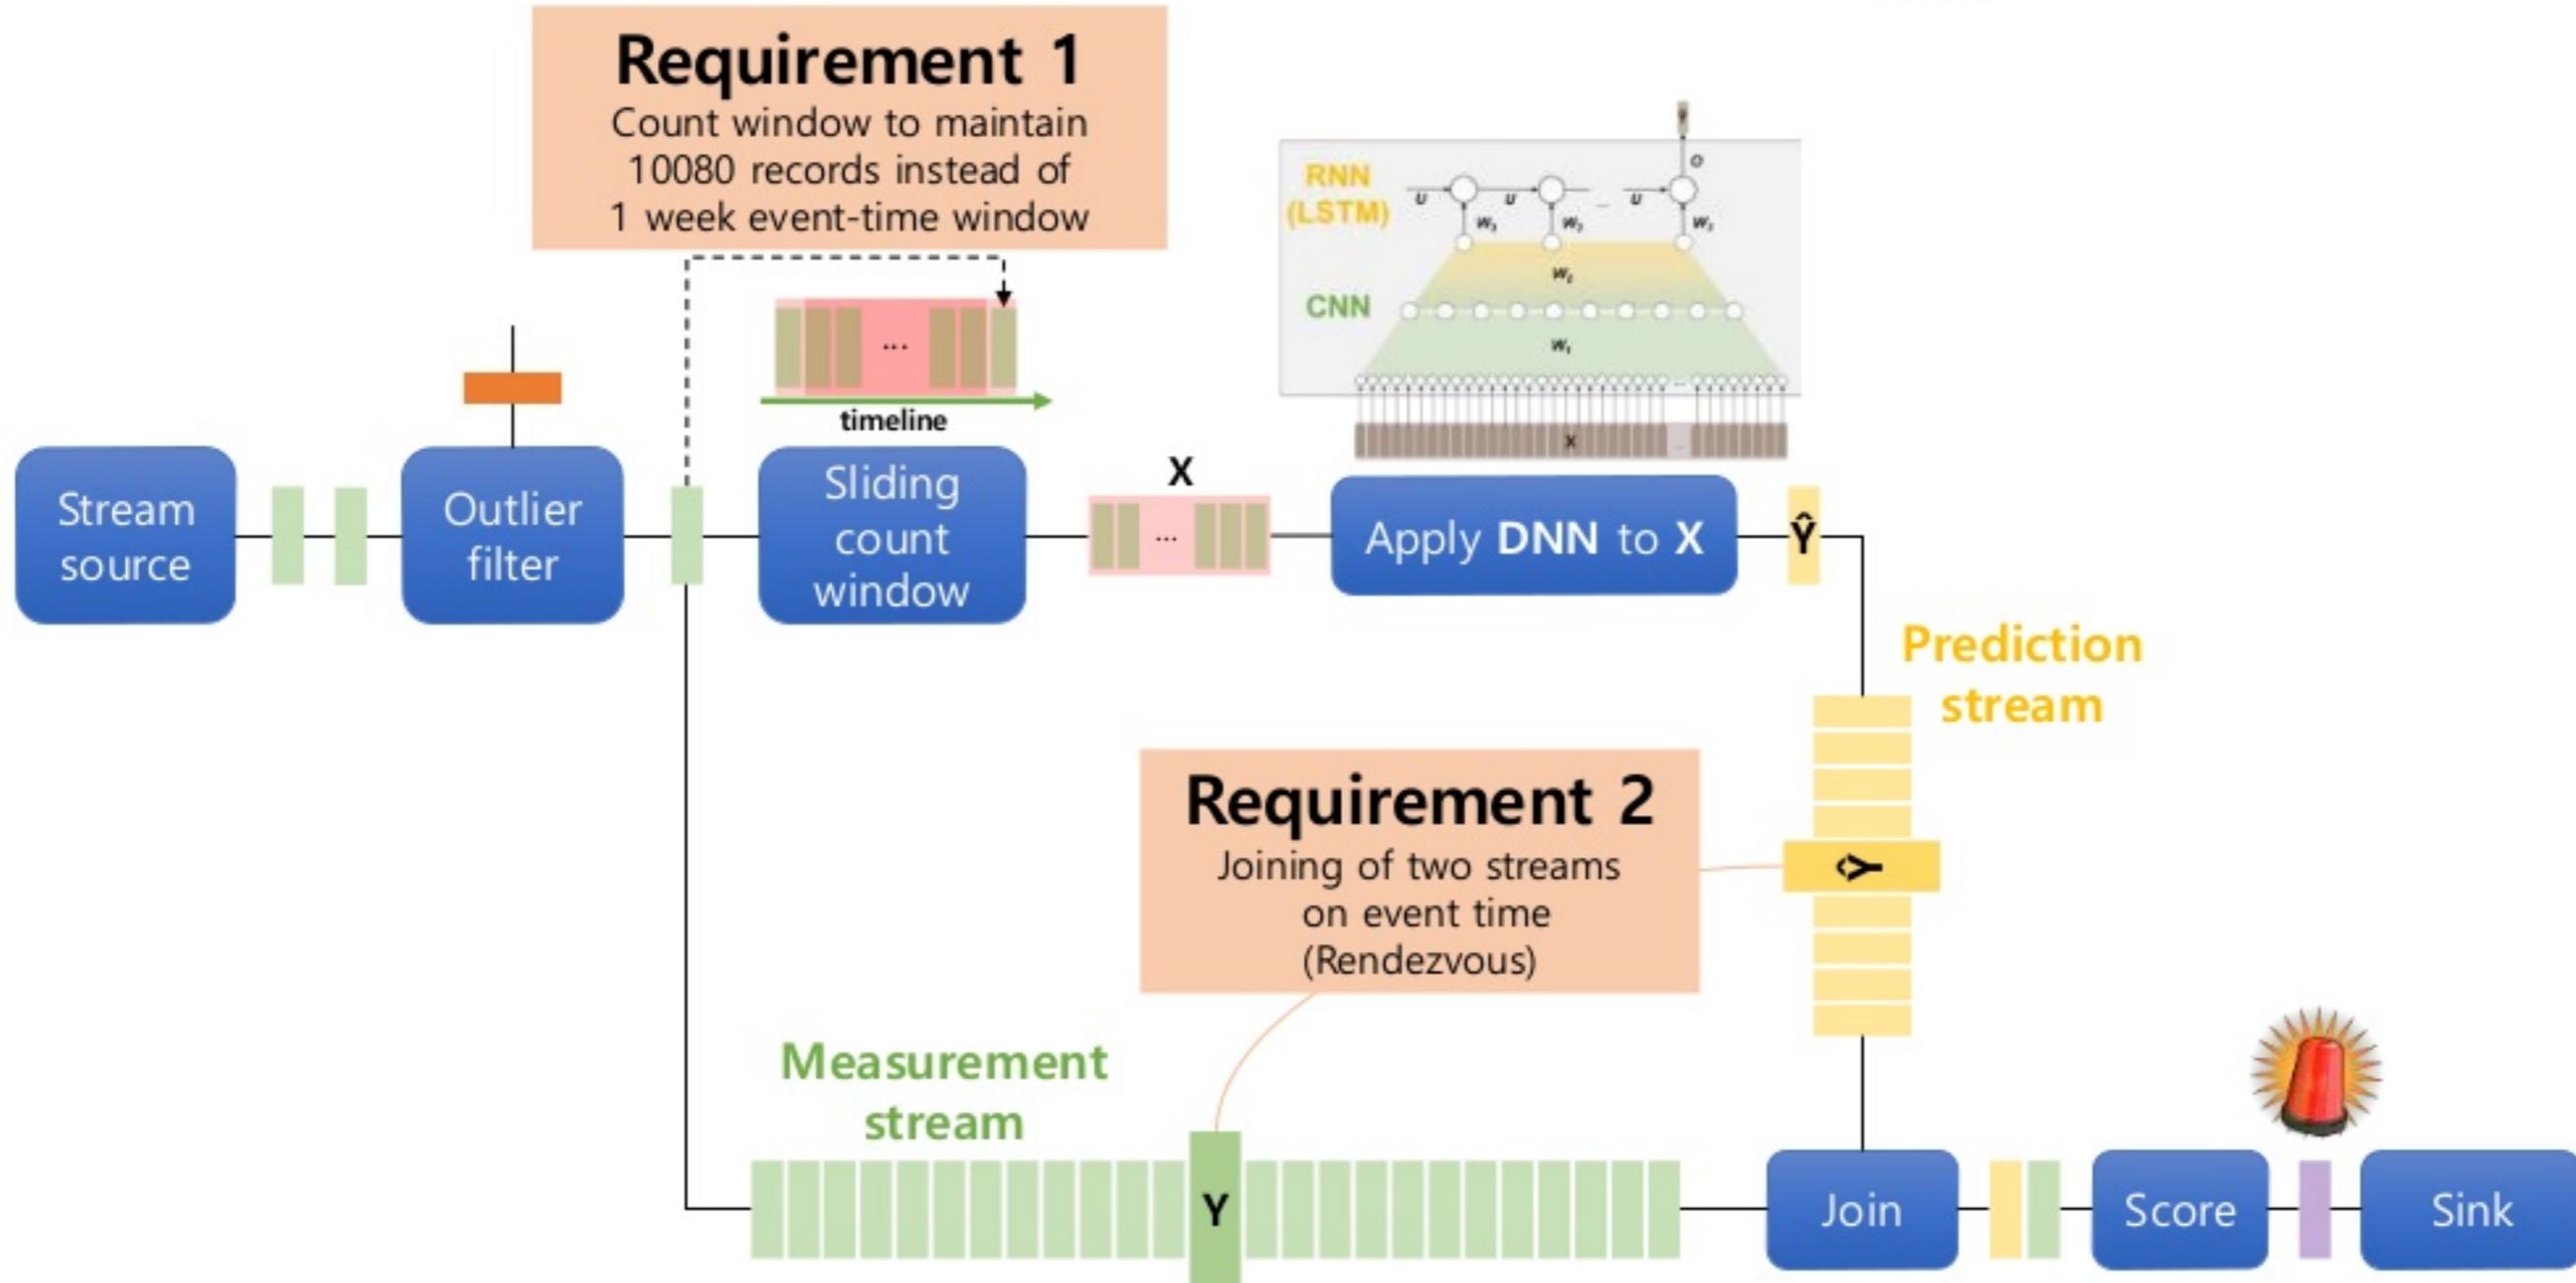




Data engineers apply **Convolutional LSTM** to live sensor data



Desired streaming topology by data engineers



Proof-of-Concept of Spark Structured Streaming

Types

```
case class Input(timestamp: Timestamp, measured: Array[Float])
case class Prediction(timestamp: Timestamp, predicted: Array[Float])
case class Score(timestamp: Timestamp, alarm: Boolean)
```

Input Streaming Dataset

```
val input: Dataset[Input] =
  spark
    .sqlContext
    .readStream
    .option("maxFilesPerTrigger", 1)
    .textFile("data/")
    .withColumn("value", udf(parseRecord).apply($"value"))
    .select("value.*")
    .as[Input]
```

generate an input stream
from local files

Prediction Streaming Dataset

```
val prediction: Dataset[Prediction] =
  input
    .groupBy(window($"timestamp", windowSize, triggerInterval))
    .agg(applyDNN.toColumn.name("predict"))
    .select("predict.*")
    .as[Prediction]
```

apply DNN to 1 week time-series,
not 10080 records
(Sliding count window is not supported)

Score Streaming Dataset

```
val score: Dataset[Score] =
  input
    .join(prediction, "timestamp")
    .withColumn("score", udf(computeScore).apply($"timestamp", $"measured", $"predicted"))
    .select("score.*")
    .as[Score]
```

joining of two streams

Inner join between two streaming Datasets is not supported in Spark Structured Streaming

Score
Dataset

```
val score: Dataset[Score] =  
  input  
    .join(prediction, "timestamp")  
    .withColumn("score", udf(computeScore).apply($"timestamp", $"measured", $"predicted"))  
    .select("score.*")  
    .as[Score]
```

```
Exception in thread "main" org.apache.spark.sql.AnalysisException: Inner join between two streaming DataFrames/Datasets is not supported;;  
Join Inner, (timestamp#10 = timestamp#33)  
:- Project [value#7.timestamp AS timestamp#10, value#7.measured AS measured#11]  
:  +- Project [UDF(value#0) AS value#7]  
:    +- Project [value#0]  
:      +- StreamingRelation DataSource<org.apache.spark.sql.SparkSession@3214bad, text, List(), None, List(), None, Map(maxFilesPerTrigger -> 1, path -> data/), None>, FileSource[data/], [value#0]  
+- Project [predict#28.timestamp AS timestamp#33, predict#28.predicted AS predicted#34]  
  +- Aggregate [window#29], [window#29 AS window#18, predictor(Predictor@74a58a06, Some(createExternalRow(staticInvoke(class org.apache.spark.sql.catalyst.util.DateTimeUtils$, ObjectType(class java.sql.Timestamp))),  
  +- Filter ((timestamp#10 >= window#29.start) && (timestamp#10 < window#29.end))  
  +- Expand [ArrayBuffer(named_struct(start, (((CEIL((cast((preciseTimestamp(timestamp#10) - 0) as double) / cast(60000000 as double))) + cast(0 as bigint)) - cast(10 as bigint)) * 60000000) + 0), e]  
  +- Project [value#7.timestamp AS timestamp#10, value#7.measured AS measured#11]  
  +- Project [UDF(value#0) AS value#7]  
  +- Project [value#0]  
      +- StreamingRelation DataSource<org.apache.spark.sql.SparkSession@3214bad, text, List(), None, List(), None, Map(maxFilesPerTrigger -> 1, path -> data/), None>, FileSource[data/], [value#0]
```

Unsupported Operations on Spark Structured Streaming (v2.2)

Unsupported Operations

There are a few DataFrame/Dataset operations that are not supported with streaming DataFrames/Datasets. Some of them are as follows.

- Multiple streaming aggregations (i.e. a chain of aggregations on a streaming DF) are not supported.
- Limit and take first N rows are not supported on streaming Datasets.
- Distinct operations on streaming Datasets are not supported.
- Sorting operations are supported on streaming Datasets only after an aggregation and in Complete Output Mode.
- Outer joins between a streaming and a static Datasets are conditionally supported.
 - Full outer join with a streaming Dataset is not supported
 - Left outer join with a streaming Dataset on the right is not supported
 - Right outer join with a streaming Dataset on the left is not supported
- Any kind of joins between two streaming Datasets is not yet supported.

Requirement 1
Count window to maintain 10080 records instead of 1 week event-time window

Requirement 2
Joining of two streams on event time (Rendezvous)

That's why we move to Flink DataStream API



Spark Structured Streaming

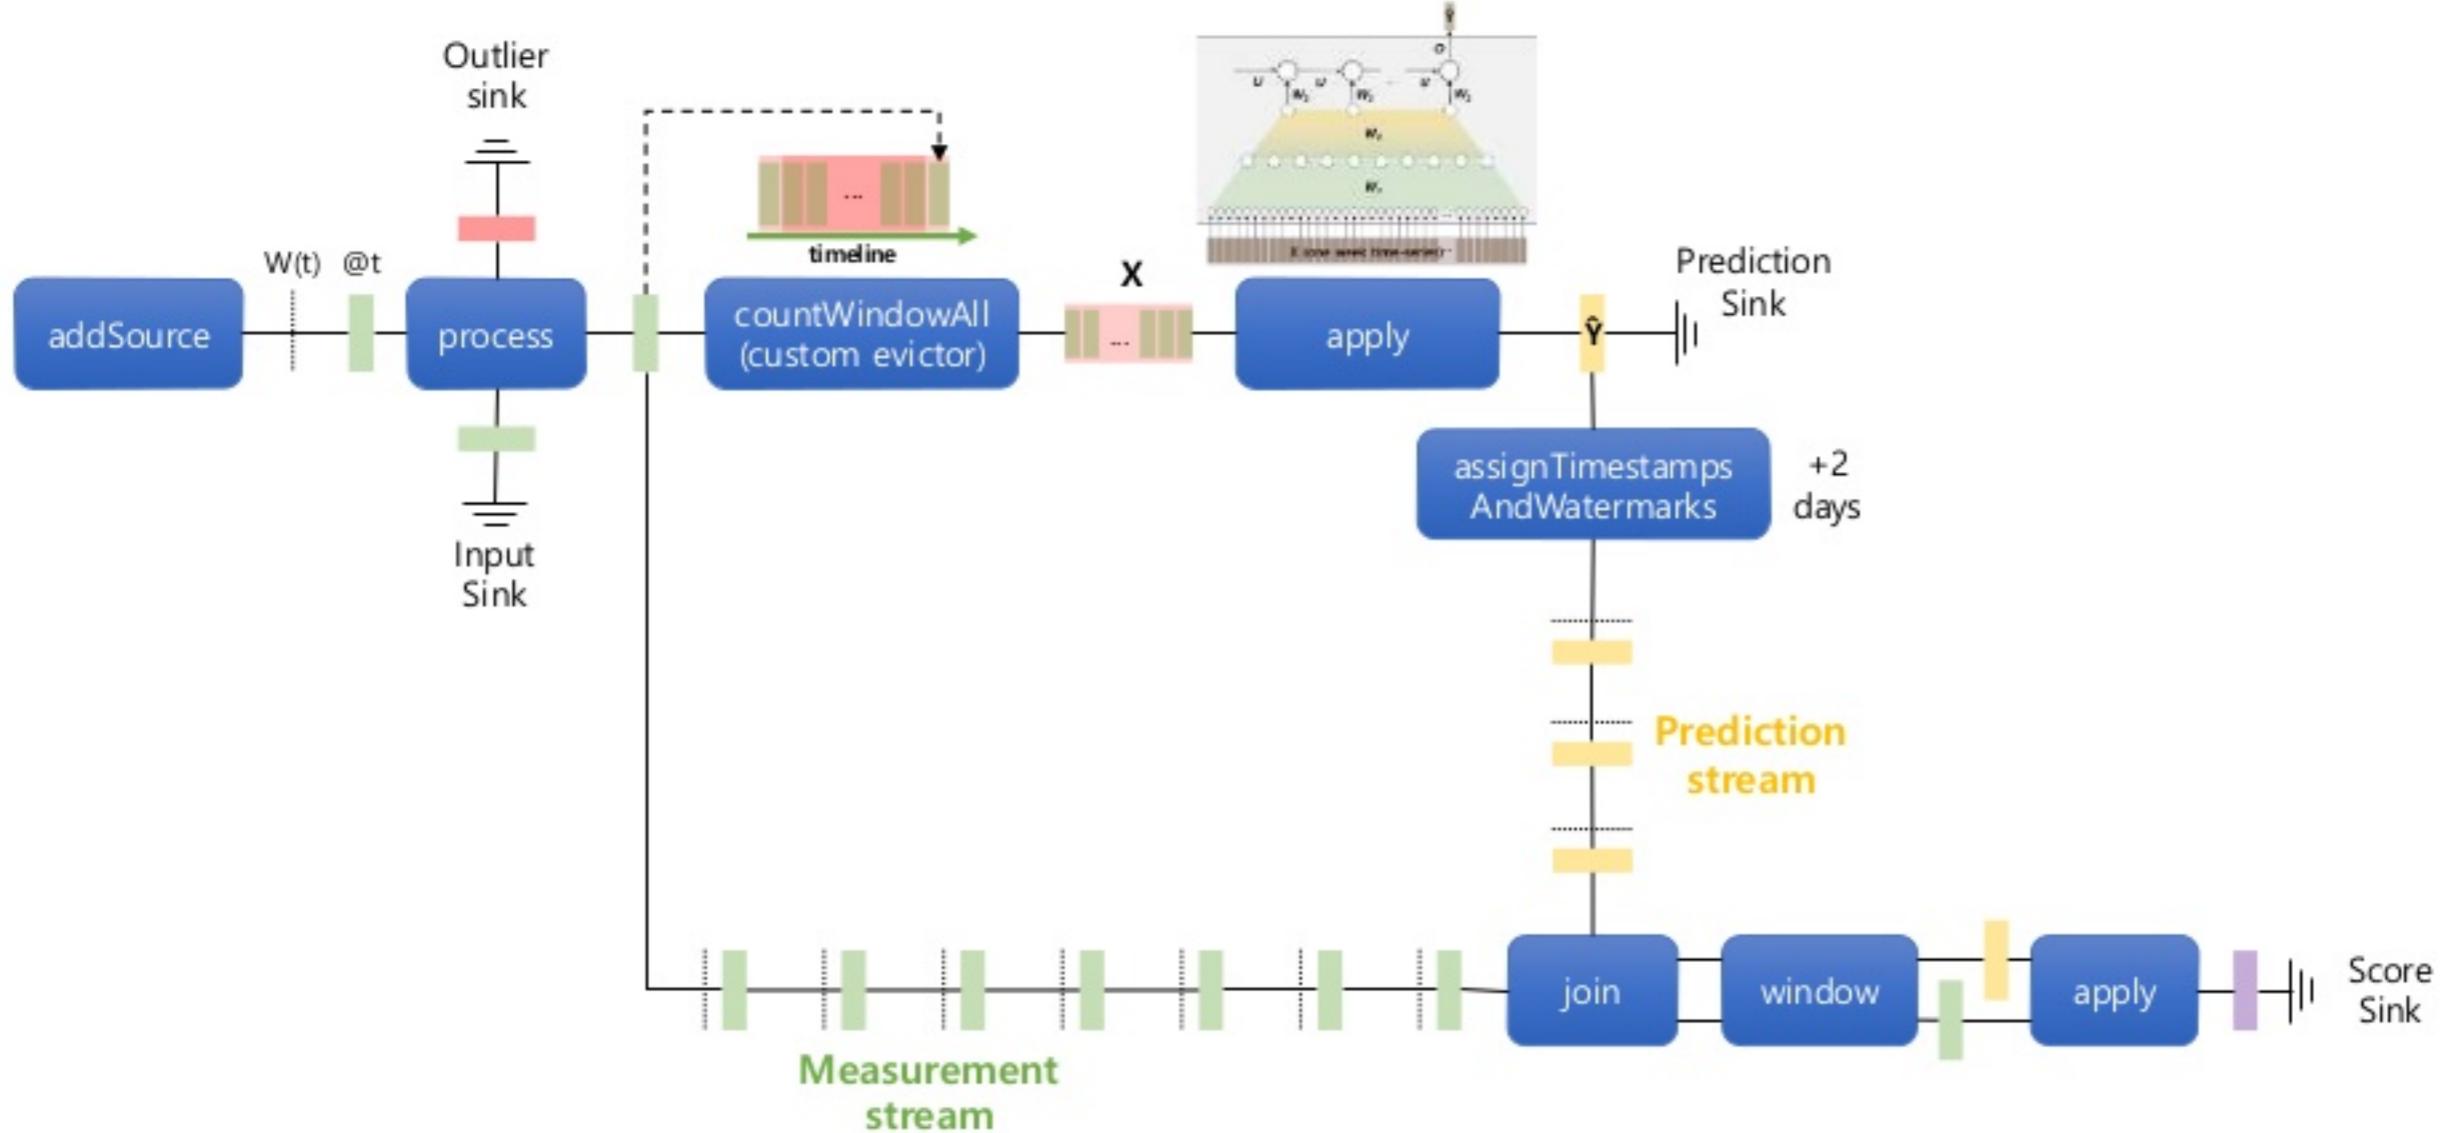
- Sliding count window : **not supported**
- Joining of two streams : **not supported**
- Micro-batch behind the scene
 - Continuous processing proposed in SPARK-20928

* it could be possible to use our **Convolutional LSTM** model using Spark Structured Stream in some other way

Flink DataStream API

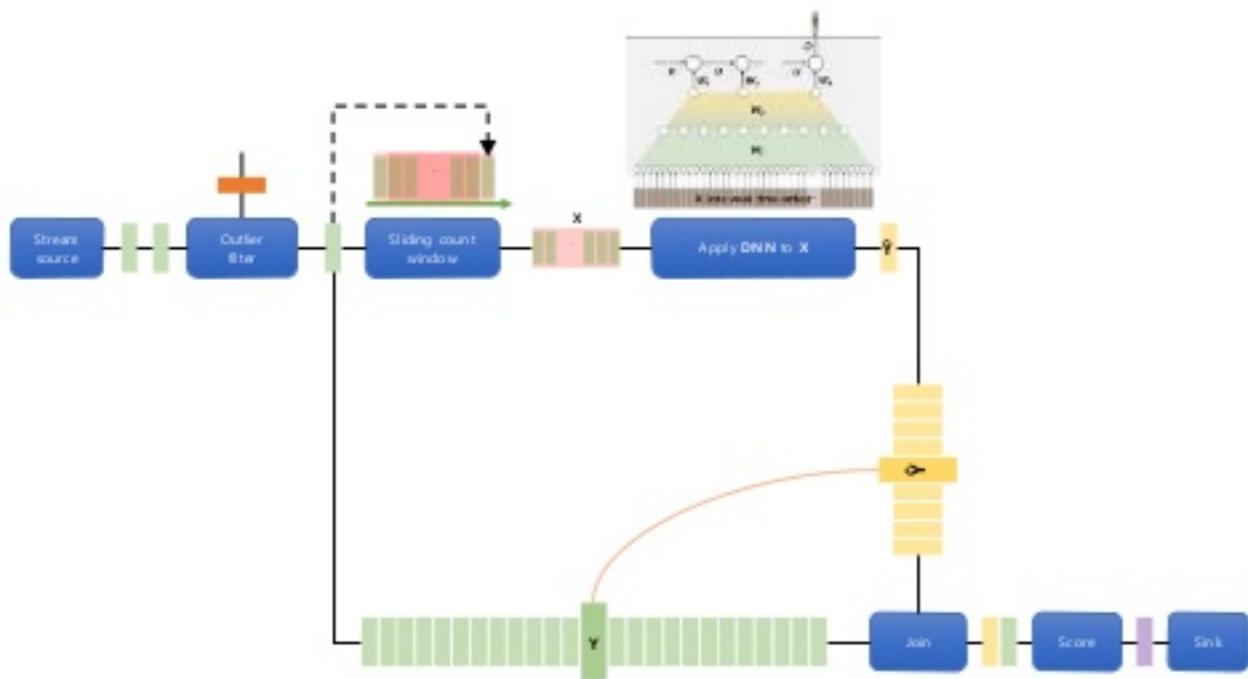
- Sliding count window : **supported**
- Joining of two streams : **supported**
- Scalability and performance proved by other use cases

Data processing pipeline with Flink DataStream API

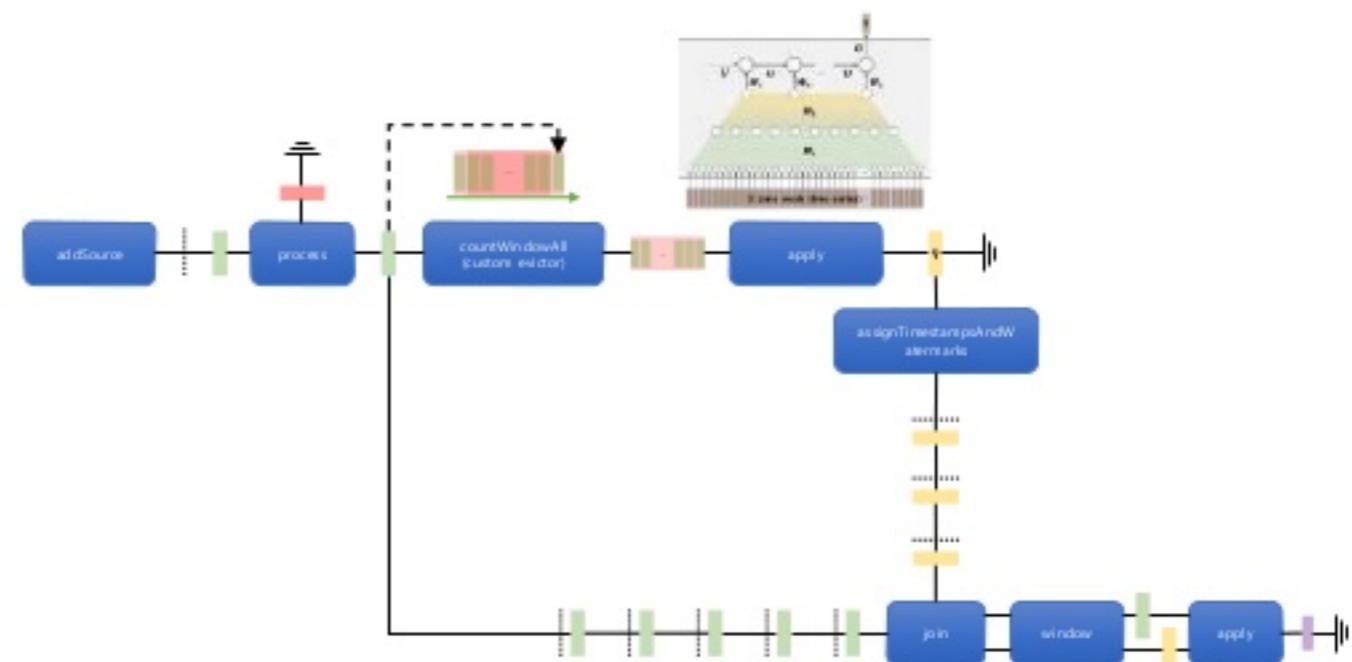


Flink can faithfully implement our streaming topology design

<Topology design>



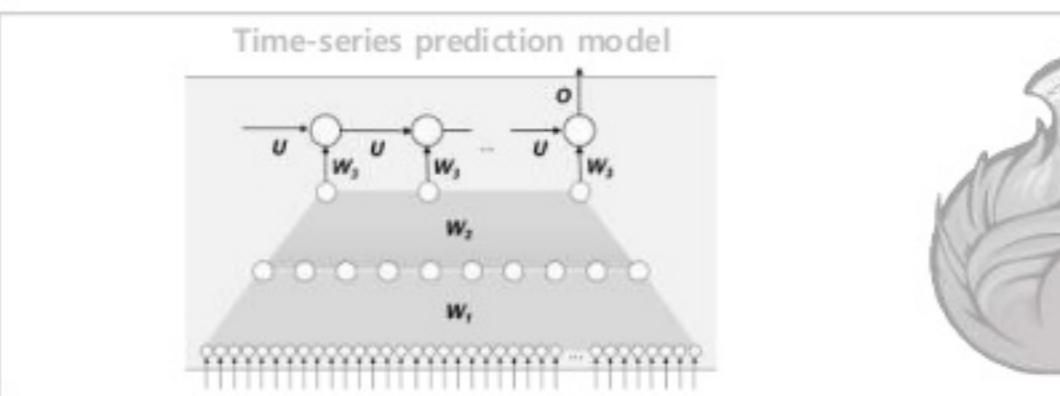
<Flink implementation>



Contents

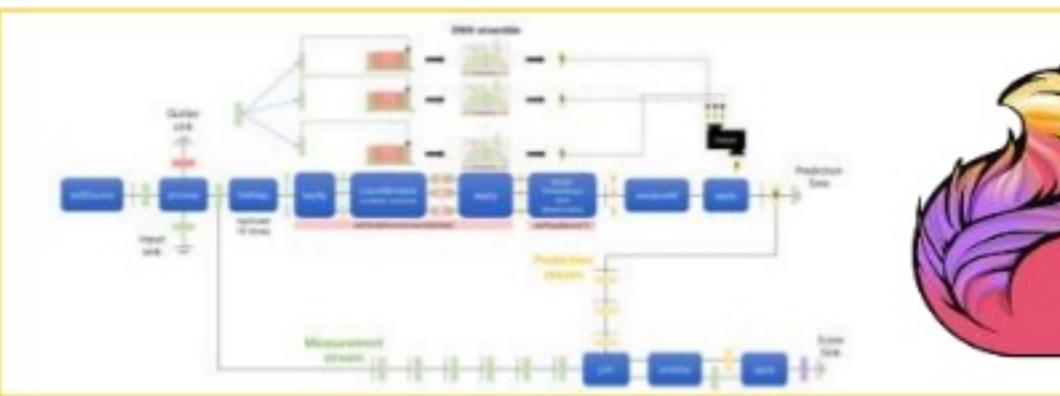
1

Why we use Flink
for our time-series prediction model



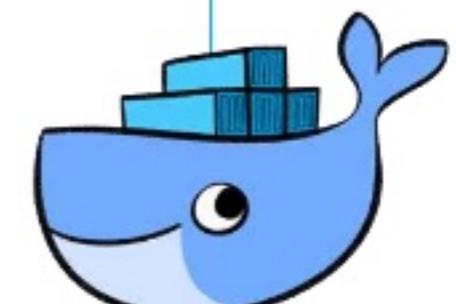
2

Flink pipeline design for
rendezvous and DNN ensemble



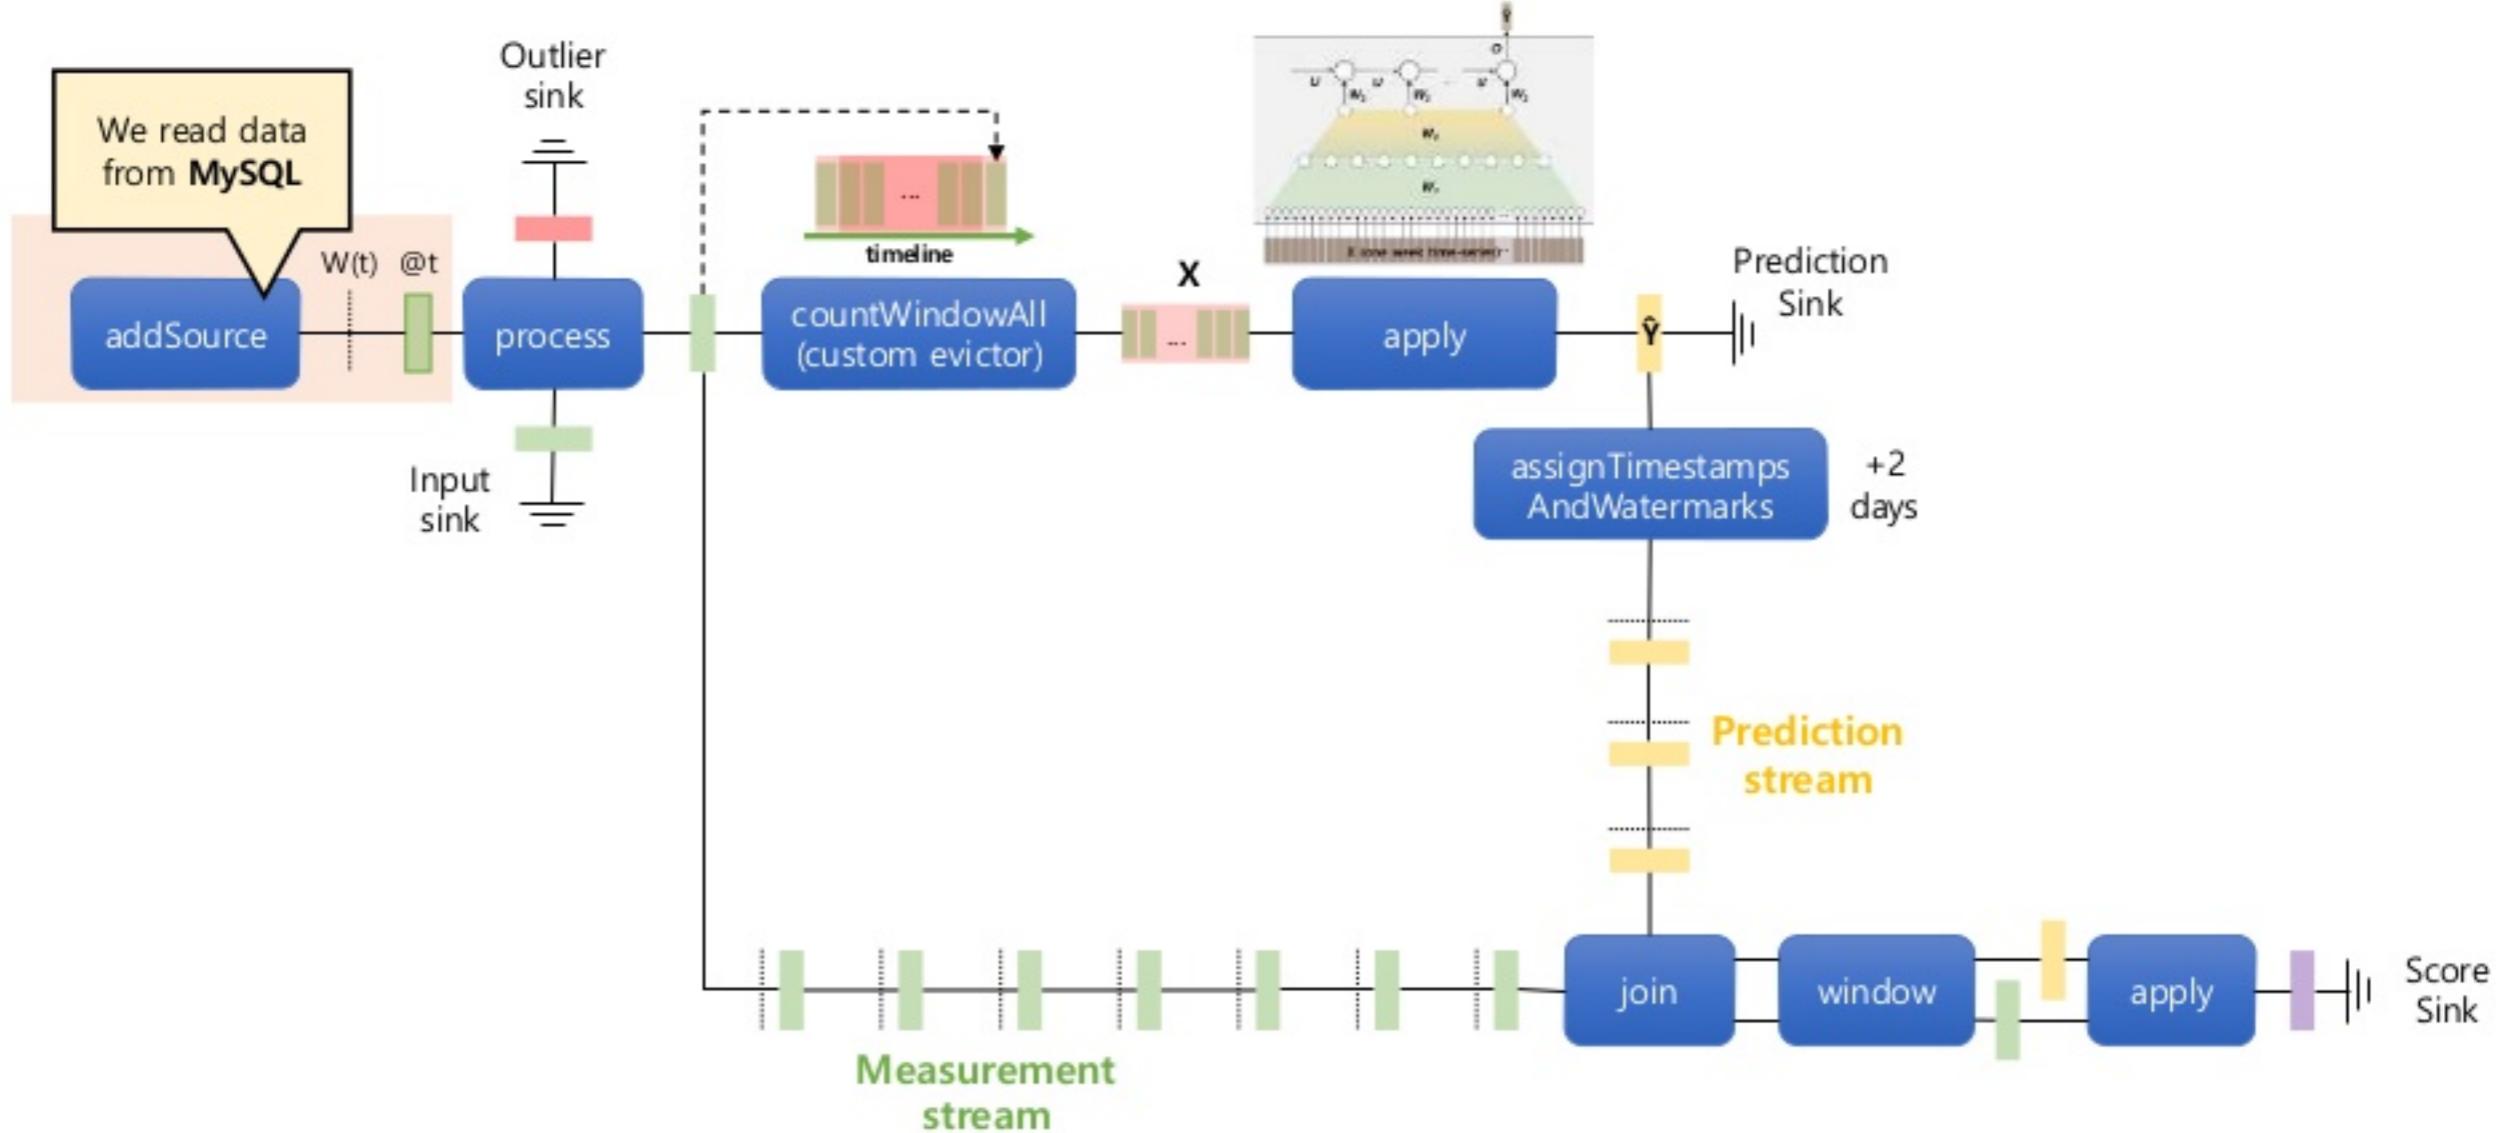
3

Solution packaging and monitoring
with Docker and Prometheus

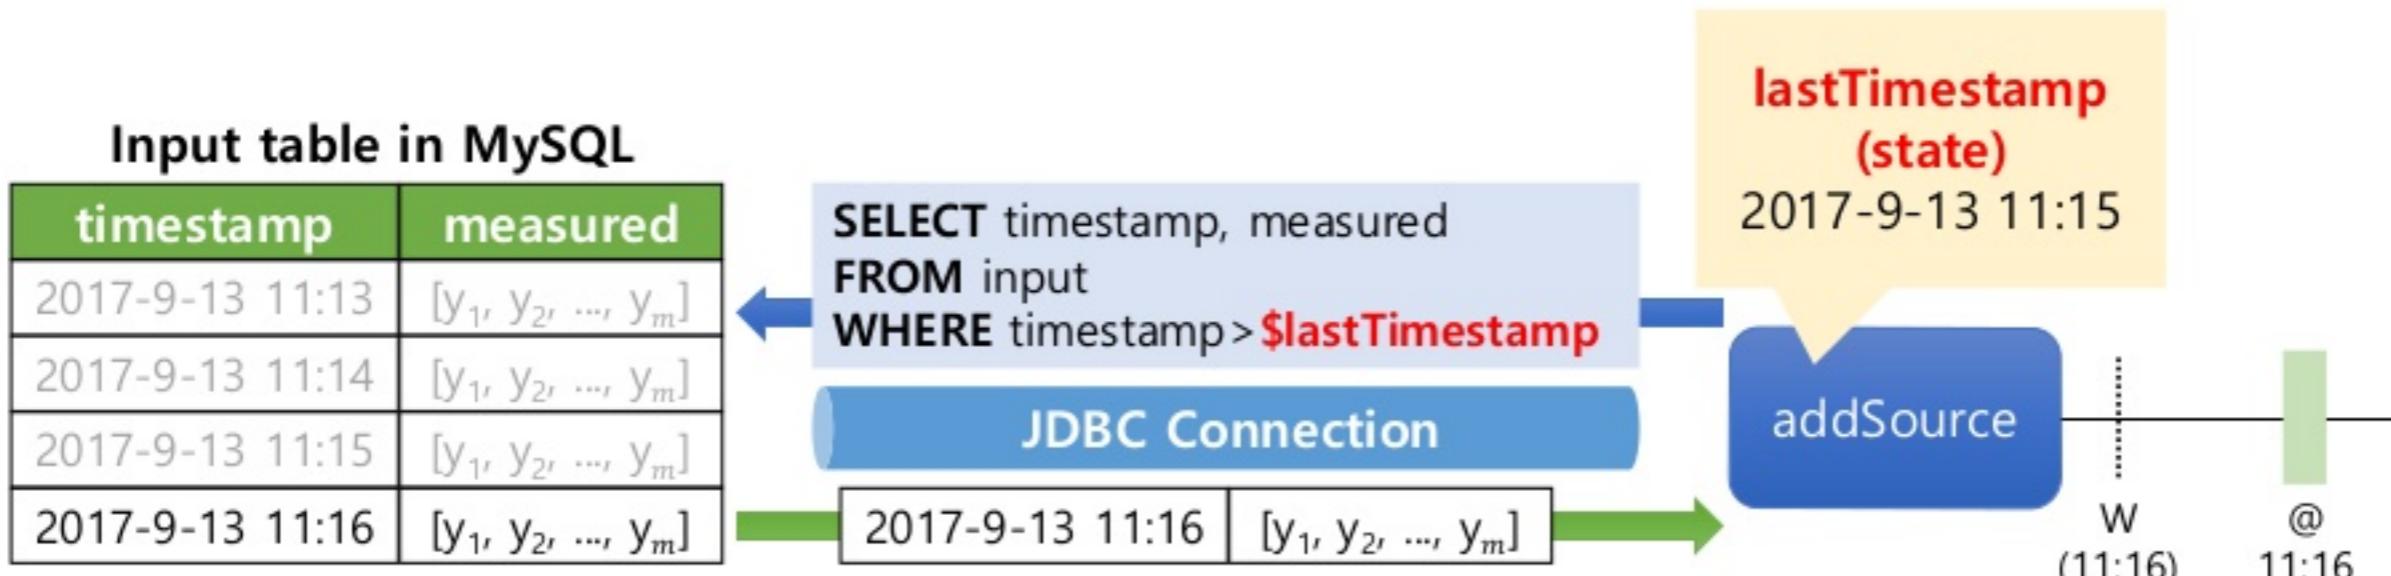




Data processing pipeline with **Flink DataStream API**



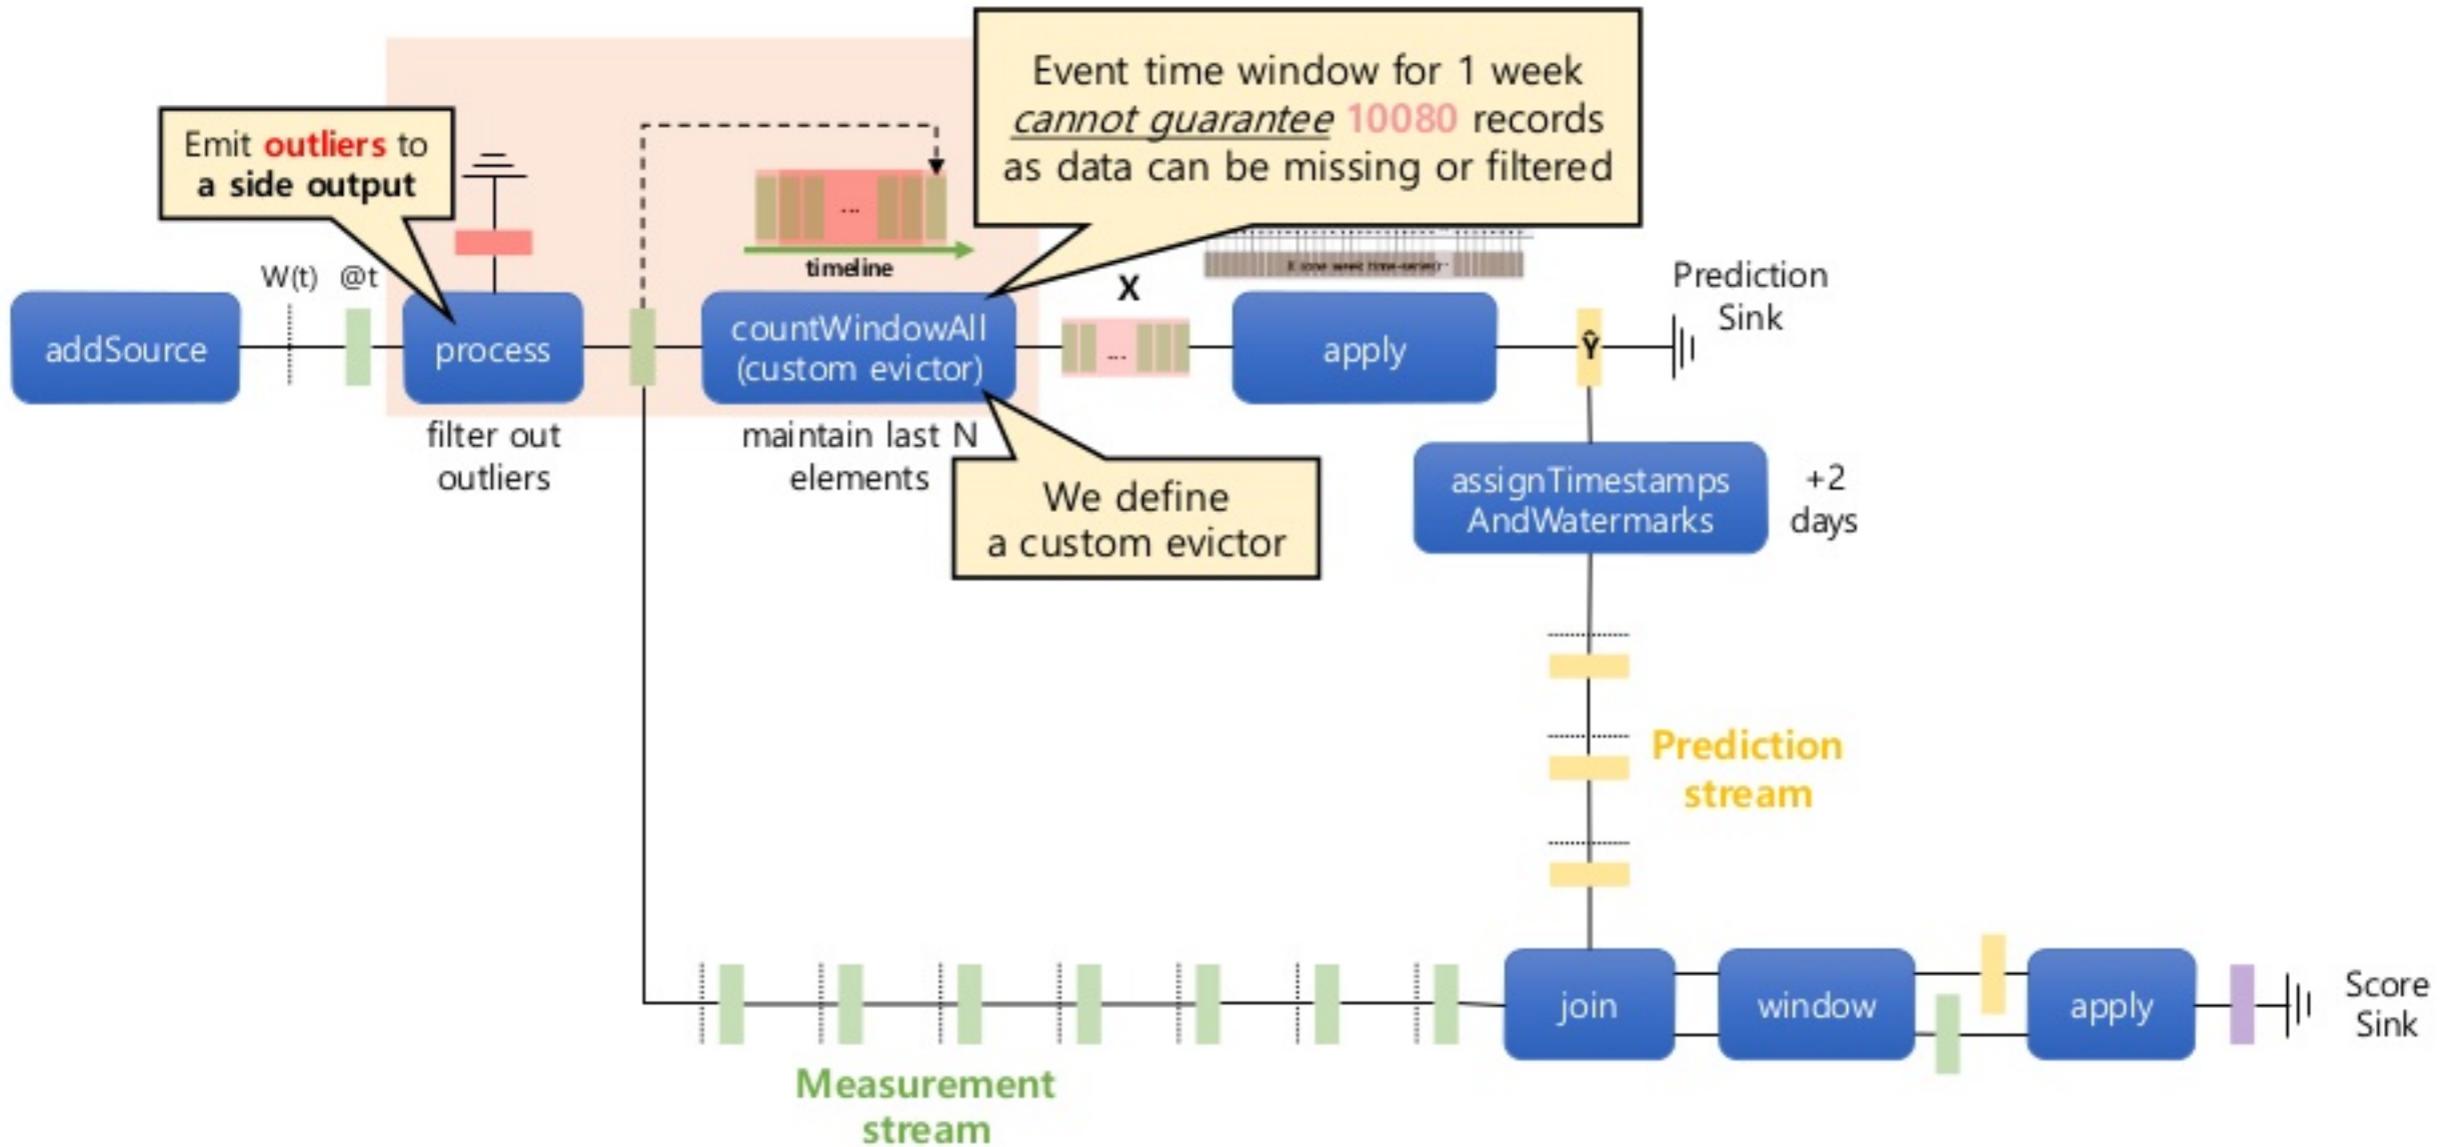
Stateful custom source to read from MySQL



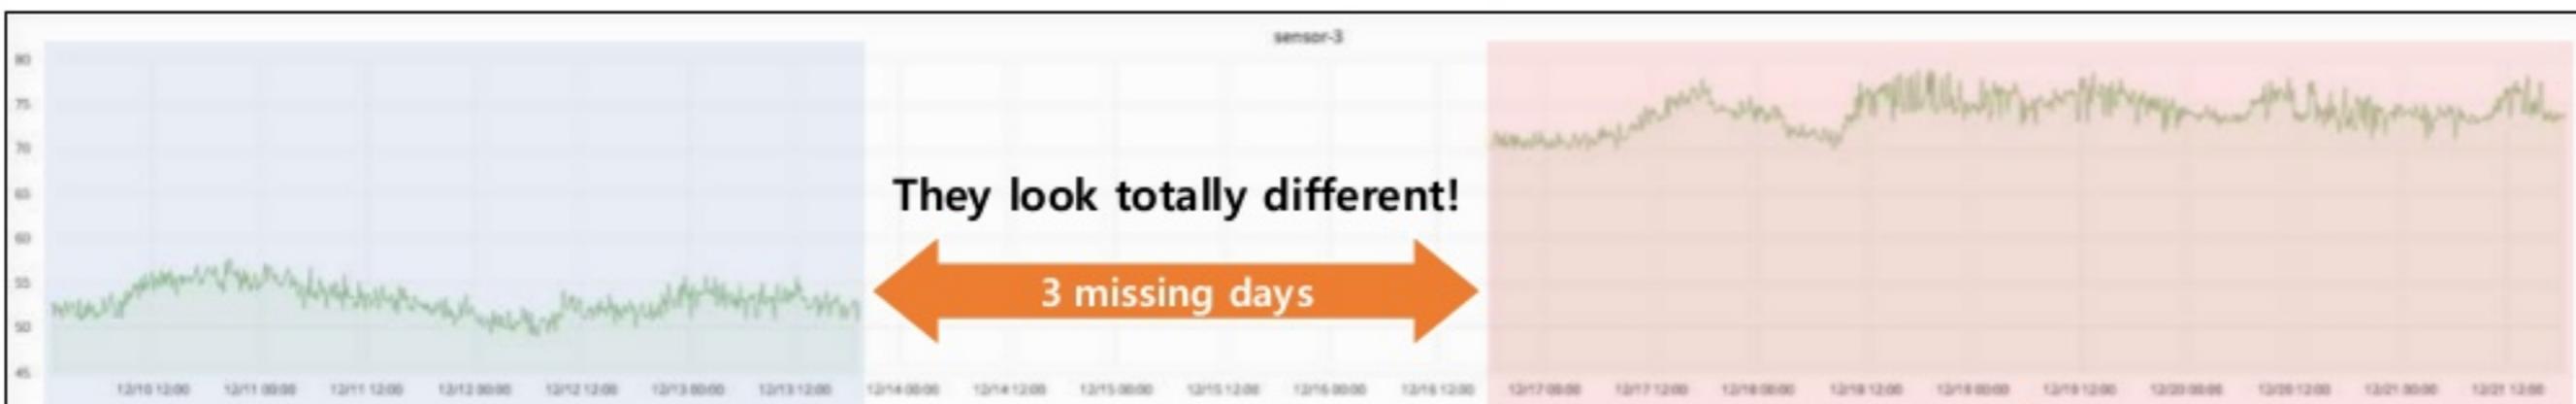
- We assume that sensor data arrive in order
 - Emit an input record and a watermark of the same time
 - Increase **lastTimestamp** afterward (11:15 → 11:16)
- Exactly-once semantics
 - Store **lastTimestamp** when taking a snapshot
 - Restore **lastTimestamp** when restarted



Data processing pipeline with Flink DataStream API

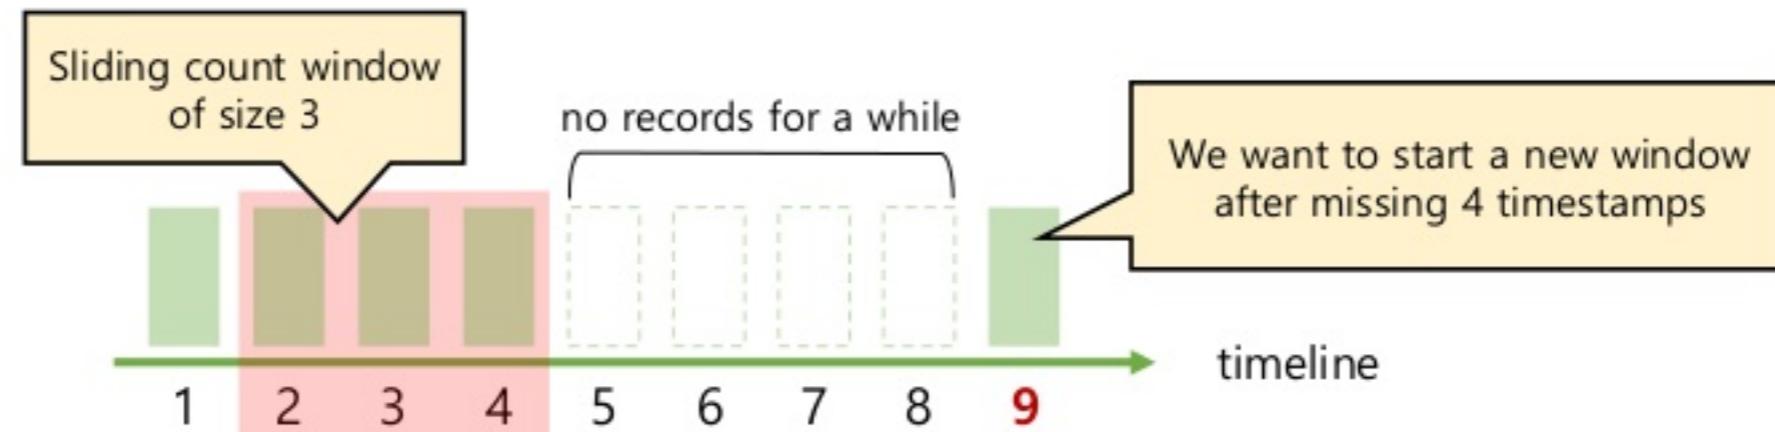


What if data is **absent** or **filtered** for a long period of time?

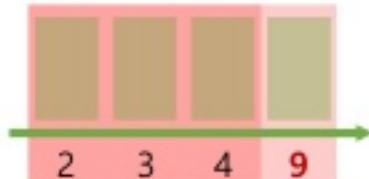


We'd better start a new sliding window
for the time-series!

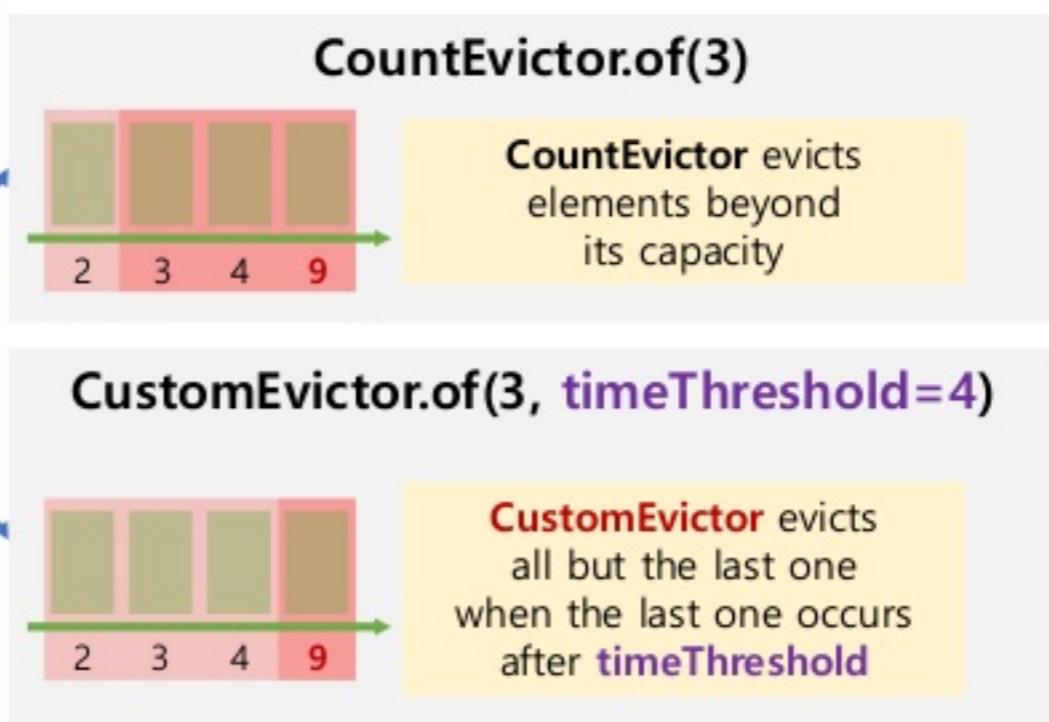
How to start a new sliding count window after a long break



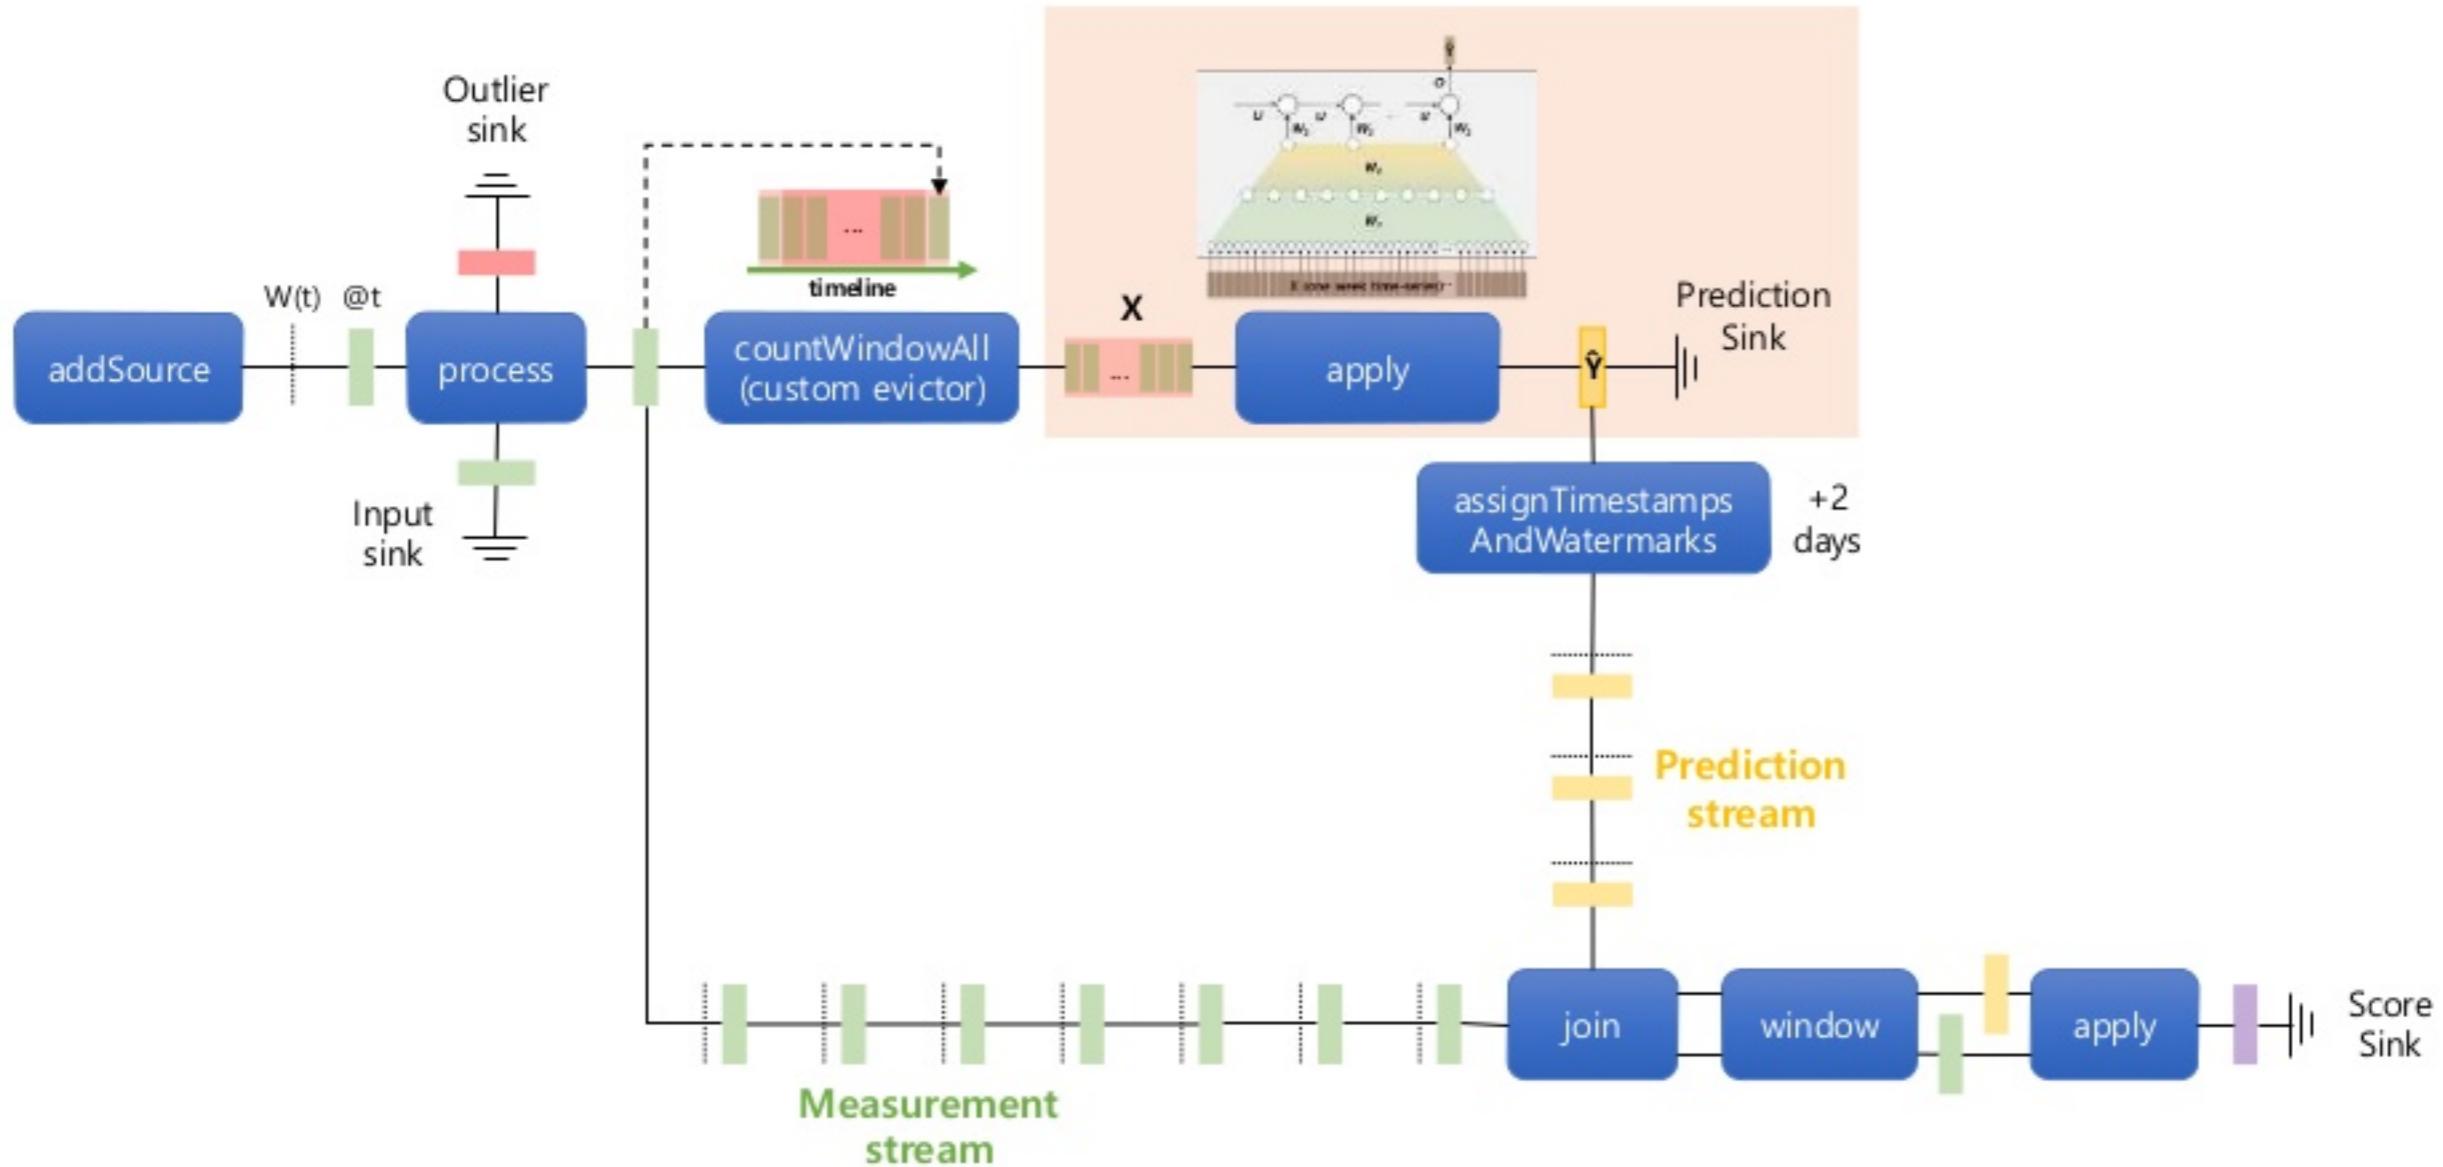
EvictingWindowOperator
adds a new input record to
InternalListState



CountTrigger.of(1)
fires every time
a record comes in



Data processing pipeline with Flink DataStream API



Working with model developers

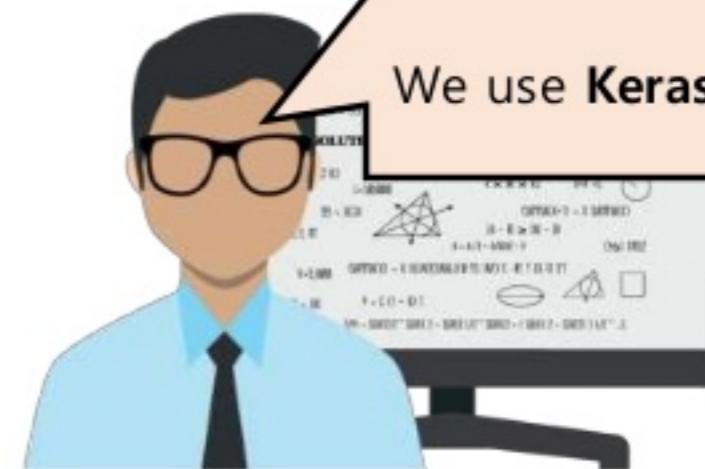
I want to develop our solution on **JVM!**

Why don't we develop models using **Deeplearning4J?**



I don't want to use **Deeplearning4J** because that's **Java...**

We use **Keras** on **Python!**



They stick to using **Python**

They develop models using a Python library called  **Keras**

How to load Keras models in Flink?

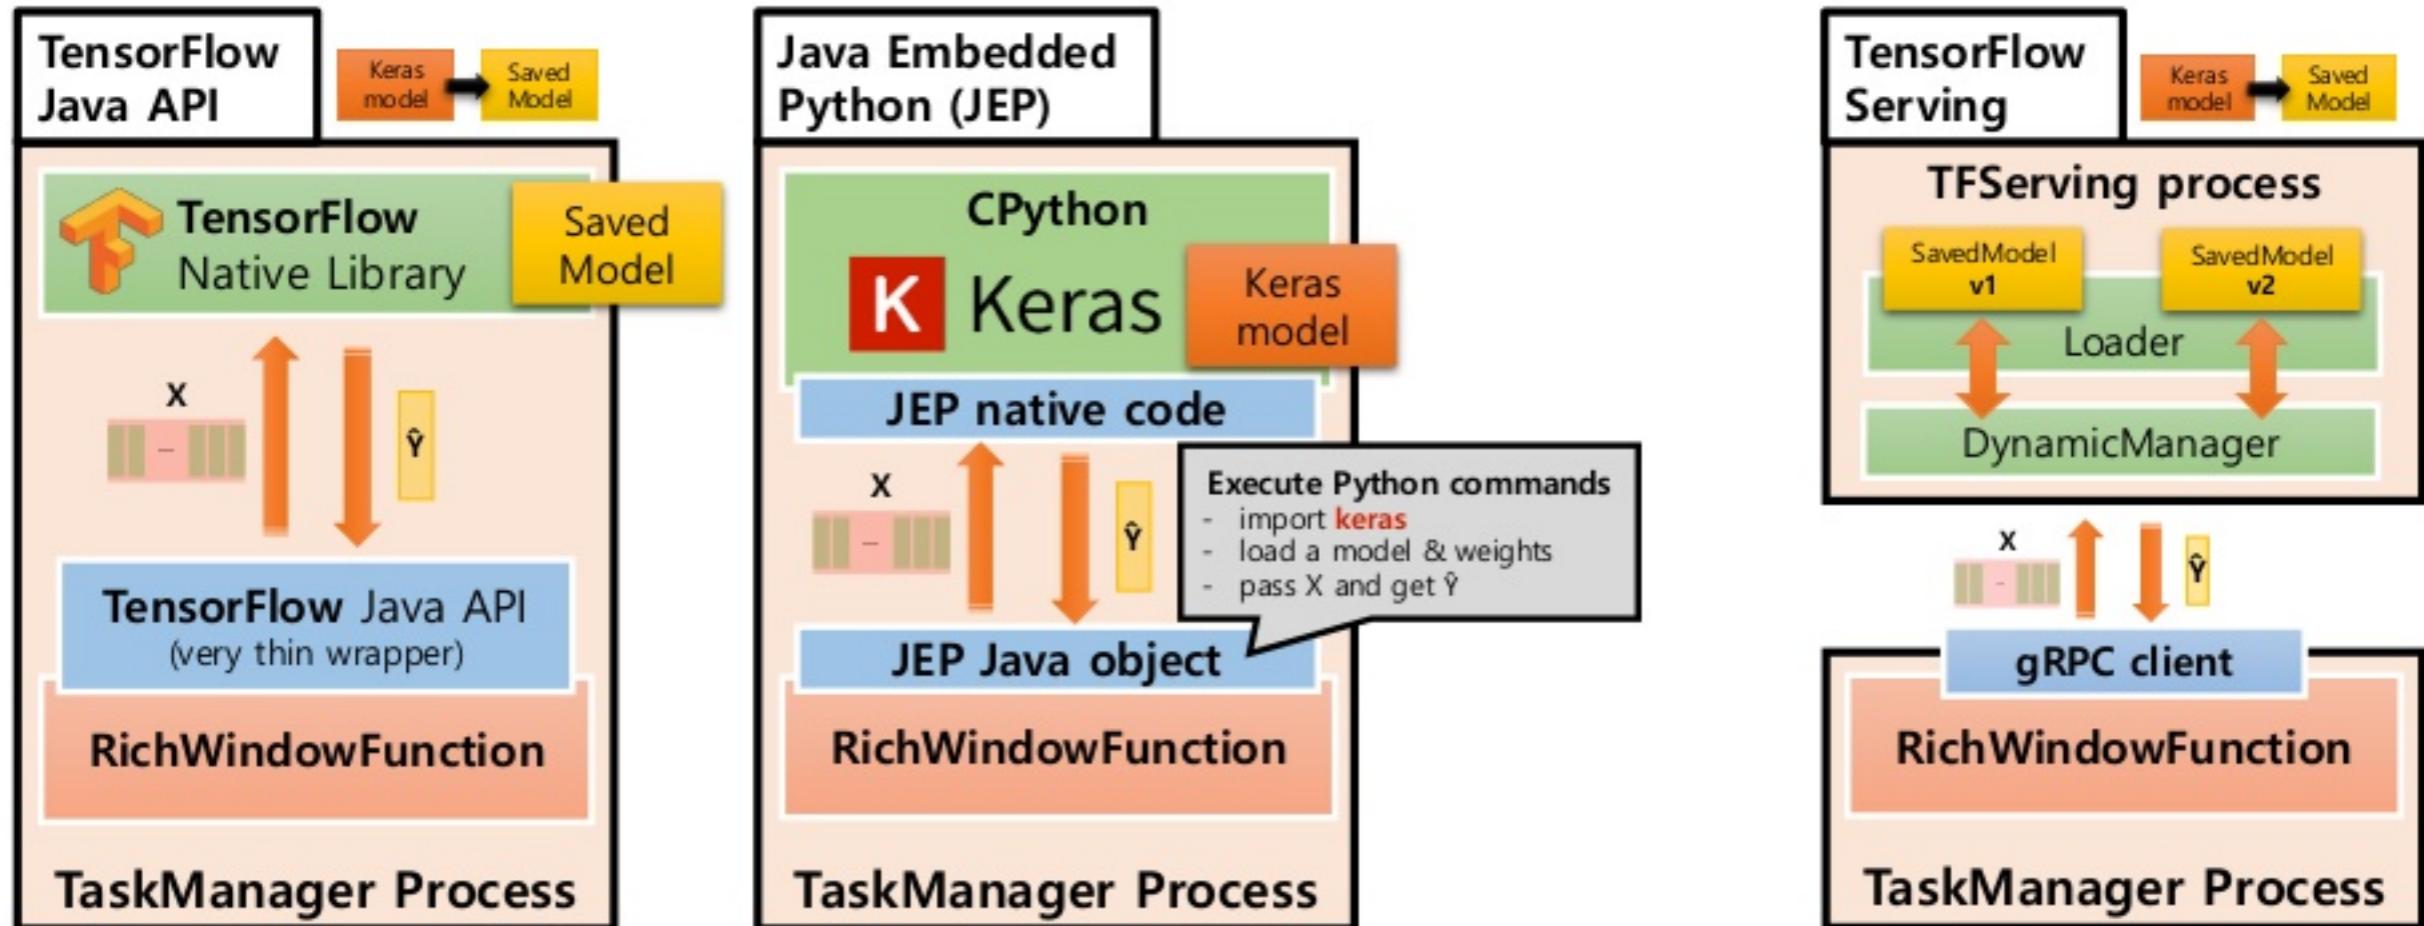


K Keras

Loading **Keras** models in JVM

- Convert **Keras** models to **TensorFlow SavedModel**
 - use tensorflow.python.saved_model.builder.**SavedModelBuilder**
 - **TensorFlow Java API** (Flink TensorFlow)
 - Do inference **inside** the JVM process
 - **TensorFlow Serving**
 - Do inference **outside** the JVM process
- Execute **Keras** through CPython inside JVM
 - Do inference **inside** the JVM process
 - **Java Embedded Python (JEP)** to ease the use of **CPython**
 - <https://github.com/ninia/jep>
- Use **KerasModellImport** from **Deeplearning4J**
 - Not mature enough

Comparison of approaches to use **Keras** models in **JVM**



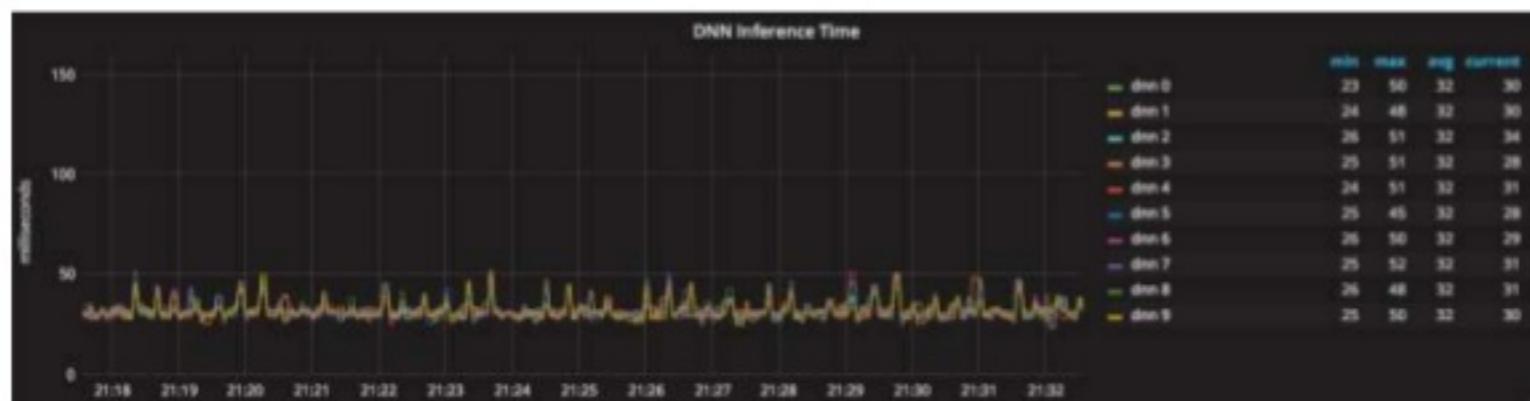
Comparison of runtime inference performance

(* We do not batch inference calls)

Keras inside CPython w/ TensorFlow backend

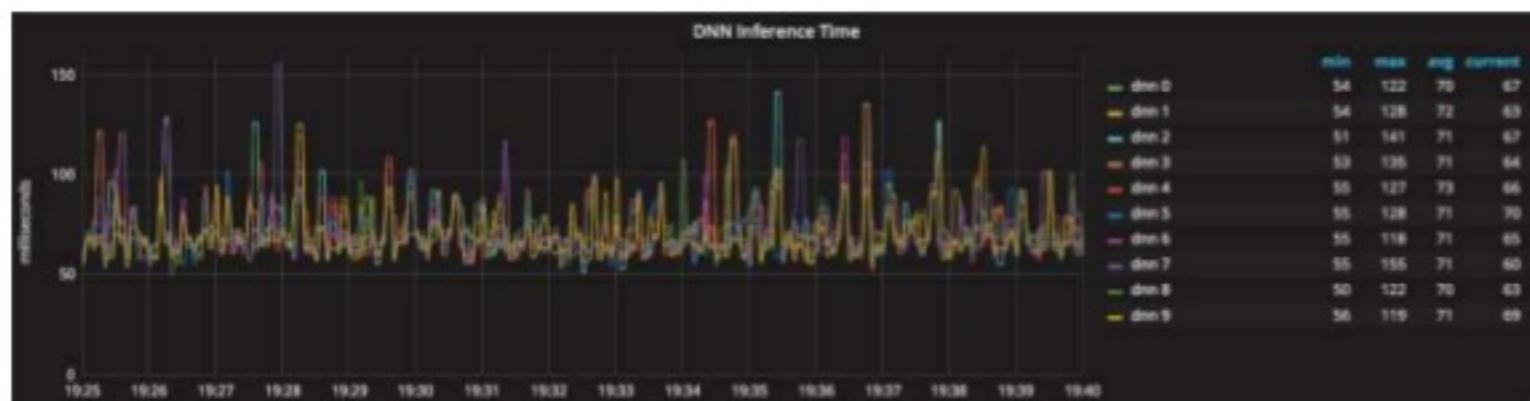
32 milliseconds per inference

(* Theano backend is extremely slow in our case)



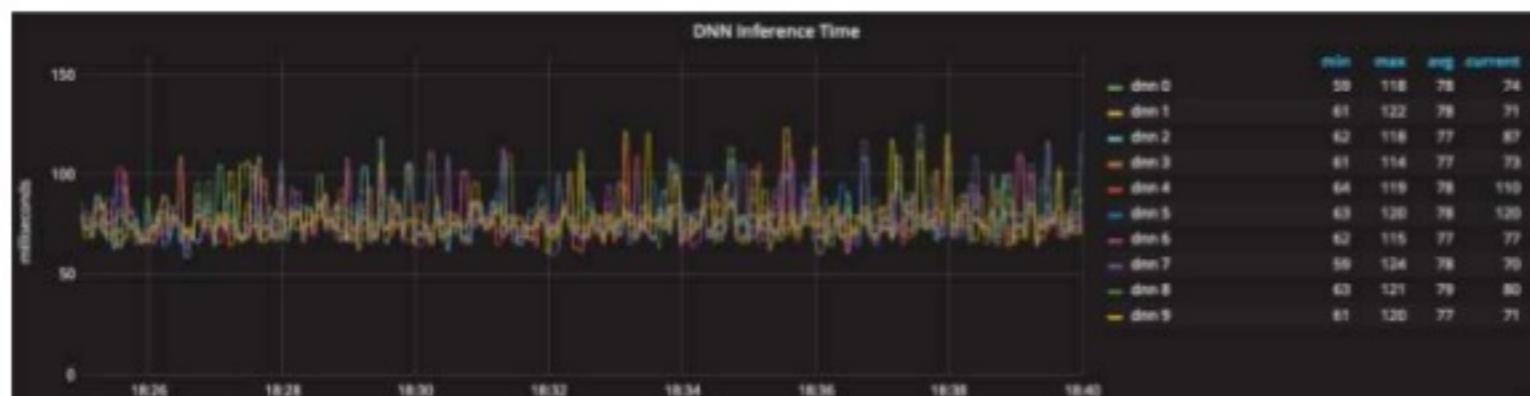
TensorFlow Serving

71.2 milliseconds per inference

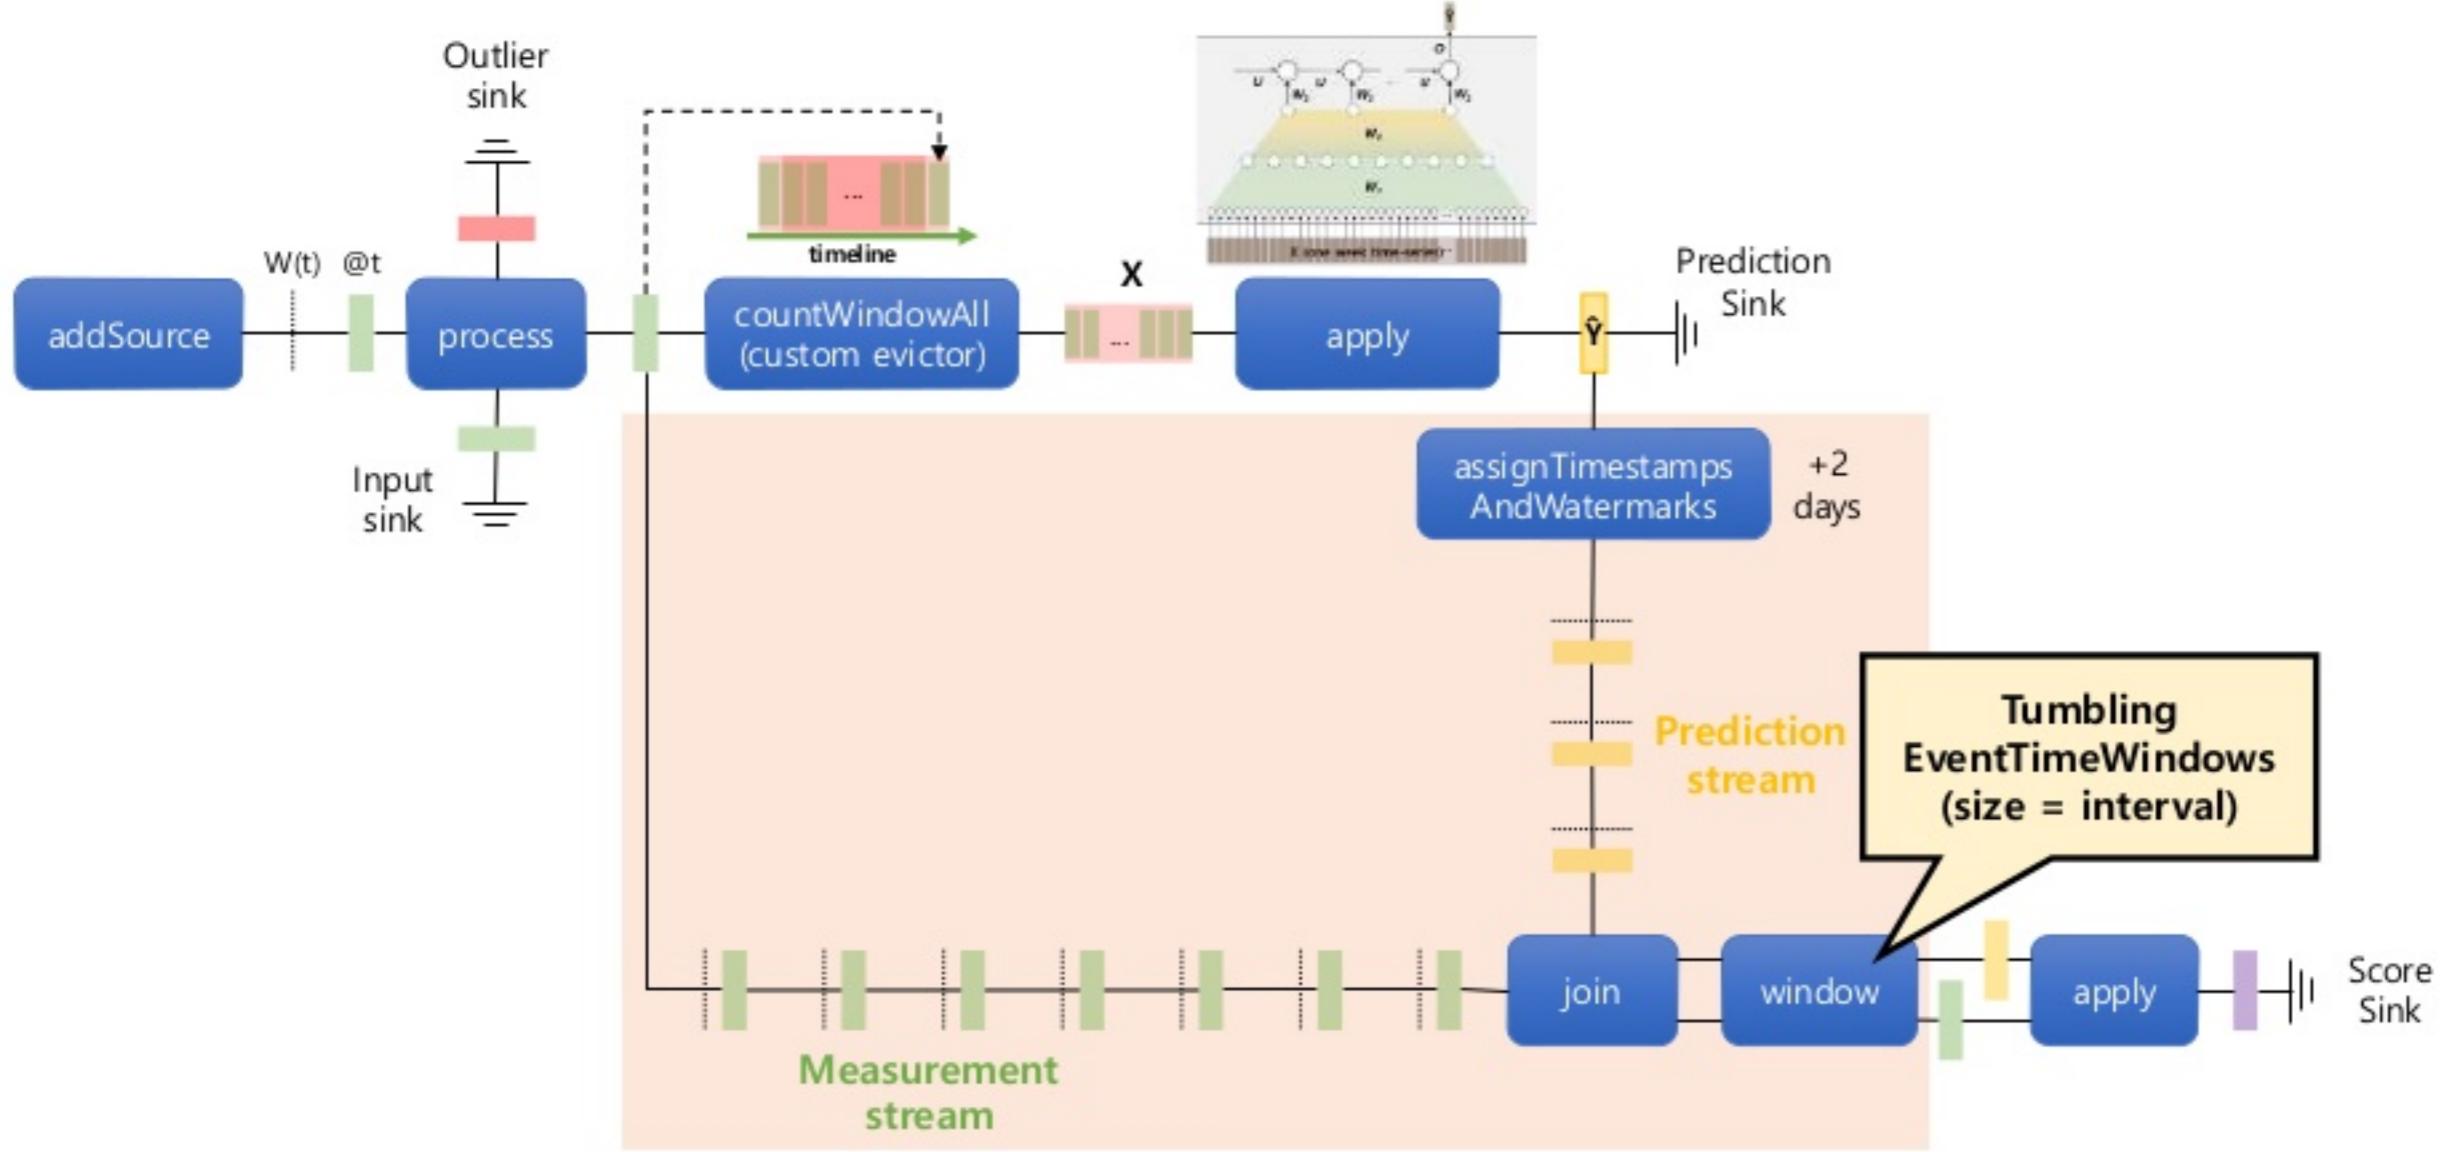


TensorFlow Java API

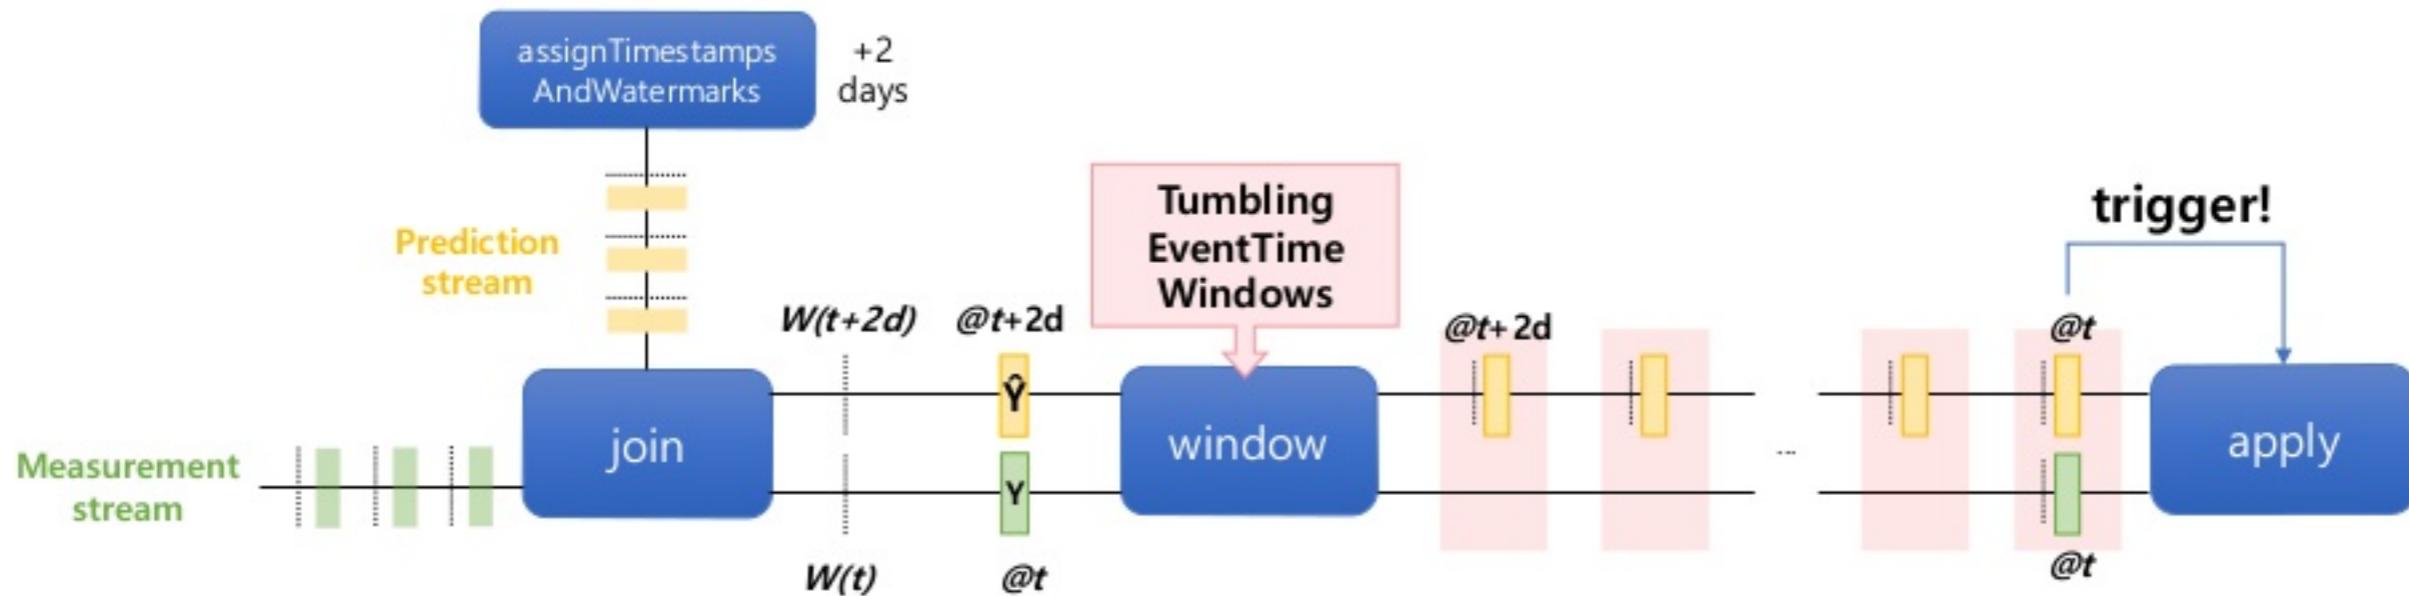
77.7 milliseconds per inference



Data processing pipeline with Flink DataStream API

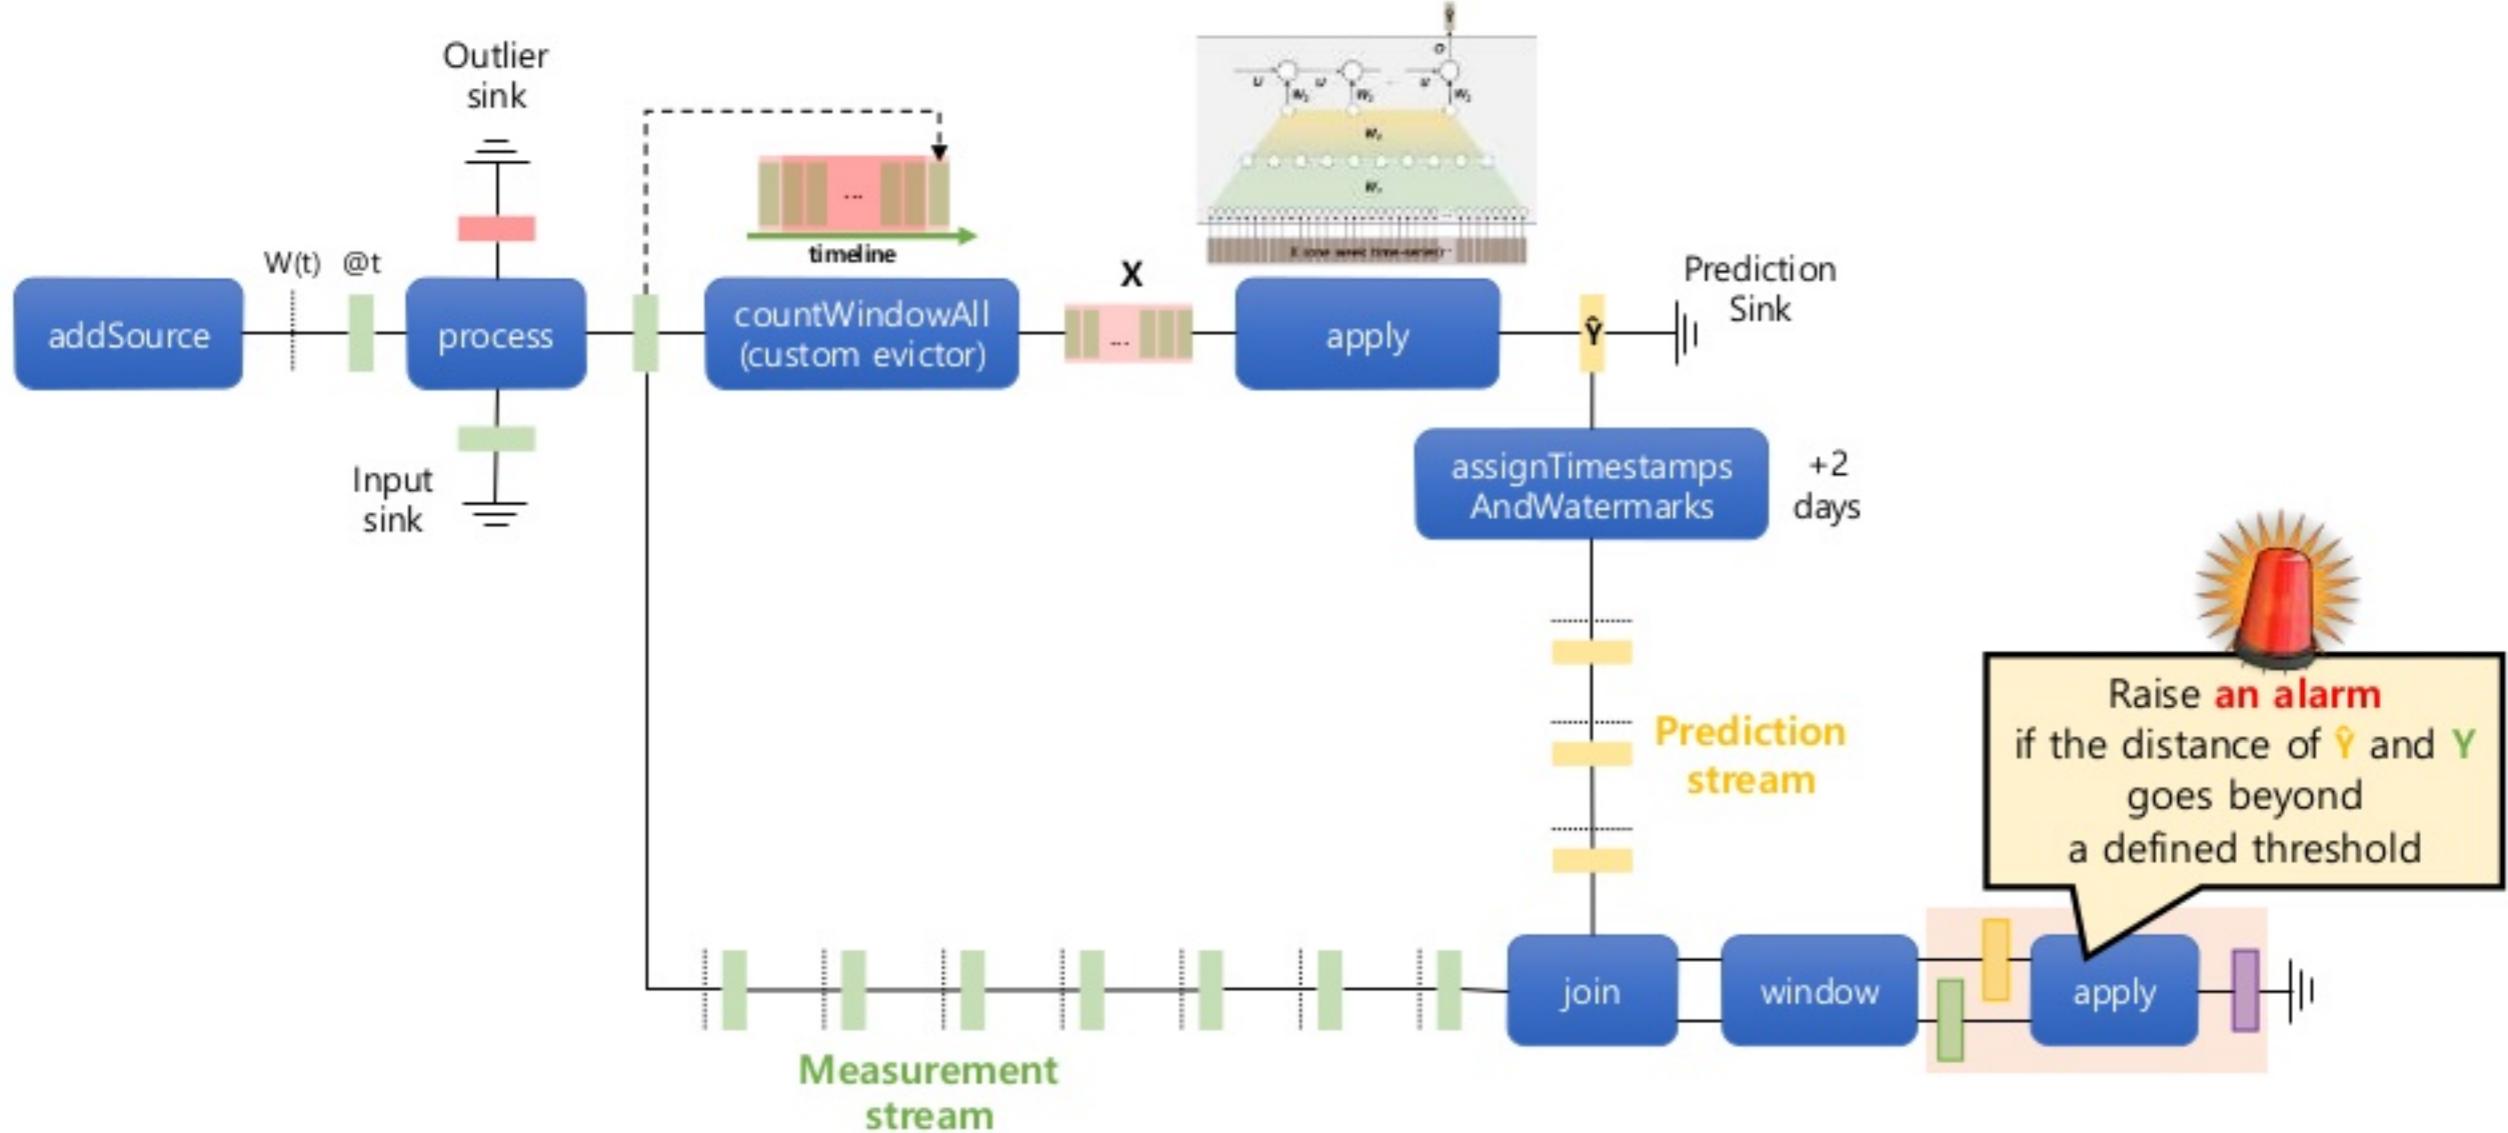


Joining two streams on event time

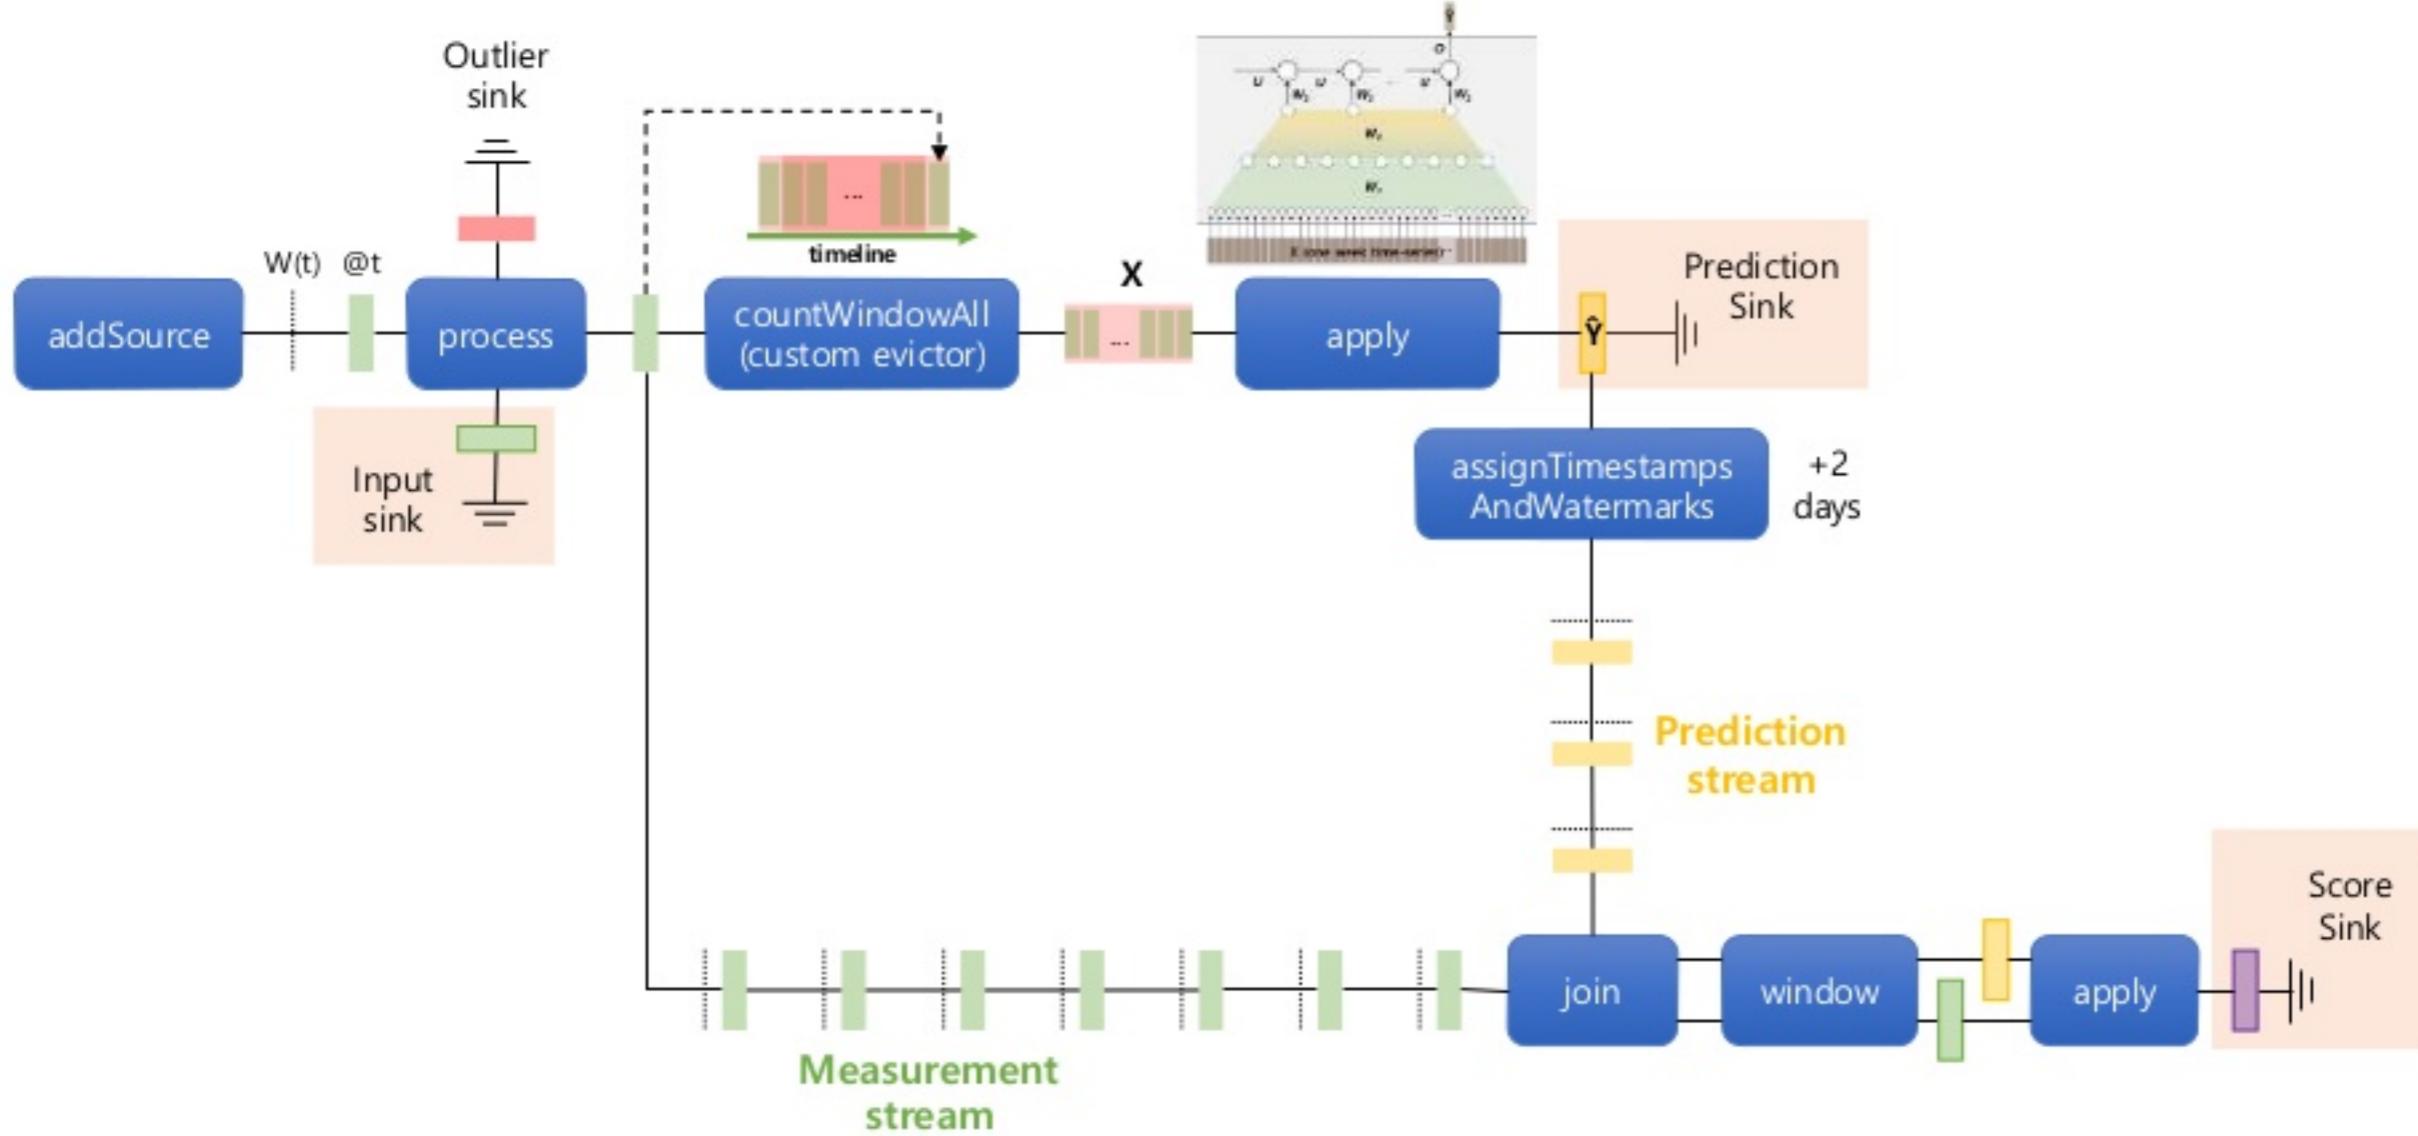


- At a certain time t ,
 - $\textcolor{violet}{Y}$ of timestamp t is arriving
 - \hat{Y} of timestamp $t+2d$ is arriving
 - \hat{Y} of timestamp t has arrived **two days ago**
- `TumblingEventTimeWindows.of(Time.seconds(timeUnit))`
 - To maintain **a window** for **a single pair of $\textcolor{violet}{Y}$ and \hat{Y}**
- A window is triggered when watermarks from both streams have arrived

Data processing pipeline with Flink DataStream API

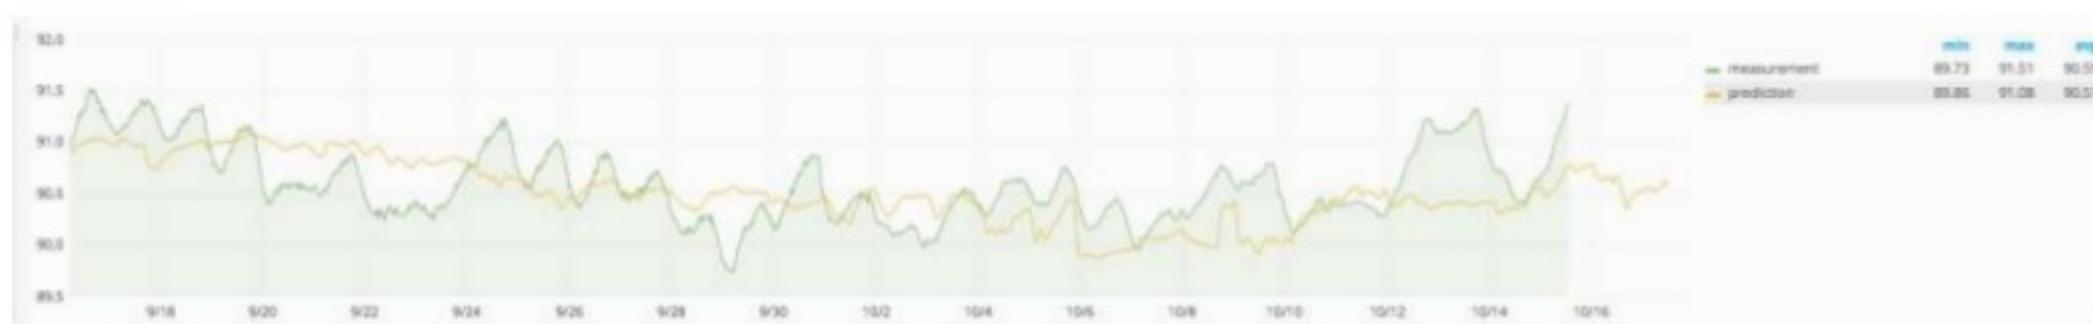
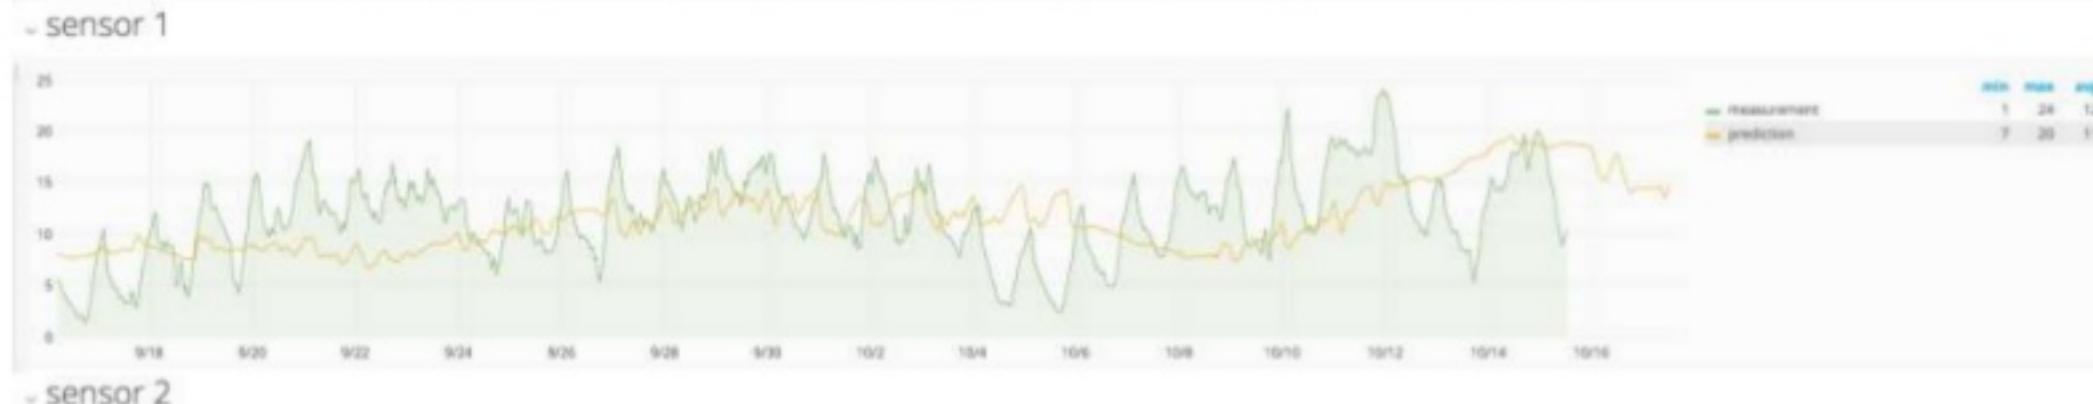
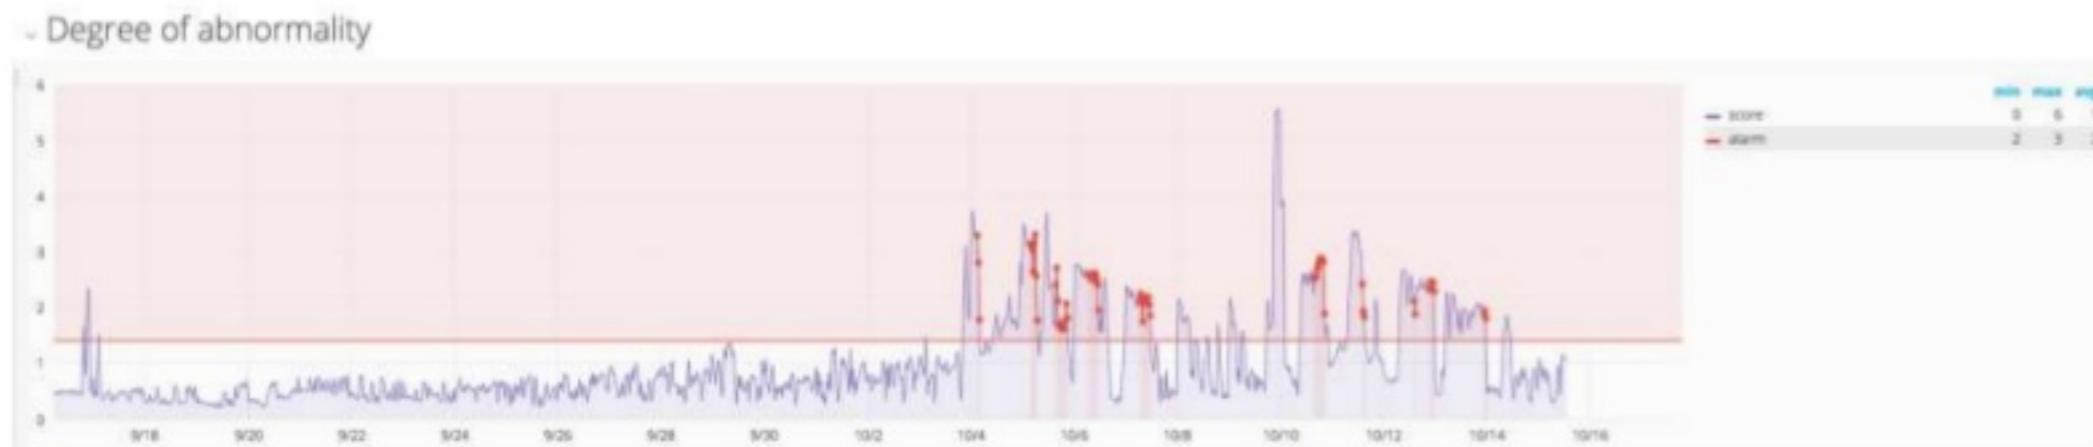


Data processing pipeline with Flink DataStream API



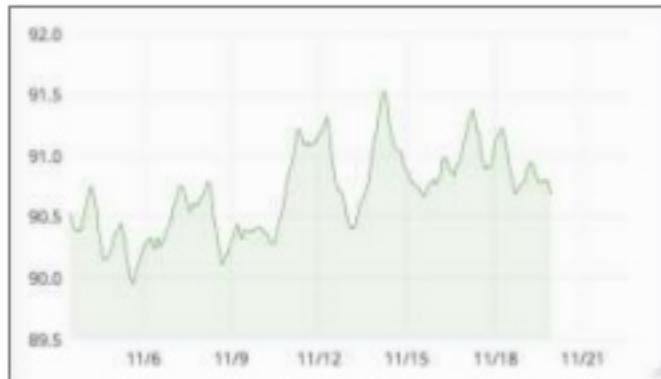
Input, Prediction, Score sinks write records to **InfluxDB**

We then plot time-series using **Grafana**

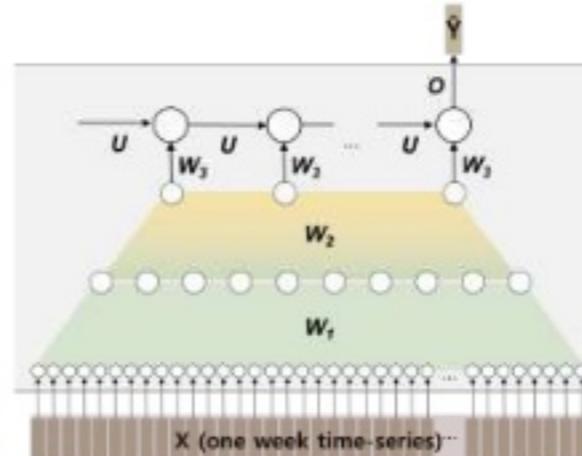


Predicting from a single DNN is not enough!

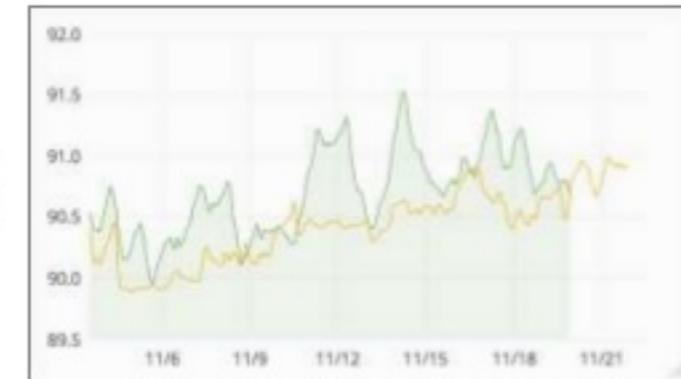
Measurement



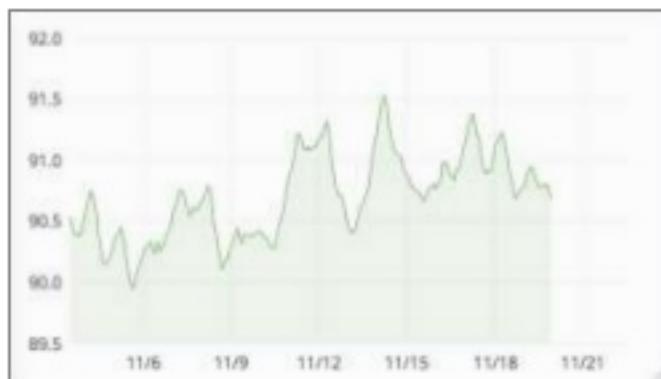
Prediction from a single DNN



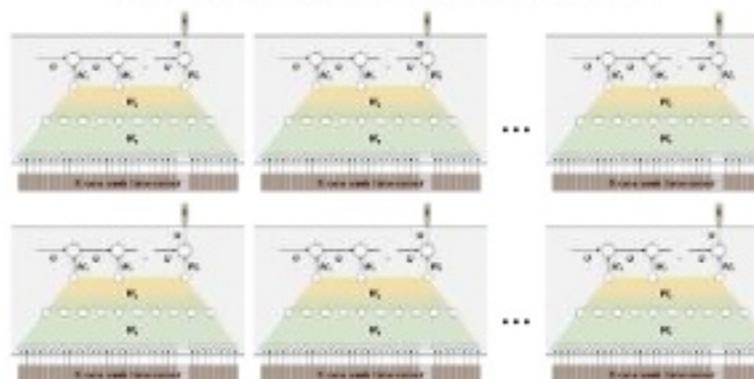
Possibly biased prediction



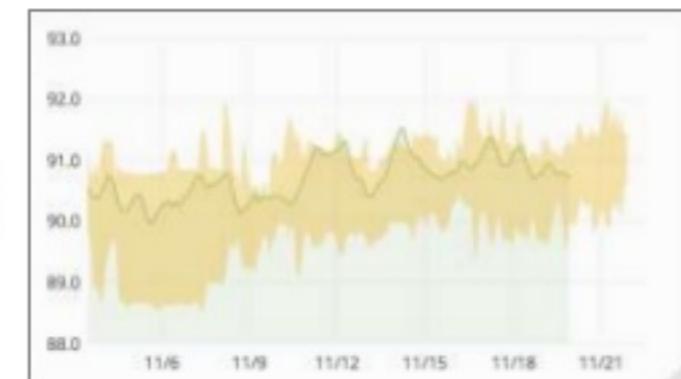
Measurement



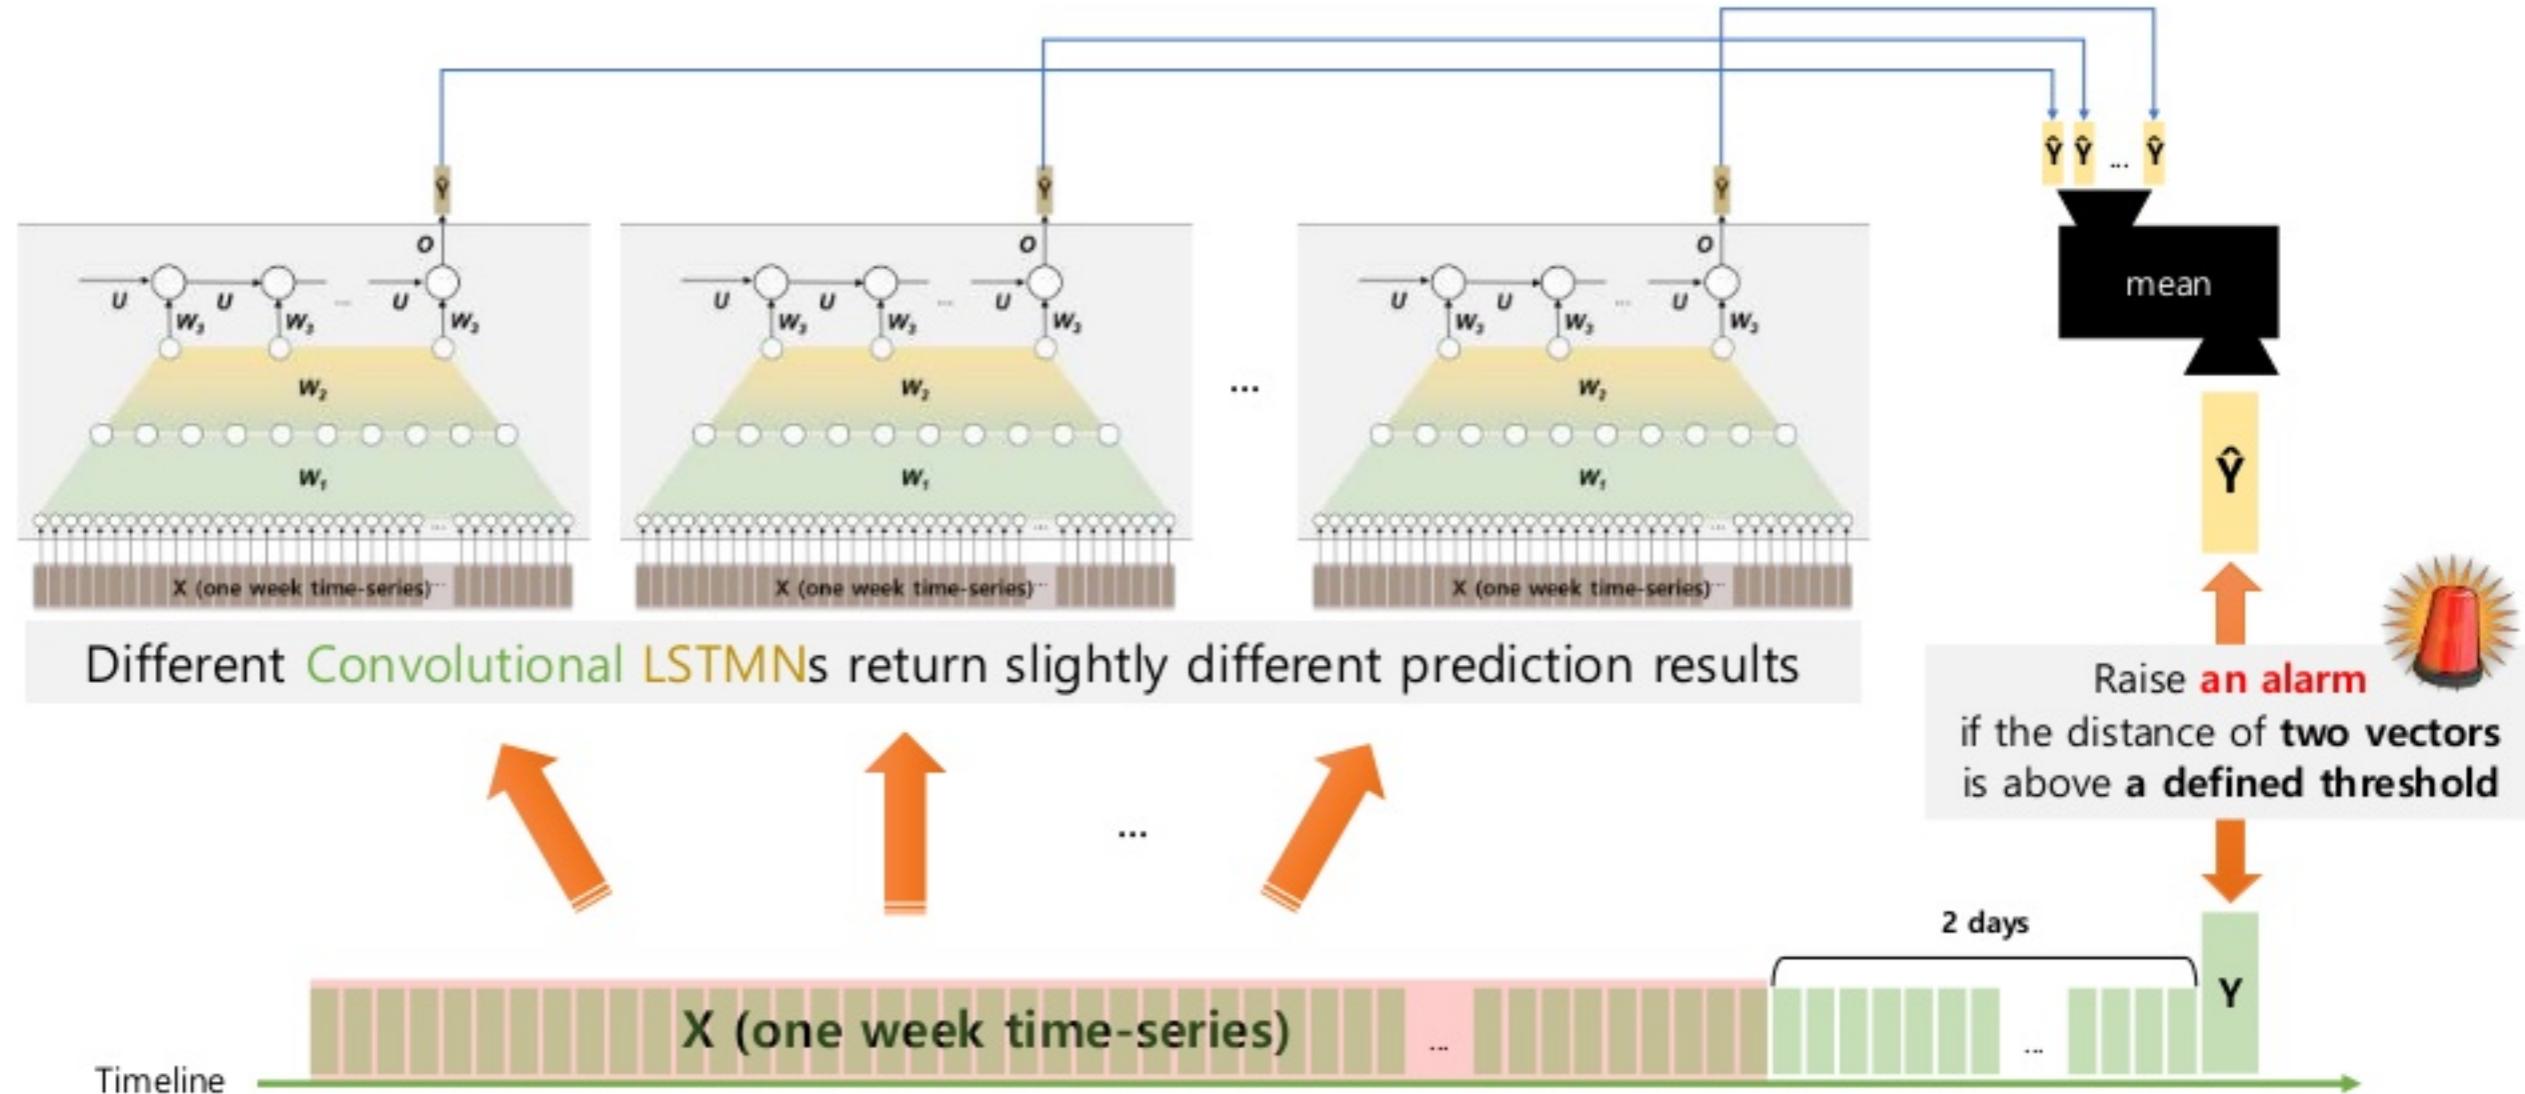
Prediction from
an ensemble of 10 DNNs



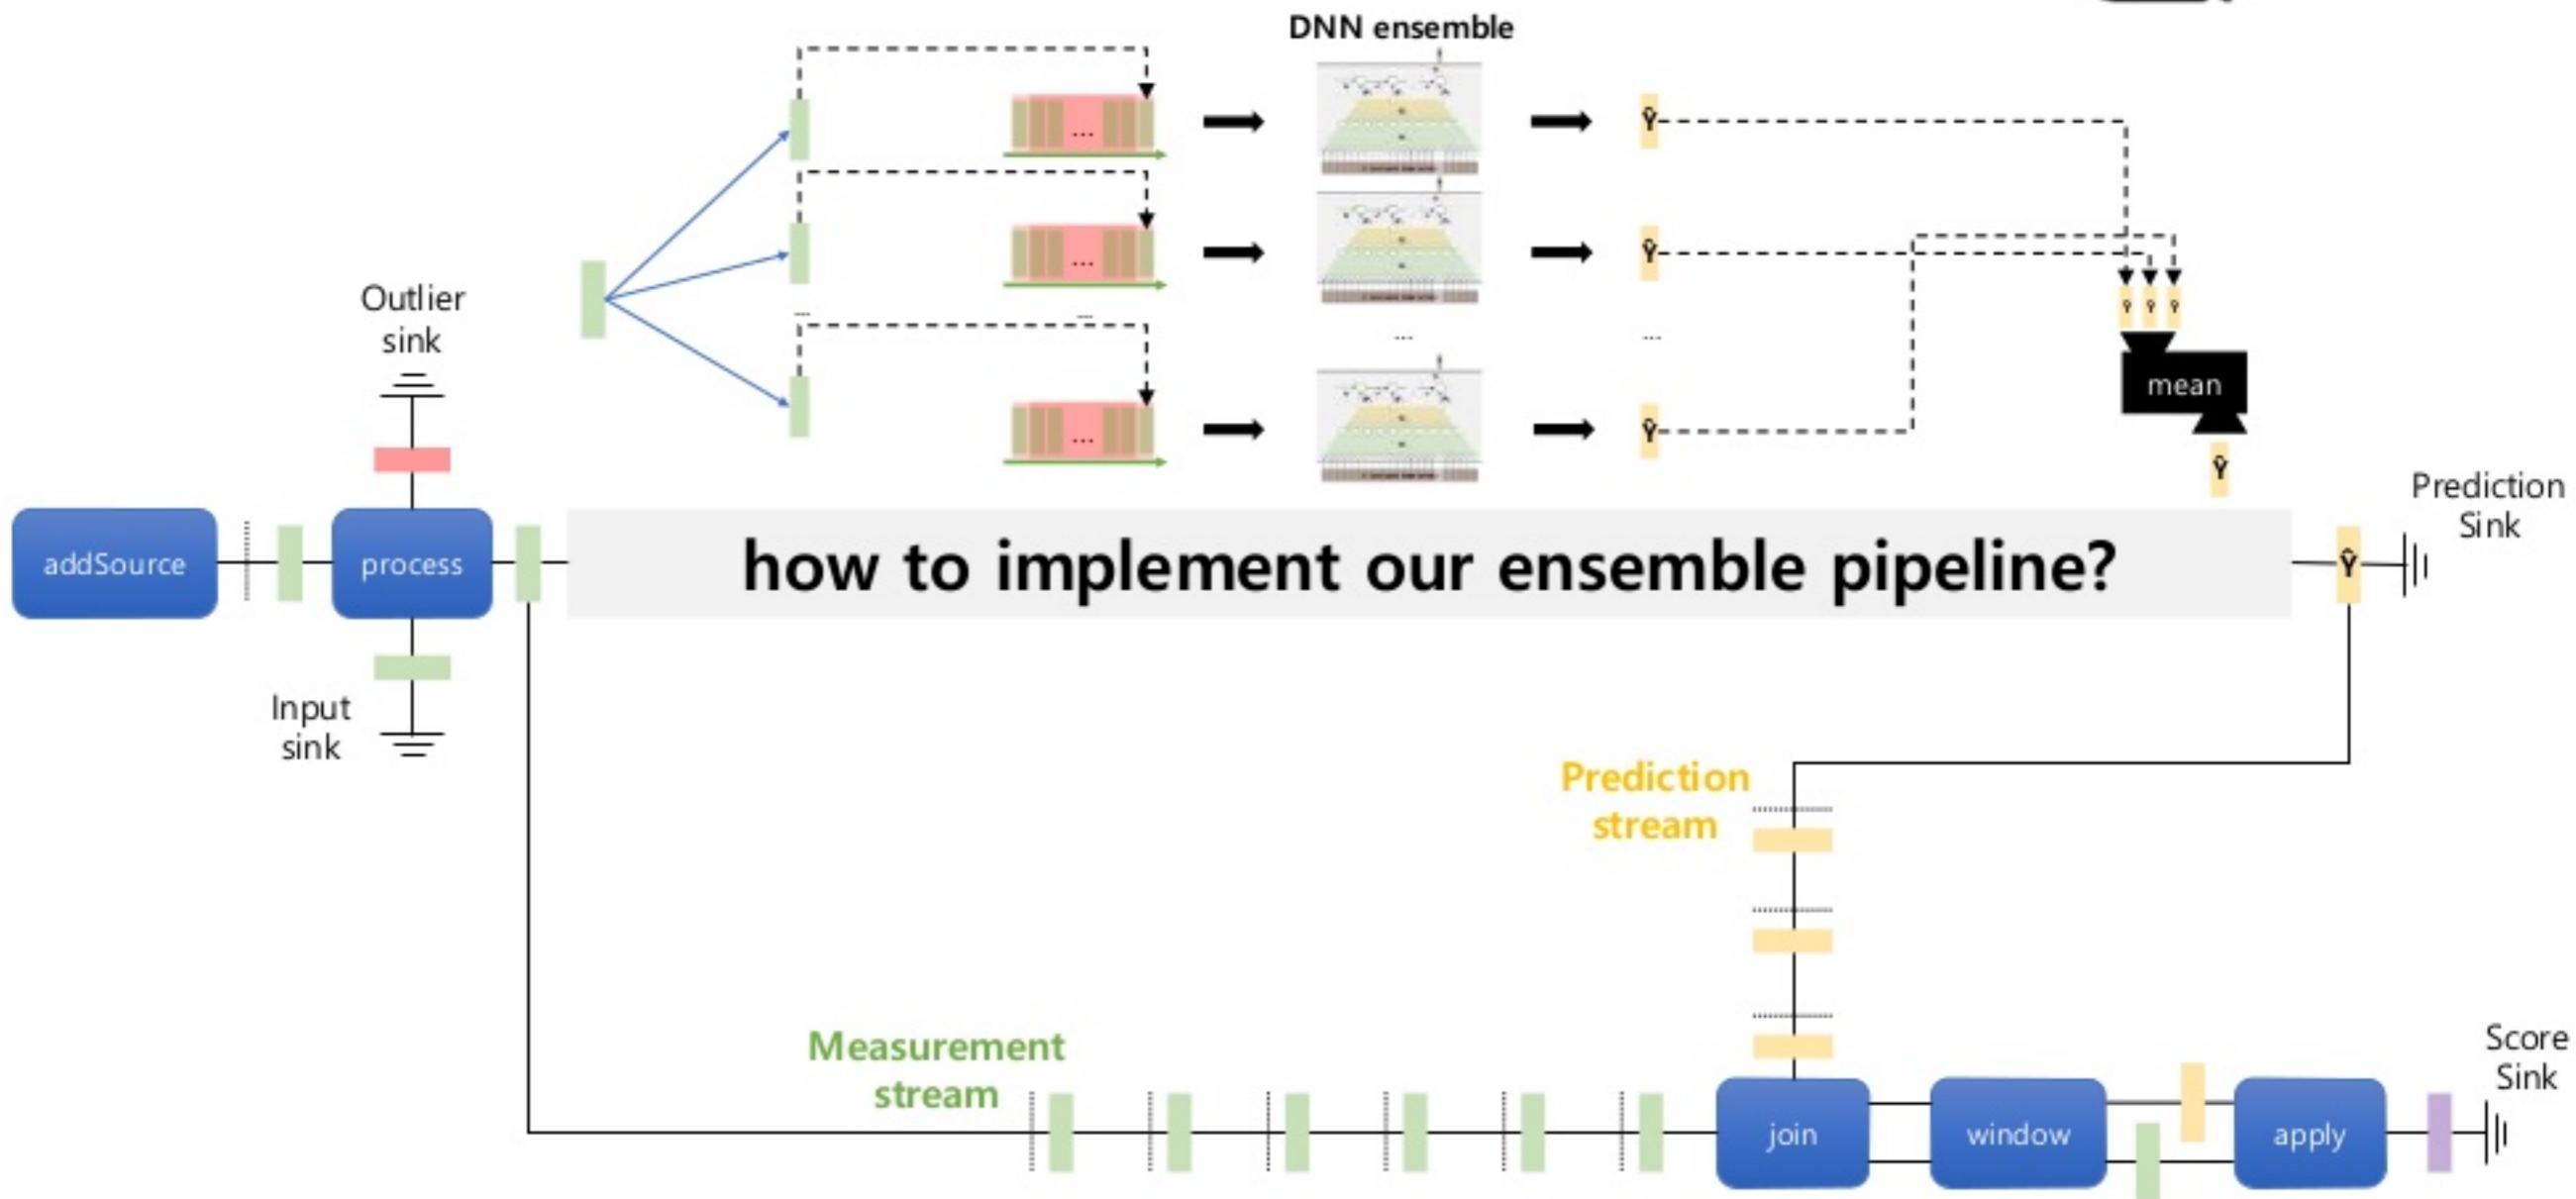
More reliable prediction!



DNN ensemble for **reliable** prediction

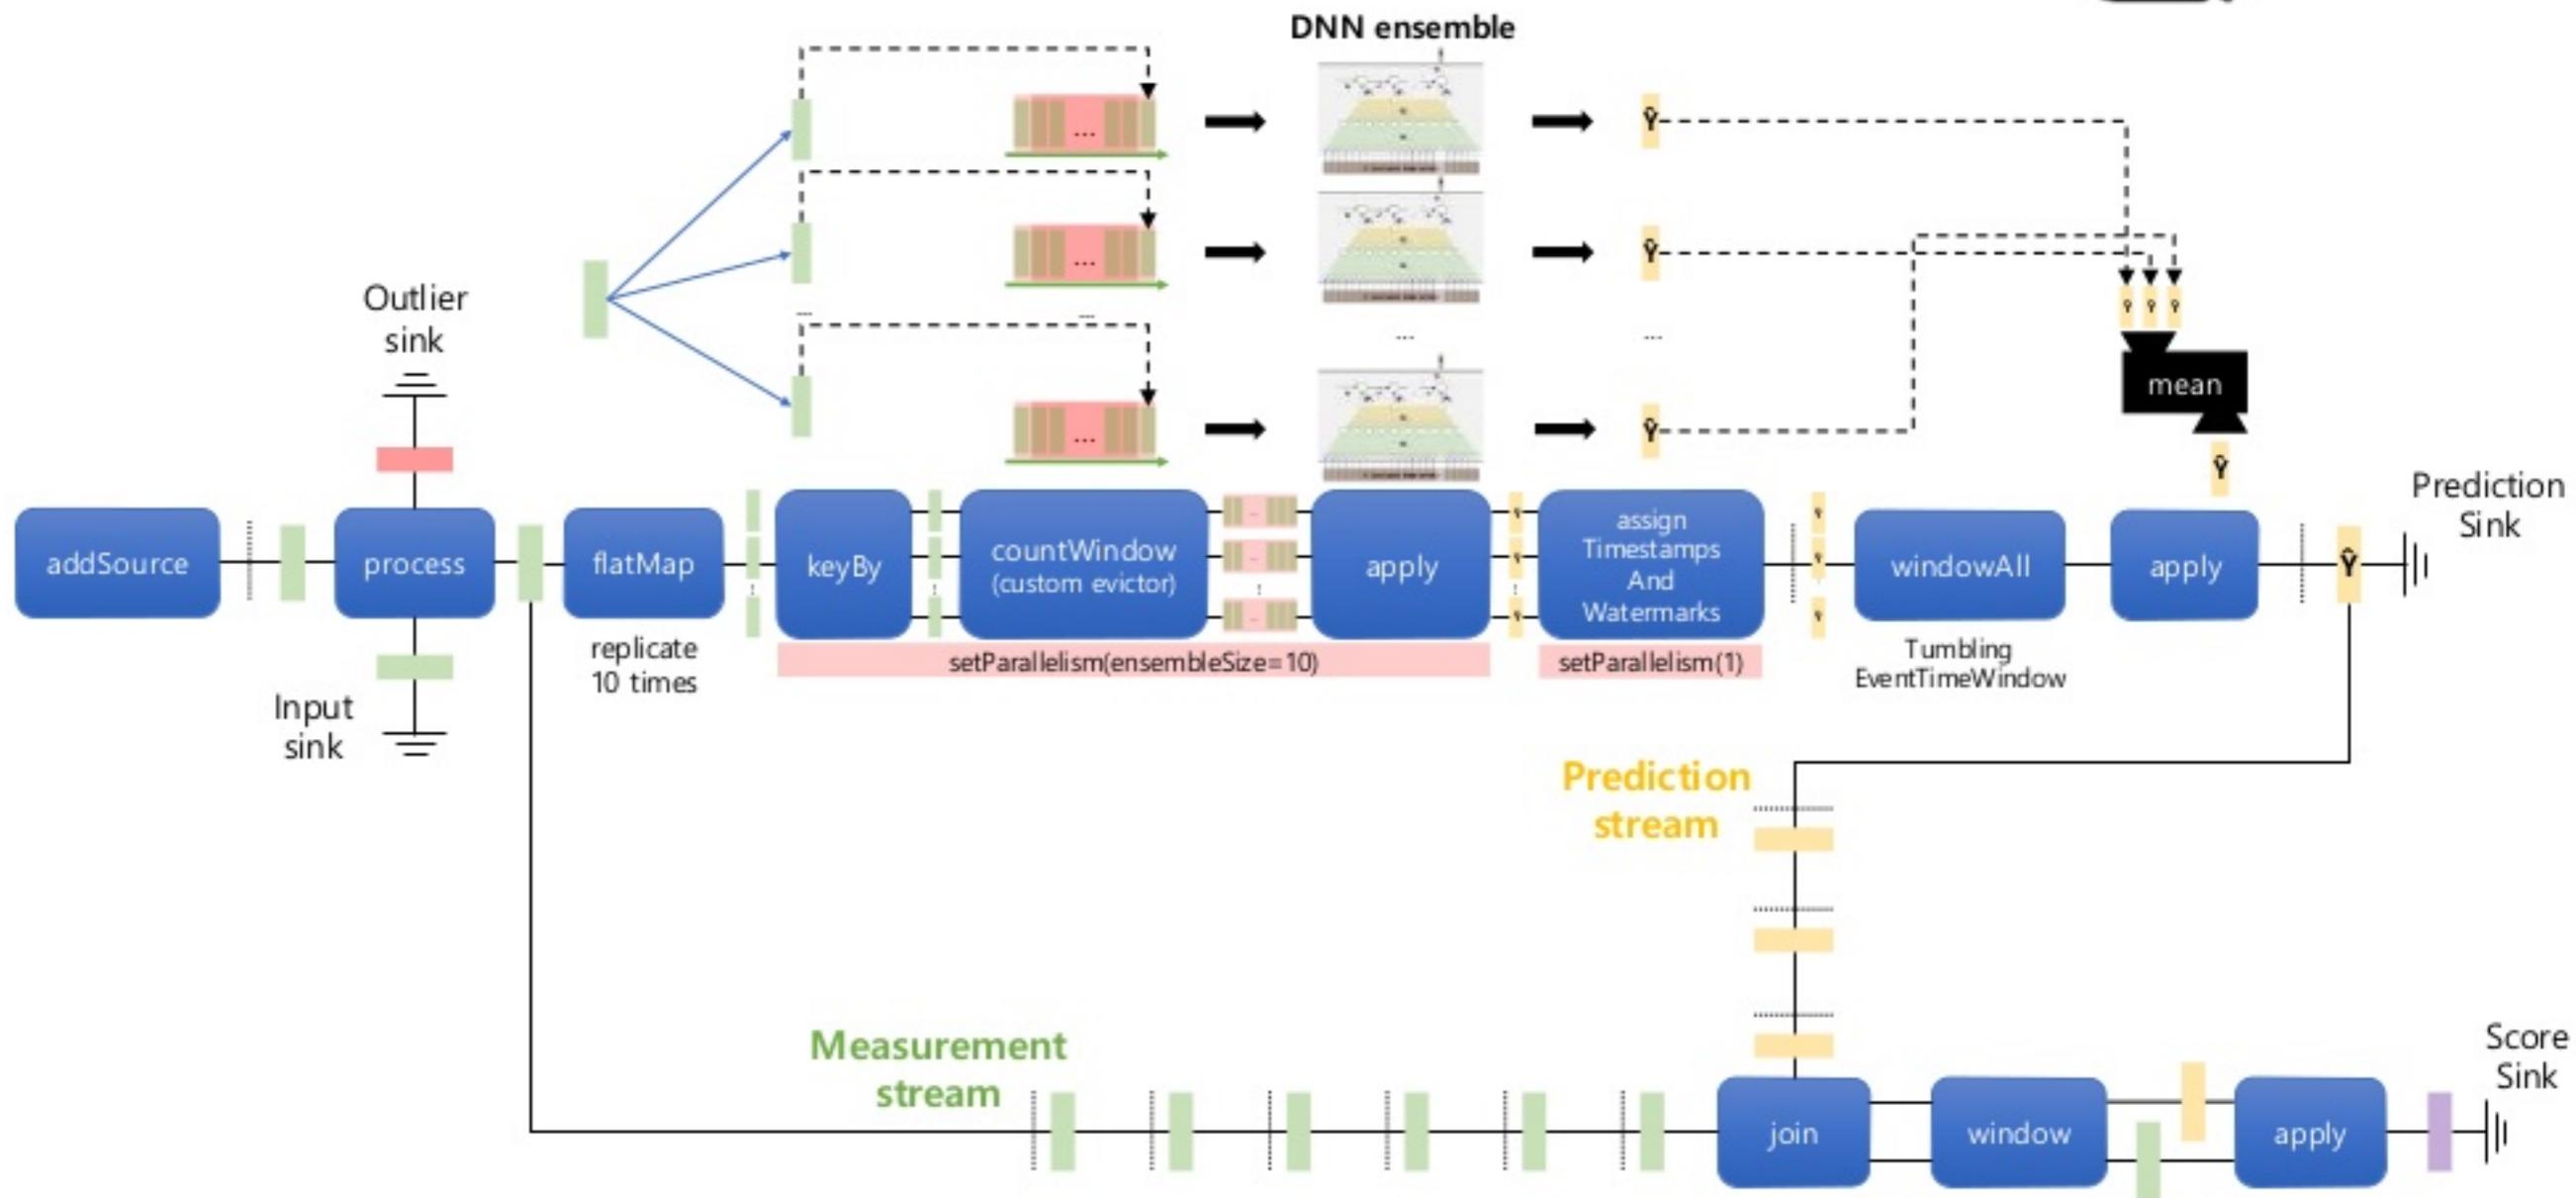


Data processing pipeline with Flink DataStream API

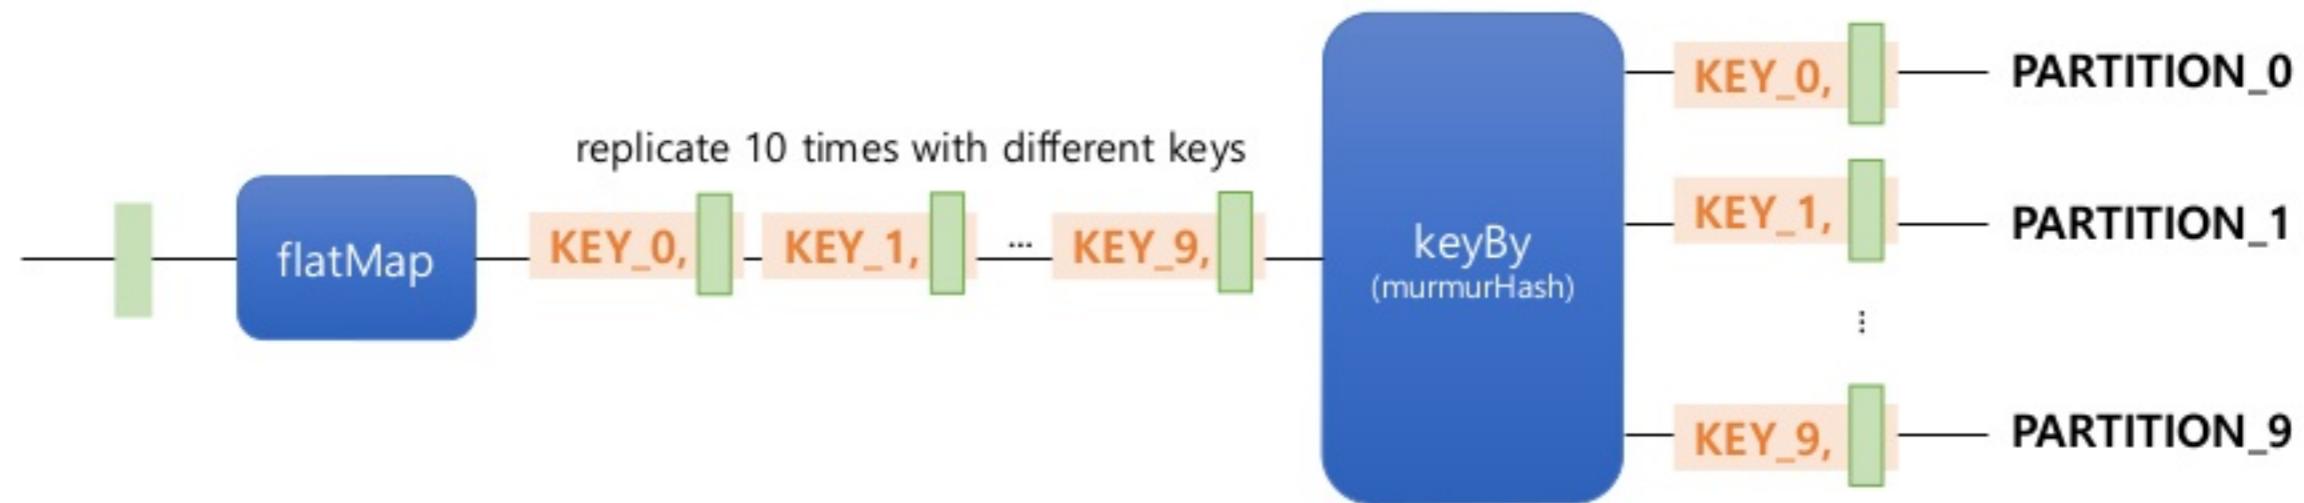




Data processing pipeline with Flink DataStream API



Distribute 10 keys evenly over 10 partitions



Carefully generate keys not to belong to the same partitions

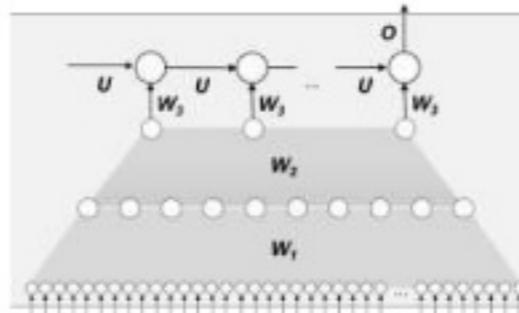
PARTITION = murmurHash(**KEY**) / maxParallalism*(parallelism/maxParallalism)

Contents

1

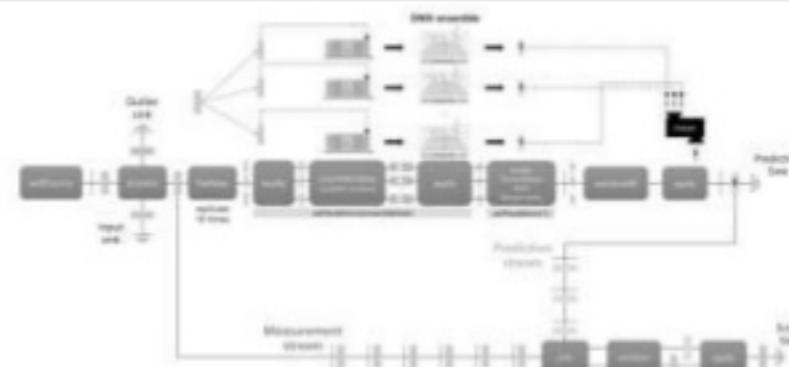
Why we use Flink
for our time-series prediction model

Time-series prediction model



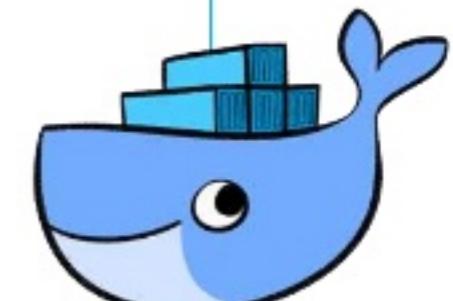
2

Flink pipeline design for
rendezvous and **DNN ensemble**

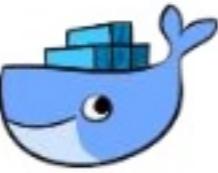


3

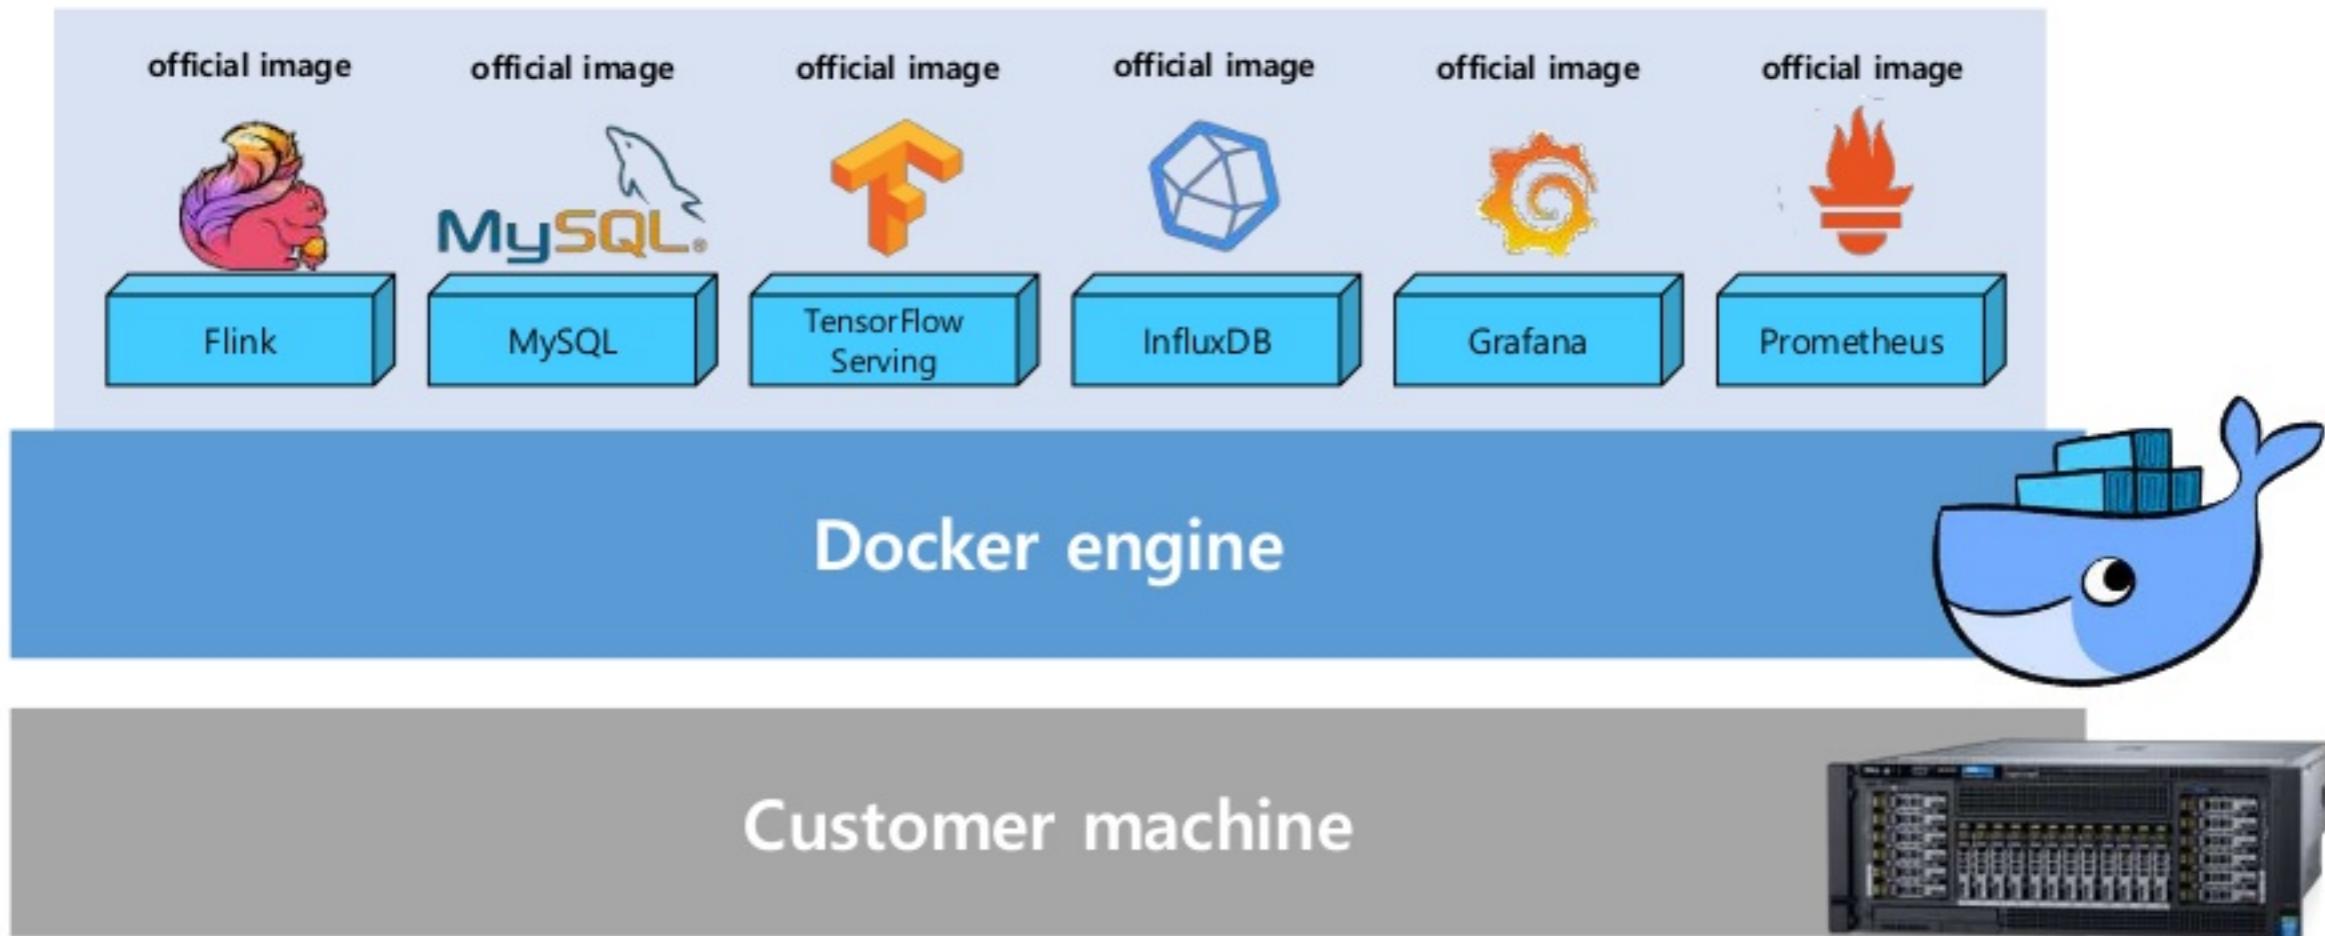
Solution packaging and monitoring
with **Docker** and **Prometheus**



A **simple** software stack on top of Docker



*No custom Docker image!
A single yml file is okay to deploy our software stack!*



Launch JobManager & TaskManager with some changes

in the official repository of the Docker image for Flink

Running a cluster using Docker Compose

With Docker Compose you can create a Flink cluster:

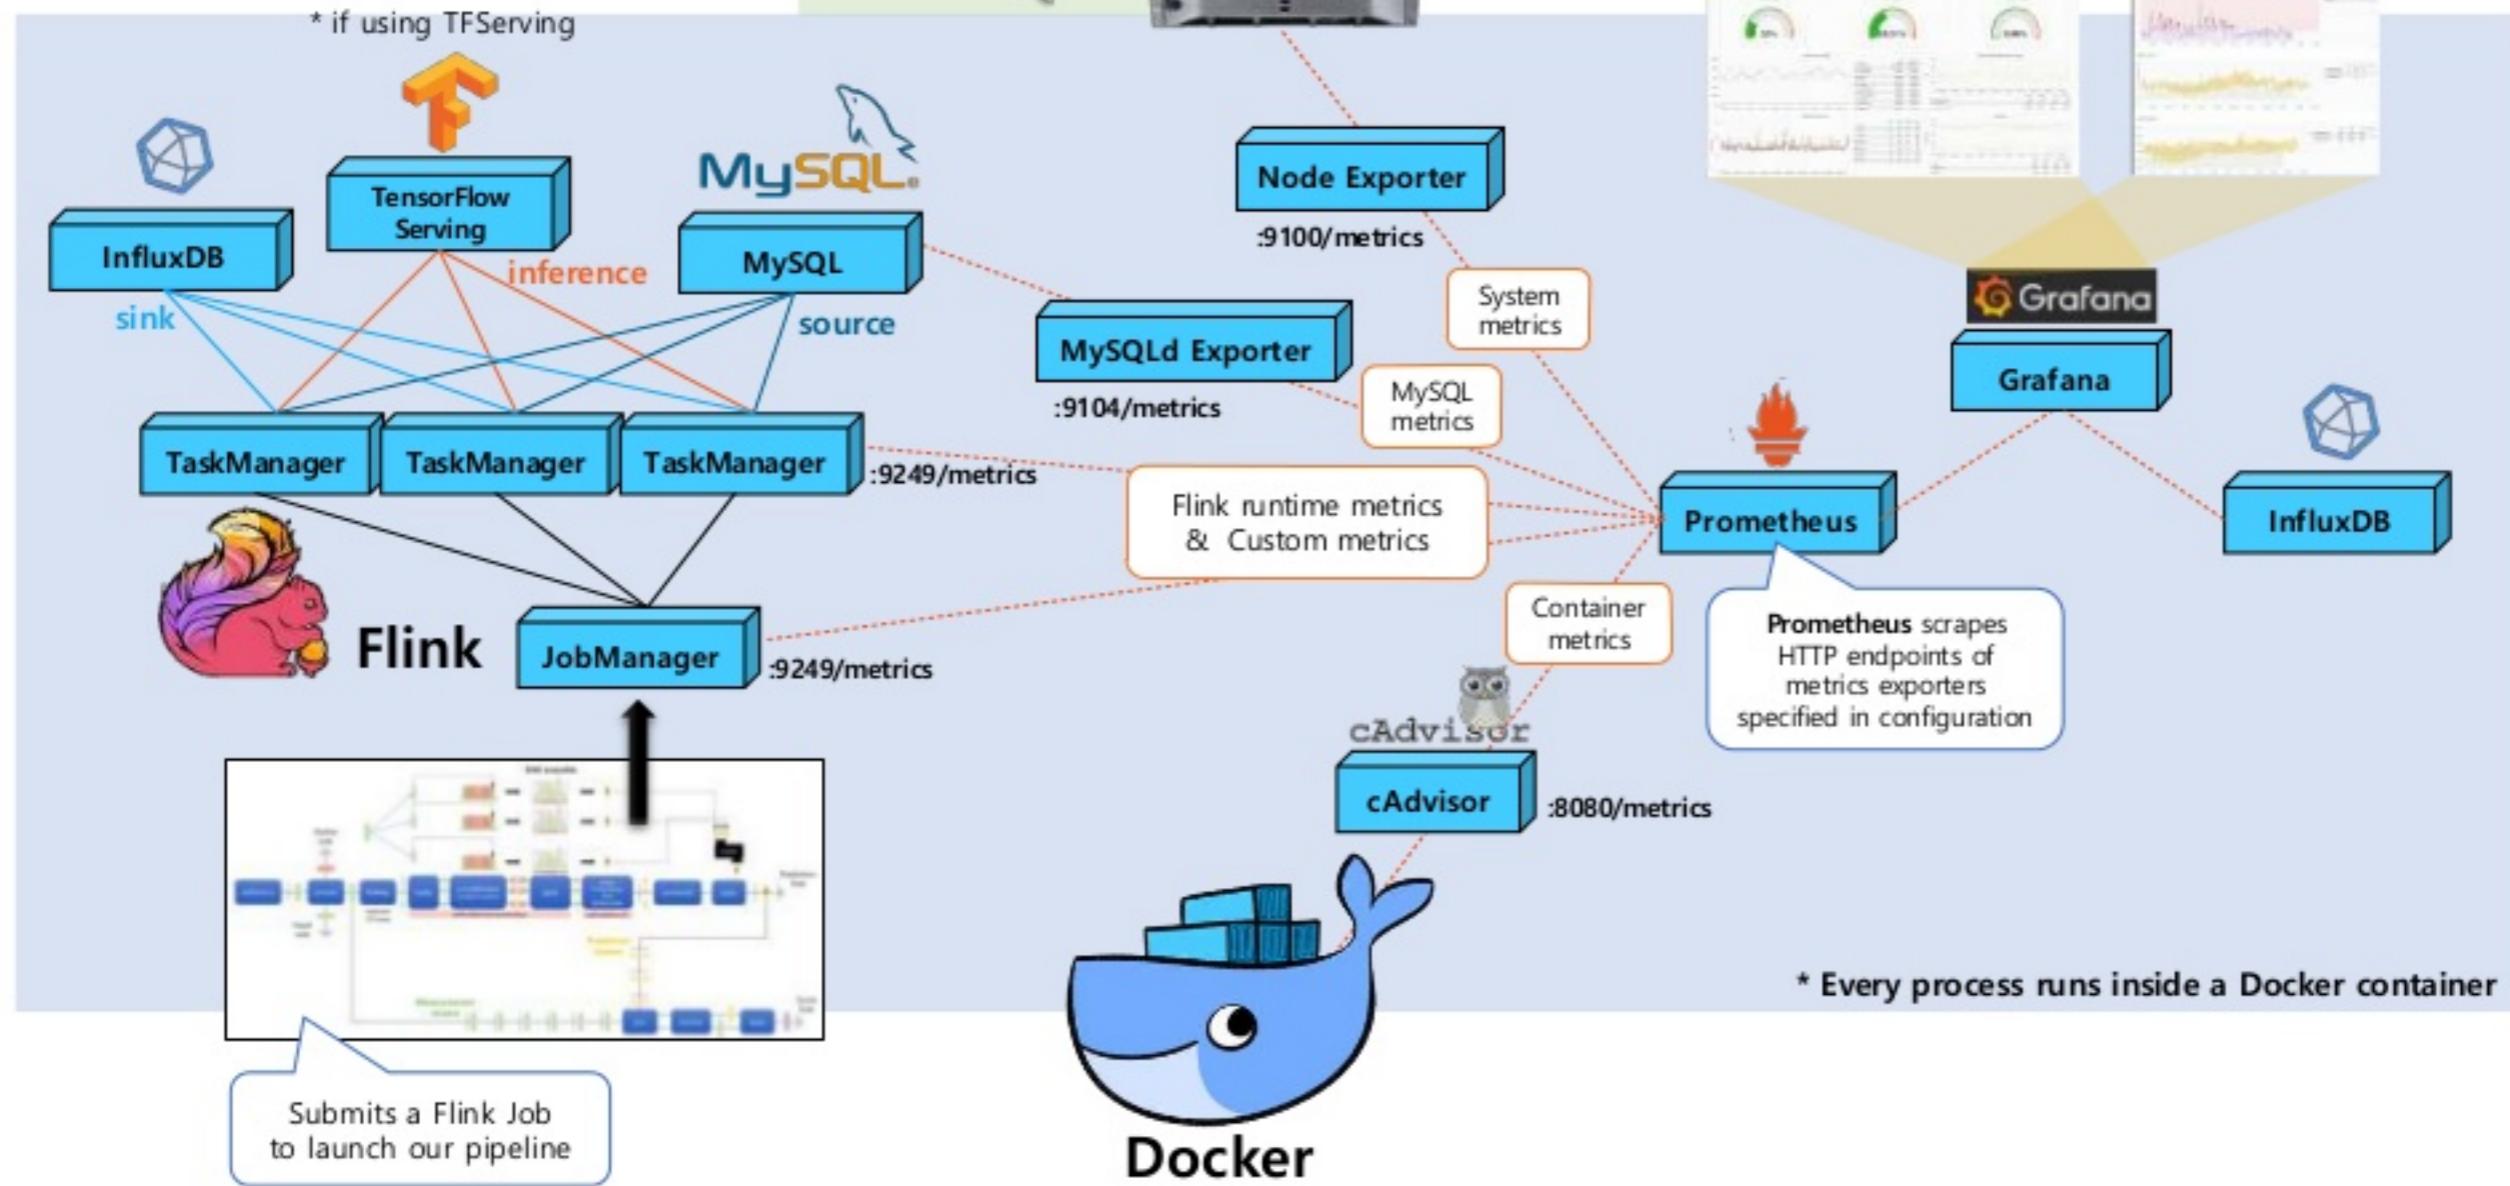
```
version: "2.1"
services:
  jobmanager:
    image: ${FLINK_DOCKER_IMAGE_NAME:-flink}
    expose:
      - "6123"
    ports:
      - "8081:8081"
    command: jobmanager
    environment:
      - JOB_MANAGER_RPC_ADDRESS=jobmanager
  taskmanager:
    image: ${FLINK_DOCKER_IMAGE_NAME:-flink}
    expose:
      - "6121"
      - "6122"
    depends_on:
      - jobmanager
    command: taskmanager
    links:
      - "jobmanager:jobmanager"
    environment:
      - JOB_MANAGER_RPC_ADDRESS=jobmanager
```

```
metrics.reporter      : prom
metrics.reporter.prom.class :
  org.apache.flink.metrics.prometheus.PrometheusReporter
jobmanager.heap.mb   : 10240
taskmanager.heap.mb  : 10240
```

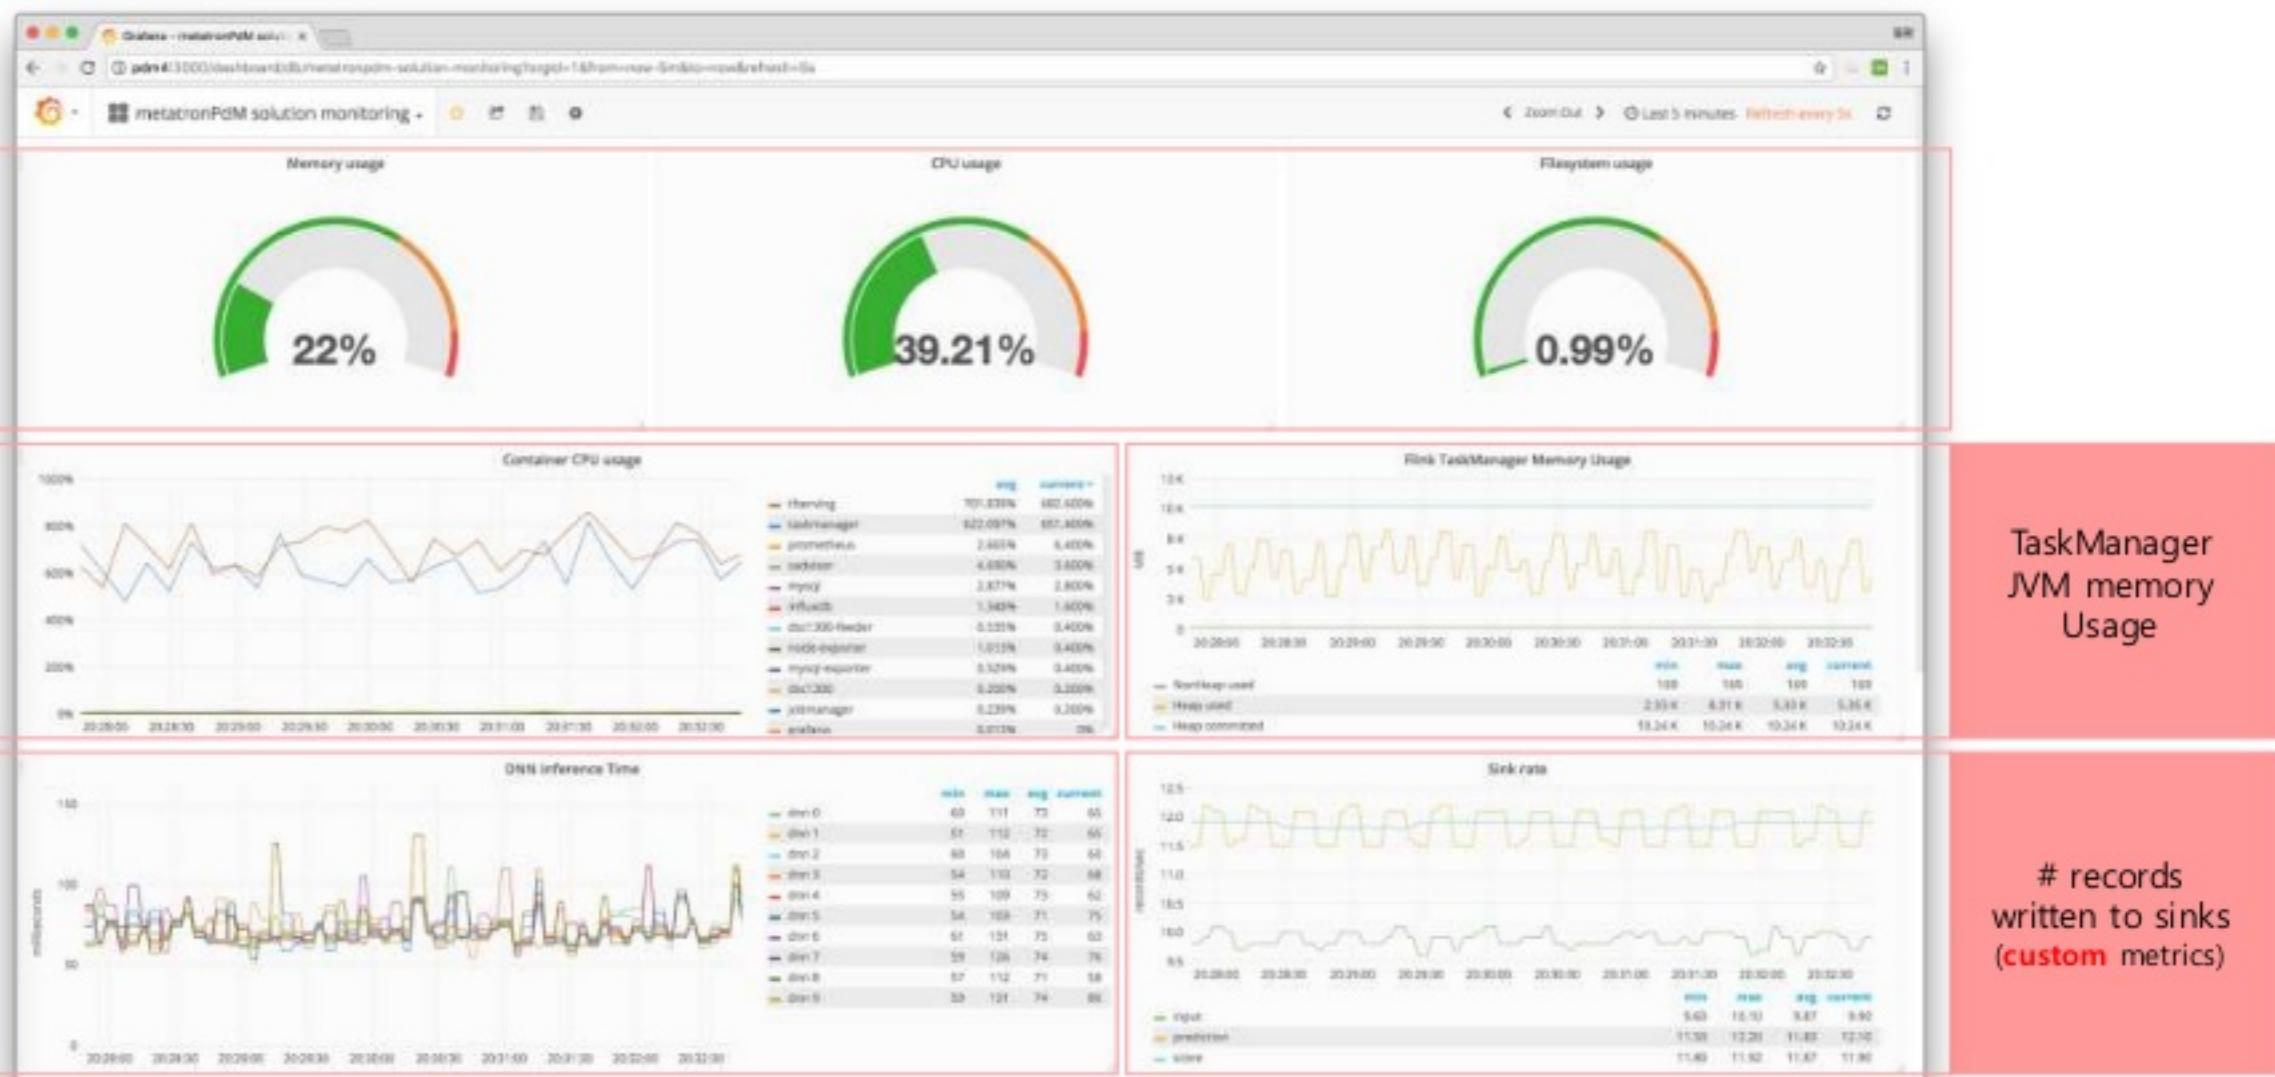
```
volumes:
  - ${PWD}/flink-conf.yaml:/opt/flink/conf/flink-conf.yaml
  - ${PWD}/flink-metrics-prometheus-1.4-SNAPSHOT.jar:/opt/flink/lib/flink-metrics-prometheus-1.4-SNAPSHOT.jar
```

You need to get
flink-metrics-prometheus-1.4-SNAPSHOT.jar
by yourself
until Flink-1.4 is officially released

Solution deployment



Solution monitoring dashboard



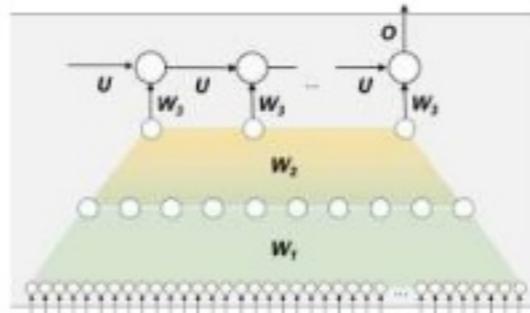
* this dashboard is based on "Docker Dashboard" by Brian Christner

Recap – contents

1

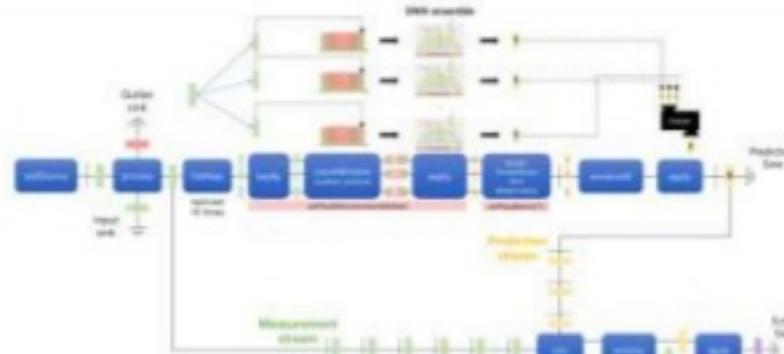
Why we use Flink
for our time-series prediction model

Time-series prediction model



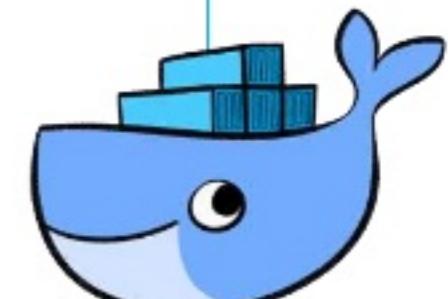
2

Flink pipeline design for
rendezvous and DNN ensemble



3

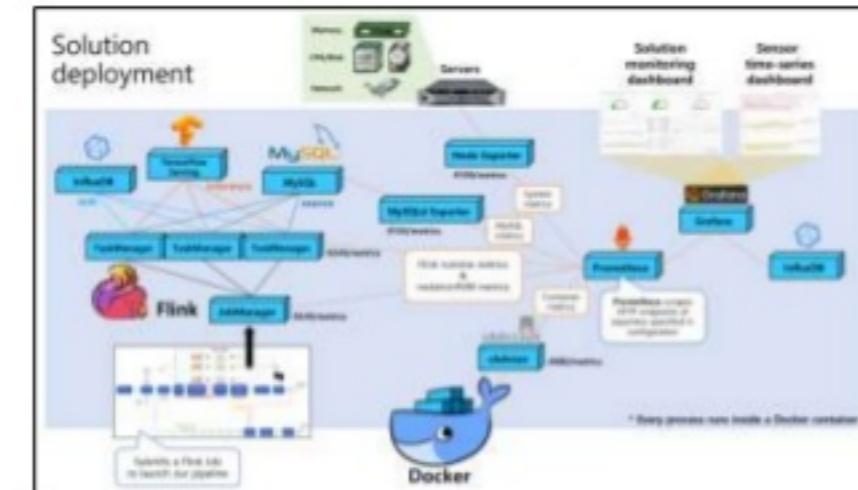
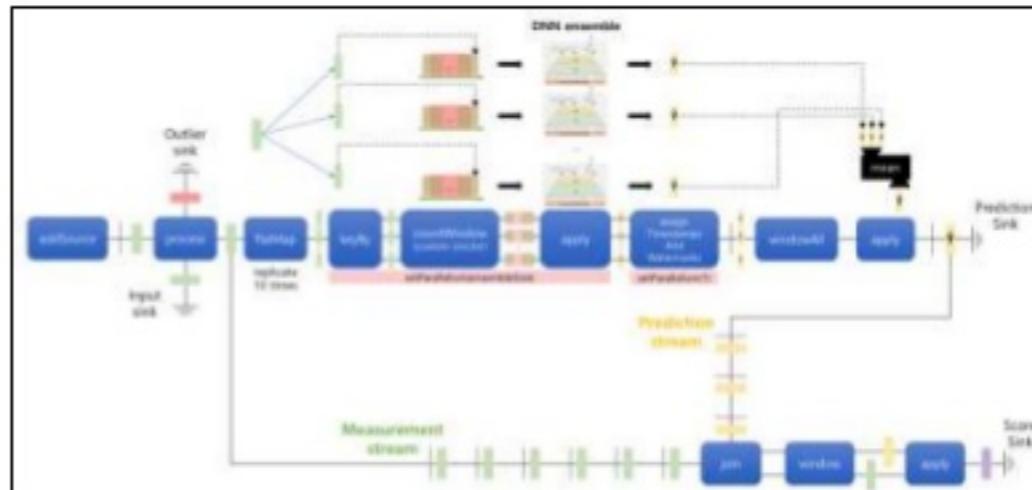
Solution packaging and monitoring
with Docker and Prometheus



Conclusion



- Flink helps us concentrate on the core logic
 - DataStream API is just like *a natural language* in presenting streaming topologies
 - flexible windowing mechanisms (count window and evictor)
 - joining of two streams on event time
 - Thanks to it, we can focus on
 - implementation of custom source/sink to meet customer requirements
 - interaction with DNN ensembles
 - It has a nice ecosystem to help build a solution
 - Docker
 - Prometheus metric reporter



THE END

You cannot use **TFServing** with Flink-1.3.2

- Netty binary incompatibility
 - **flink-runtime_2.11:1.3.2**
 - depends on **io.netty:4.0.27.Final**
 - **grpc-netty:4.1.14**
 - depends on **io.netty:4.1.14.Final**
- From “New and noteworthy in 4.1” page of Netty project
4.1 contains multiple additions which might not be fully backward-compatible with 4.0
- You could use **grpc-okhttp** instead of **grpc-netty**
 - **grpc-okhttp** conflicts with another library (influxdb client for Java)
- You can use **TFServing** with **FLINK-1.4-SNAPSHOT**
 - [FLINK-7013] Add shaded netty dependency
 - **io.netty.*** is recently relocated to **org.apache.flink.shaded.netty4.***

Looking forward to the official release of v1.4

Flink / FLINK-6221
Add Prometheus support to metrics

Agile Board Export

Details		People	
Type:	Improvement	Status:	CLOSED
Priority:	Minor	Resolution:	Fixed
Affects Version/s:	1.2.0	Fix Version/s:	1.4.0
Components:	Metrics	Assignee:	Maximilian Bode
Labels:	None	Reporter:	Joshua Griffith
Description			
Prometheus is becoming popular for metrics and alerting. It's possible to use statsd-exporter to load Flink metrics into Prometheus but it would be far easier if Flink supported Prometheus as a metrics reporter. A dropwizard client exists that could be integrated into the existing metrics system.			
		Dates	
		Created:	30/Mar/17 16:15
		Updated:	02/Jul/17 12:38
		Resolved:	01/Jul/17 10:13

Flink / FLINK-6529 Rework the shading model in Flink / FLINK-7013
Add shaded netty dependency

Agile Board Export

Details		People	
Type:	Sub-task	Status:	CLOSED
Priority:	Major	Resolution:	Fixed
Affects Version/s:	None	Fix Version/s:	None
Components:	Network	Assignee:	Chesney Schepler
Labels:	None	Reporter:	Chesney Schepler
Issue Links			
links to	GitHub Pull Request #4452		
		Dates	
		Created:	27/Jun/17 11:18
		Updated:	07/Aug/17 13:18
		Resolved:	28/Jun/17 09:00

No more 1.4-SNAPSHOT on the customer site!