

The background of the slide features a large, semi-transparent watermark of the Flink Forward Berlin 2017 logo. The logo consists of the word "FLINK" in white, bold, sans-serif capital letters at the top, and "FORWARD" in a larger, dark blue, bold, sans-serif font below it. A stylized orange and yellow globe graphic is positioned behind the text.

FLINK

Moving Beyond Moving Bytes

Joey Frazee
Suneel Marthi

September 12, 2017
Flink Forward, Berlin, Germany

\$WhoAreWe

Joey Frazee

 @jfrazee

- Product Solutions Architect, Hortonworks
- Committer on Apache NiFi, and PMC on Apache Streams

Suneel Marthi

 @suneelmarthi

- Principal Software Engineer, Office of Technology, Red Hat
- Member of Apache Software Foundation
- Committer and PMC on Apache Mahout, Apache OpenNLP, Apache Streams

Agenda

- What is a Schema Registry?
- Why should you Care?
- What Exists Today?
- Different Wire Formats
- Using a Schema Registry
- Using a Schema Registry across a Data pipeline
- Implementation with Flink Deserialization Schemas
- Demo

What is a Schema Registry?

What is a Schema?

- information about a record
- field names, field types, default values and aliases

A schema registry is:

- a centralized, versioned schema repository service
- that supports de-centralized schema-based serialization and deserialization

Why should you care?

Because, Real-time stream processing mandates that you know the semantics of your data:

- Interesting operations on streaming
 - joins, projection, aggregation, filtering, streaming SQL all require prior knowledge of the types and values of data
- Otherwise, you're just moving bytes and counting anonymous things (you don't need something as powerful as Flink to do that)

And, using embedded schemas is an (unnecessary) overhead:

- The schema can be larger than data
- And it introduces a copy of the schema for every message or topic

And, including schemas in your project is bad:

- Is not recommended b/c it tightly couples the project to your data
- Creates practical scalability issues across system and application boundaries

In general a schema registry offers or implements :

- Schema database
- Version control strategy
- Application API for serialization and deserialization according to the schema
- API service (e.g., REST) for schema management
- Way to acquire, include, or pull in binary artifacts (e.g., SerDes) from the service
- Wire format that encodes a schema *identifier* along with contents in serialized objects

What Exists Today

3 Options

Cask Schema Registry

- A schema serving layer for Avro and Protobuf to support data preparation and validation in Cask CDAP Wrangler

Confluent Schema Registry

- An interface for storing and retrieving Avro schemas for efficient serialization in Kafka and interop with Kafka Streams

Hortonworks Registry

- Shared repository of schemas and SerDes to support Avro schema sharing, record processing and serialization in and across applications (e.g., Apache NiFi, Hortonworks Streaming Analytics Manager)

Wire Formats

Cask (N/A?)

Confluent (5 byte header)

- Magic byte/protocol version (byte): 0
- Schema id (int): 1-4

Hortonworks (13 byte header)

- Magic byte/protocol version (byte): 0
- Schema id (long): 1-8
- Schema version (int): 9-12

Feature Comparison

	REST API	Schemas	Custom SerDes	Storage	HA	UI	Maven Plugin	Schema Compatibility Checking	Kafka Integration
Cask	Y	Avro, Protobuf		Cask CDAP DataSet	?		Y		?
Confluent	Y	Avro		Kafka Topic	master/slave		Y	Y	Y
Hortonworks	Y	Avro	Y	JDBC, HDFS	storage + load balancer/proxy				Y

Using a Schema Registry

Add a New Schema

Add New Schema

NAME * ✓ | CLEAR

DESCRIPTION *

TYPE * x ▾

SCHEMA GROUP *

COMPATIBILITY ▼

EVOLVE

SCHEMA TEXT *

```

1  {
2      "namespace": "twitter",
3      "type": "record",
4      "name": "Tweet",
5      "fields": [
6          {"name": "id", "type": "long"},
7          {"name": "id_str", "type": "string"},
8          {"name": "text", "type": "string"},
9          {"name": "lang", "type": "string"},
10         {"name": "favorite_count", "type": "int"},
11         {"name": "created_at", "type": "string"},
12         {"name": "timestamp_ms", "type": "long"},
13         {"name": "user", "type": {
14             "type": "record",
15             "name": "User",
16             "fields": [
17                 {"name": "id", "type": "long"},
18                 {"name": "id_str", "type": "string"}]}},
19     ]
20 }
```

CANCEL SAVE

Schema Entry

twitter.Tweet FORWARD	TYPE avro	GROUP Kafka	VERSION 1	SERIALIZER & DESERIALIZER 0
DESCRIPTION v1	<pre>1 { 2 "namespace": "twitter", 3 "type": "record", 4 "name": "Tweet", 5 "fields": [6 { 7 "name": "id", 8 "type": "long" 9 }, 10 { 11 "name": "id_str", 12 "type": "string" 13 }, 14 { 15 "name": "text" 16 } 17] 18}</pre>	VERSION 1	CHANGE LOG v1 30s ago CREATED	

Edit Schema

Edit Version

DESCRIPTION *

Added user.statuses_count|

SCHEMA TEXT *

CLEAR

```
37         "fields": [
38             {"name": "text", "type": "string"
39         ]
40     }
41 },
42 {"name": "user_mentions", "type": {
43     "type": "array", "items": {
44         "type": "record",
45         "name": "UserMention",
46         "fields": [
47             {"name": "id", "type": "long"},
48             {"name": "id_str", "type": "string"
49             {"name": "screen_name", "type": "string"
50             {"name": "name", "type": "string"
51         ]
52     }
53 }
54 ]
55 }
56 ]
57 }
58 ]
```

CANCEL SAVE

Schema Version

twitter.Tweet
FORWARD

TYPE	GROUP	VERSION	SERIALIZER & DESERIALIZER
avro	Kafka	2	0

DESCRIPTION
Added user.

```
1 {  
2   "namespace": "twitter",  
3   "type": "record",  
4   "name": "Tweet",  
5   "fields": [  
6     {  
7       "name": "id",  
8       "type": "long"  
9     },  
10    {  
11      "name": "id_str",  
12      "type": "string"  
13    },  
14    {  
15      "name": "text"  
16    }  
17  ]  
18}
```

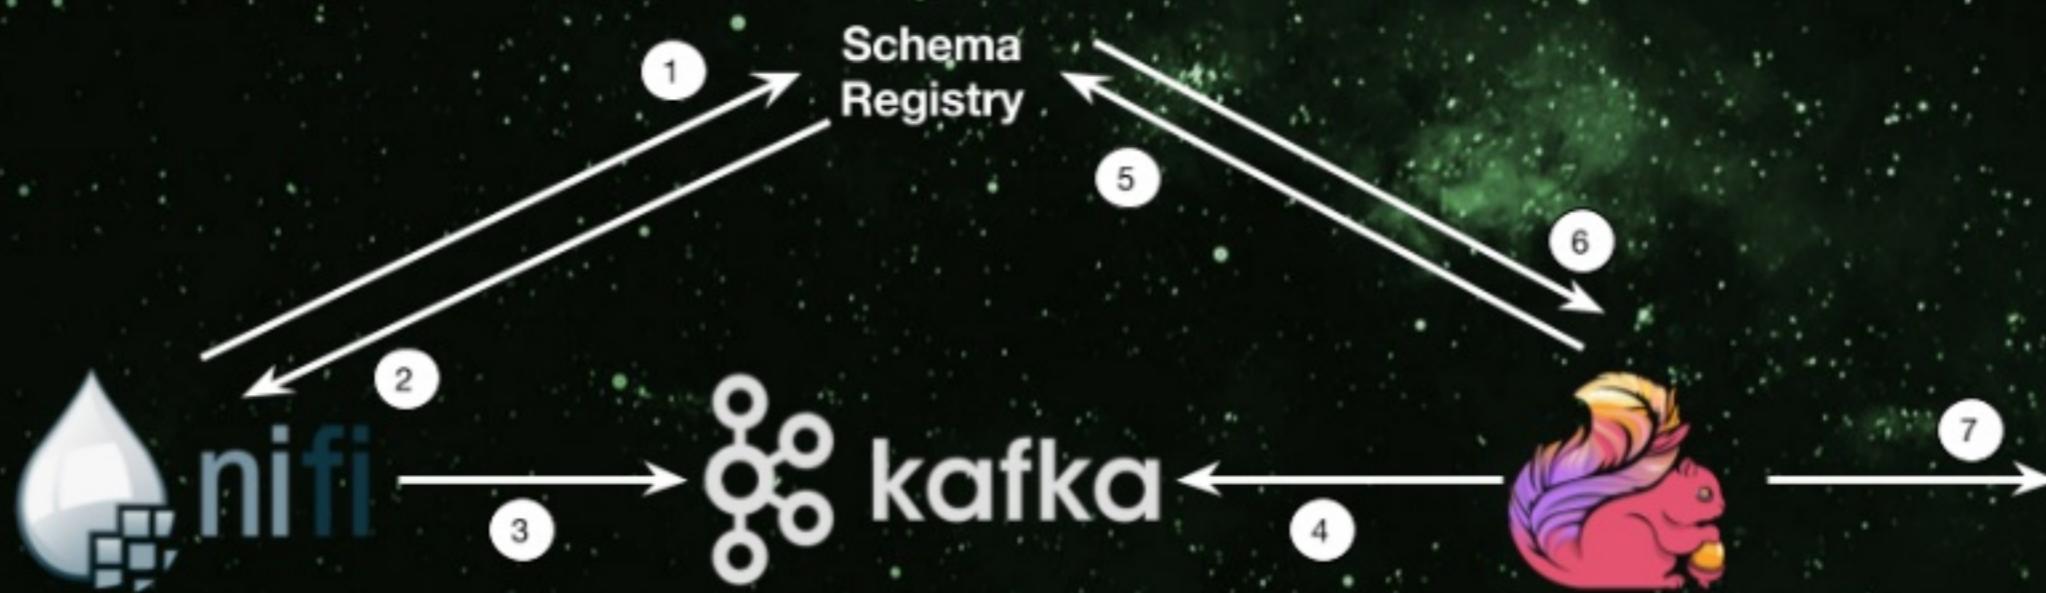
VERSION 2

CHANGE LOG

- v2 7s ago EDITED
- v1 1m 25s ago CREATED

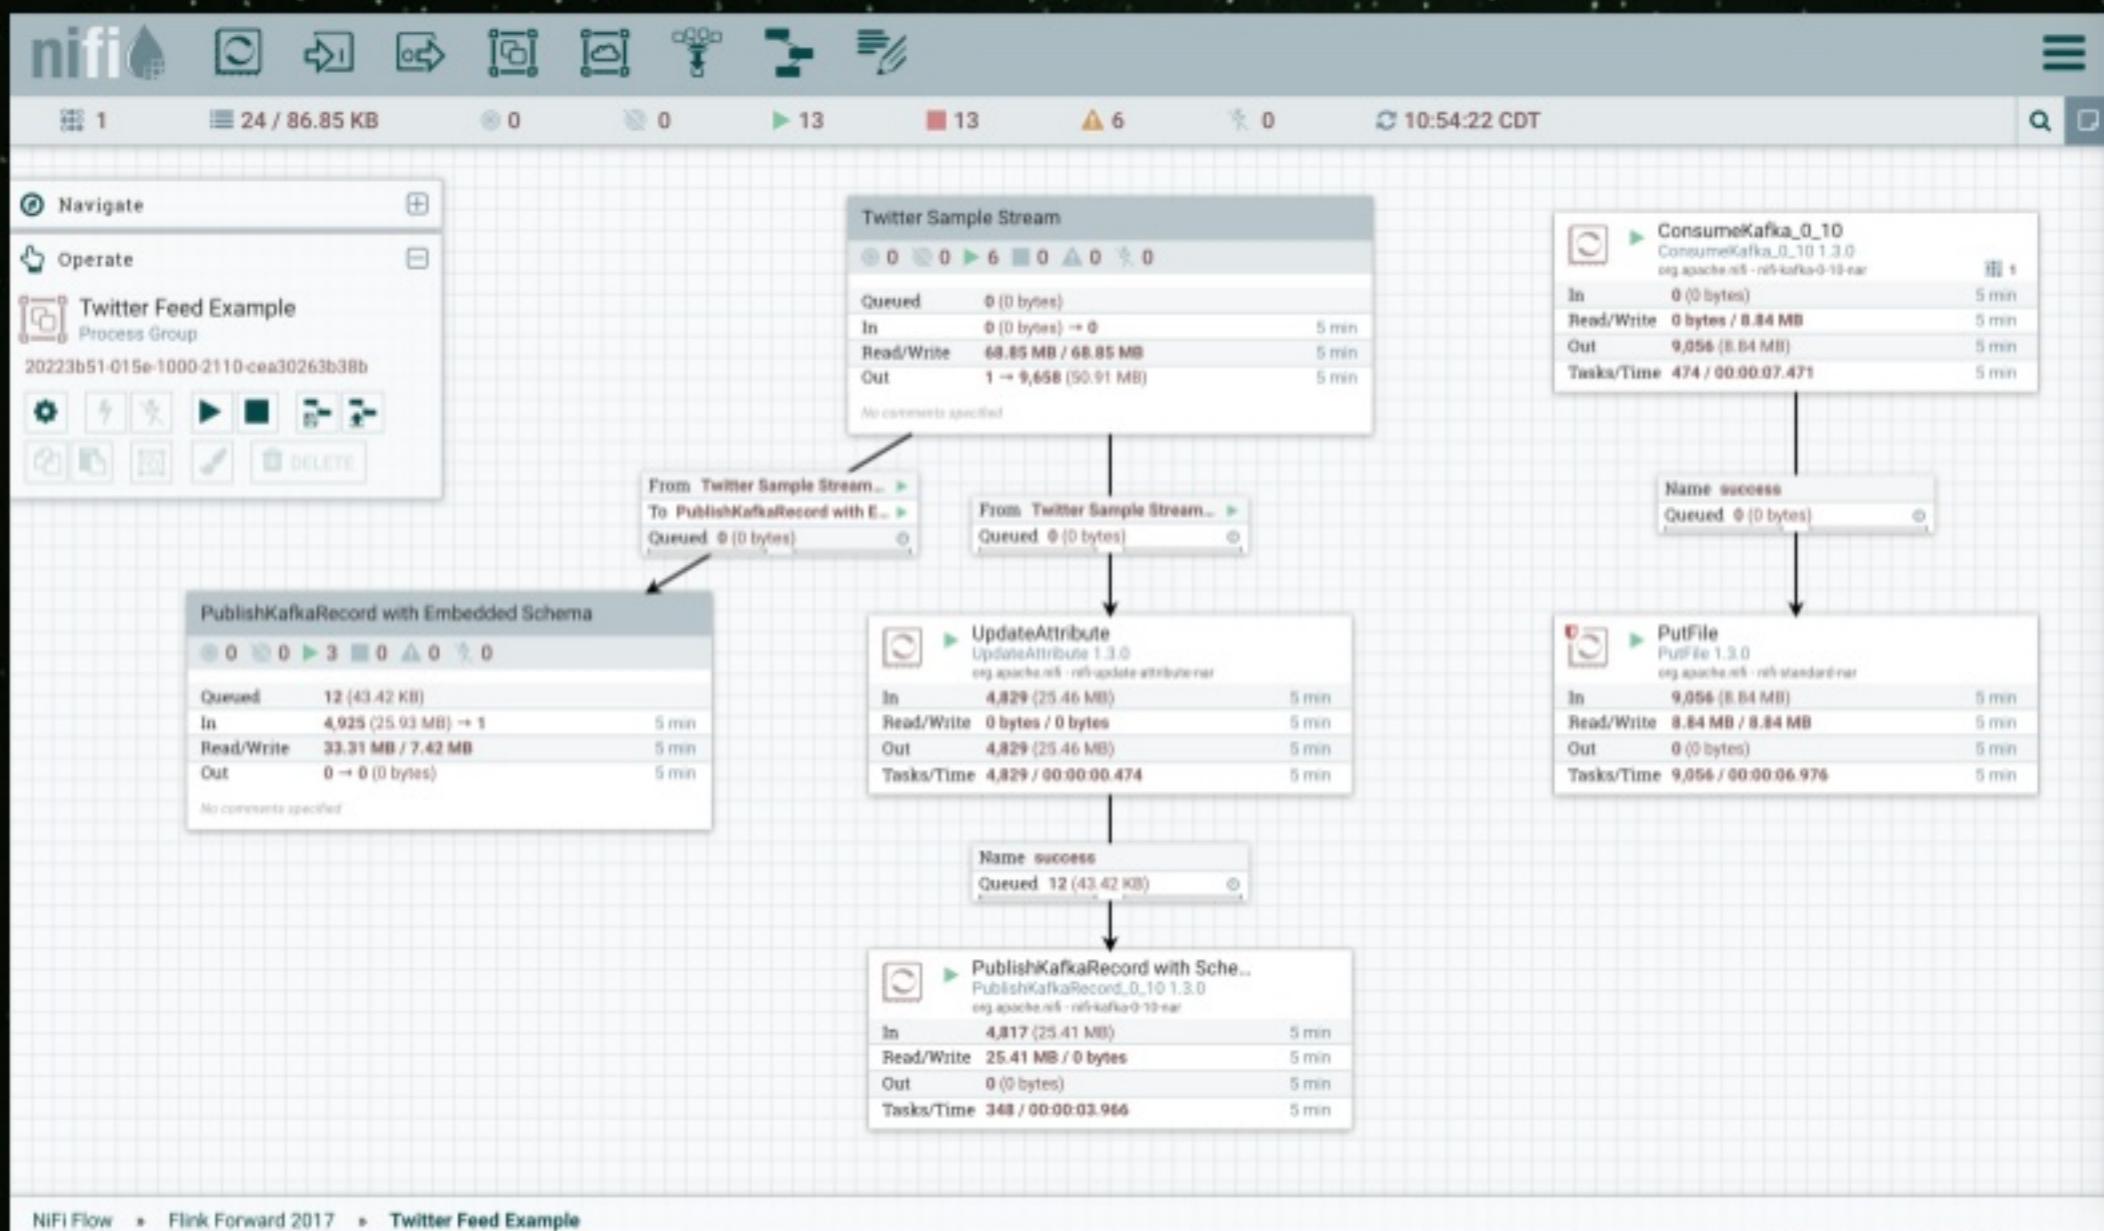
Using a Schema Registry across a Data pipeline

Example Data Pipeline



1. Request schema from schema registry service via schema id
2. Receive the associated schema
3. Serialize the message contents according to the schema, packed with the encoded schema metadata, and publish to Kafka
4. Consume from Kafka and decode the message into its schema metadata and contents
5. Request the schema from schema registry service via schema id
6. Receive the associated schema
7. Deserialize the contents according to the schema and do cool stuff

Apache NiFi Twitter Feed Example



Schema Access Strategies

Embedded schema:

- Whole schema is written out with the message contents (in Avro this corresponds to DataFileReader/Writer)

Schema metadata reference:

- Schema id and other metadata are written as a header with the contents

Implicit schema:

- No schema is presented and application must know what it's expecting or iterate through the universe of possibilities

Serialization with Embedded Schema

0000000	o b j 001 002 026 a v r o . s c h e m a
0000020	a 232 022 { " t y p e " : " r e c o r d "
0000040	r d " , " n a m e " : " T w e e d "
0000060	t " , " n a m e s p a c e " : "
0000100	t w i t t e r " , " f i e l d s
0000120	" : [{ " n a m e " : " i d ",
0000140	" t y p e " : " l o n g " } ,
0000160	" n a m e " : " i d _ s t r " ,
0000200	" t y p e " : " s t r i n g "]
0000220	, { " n a m e " : " t e x t " ,
0000240	" t y p e " : " s t r i n g "]
0000260	, { " n a m e " : " l a n g " ,
0000300	" t y p e " : " s t r i n g "]
0000320	, { " n a m e " : " f a v o r i t e _ c o u n t " ,
0000340	t e _ c o u n t " , " t y p e " :
0000360	: " l o n g " } , { " n a m e "
0000400	...

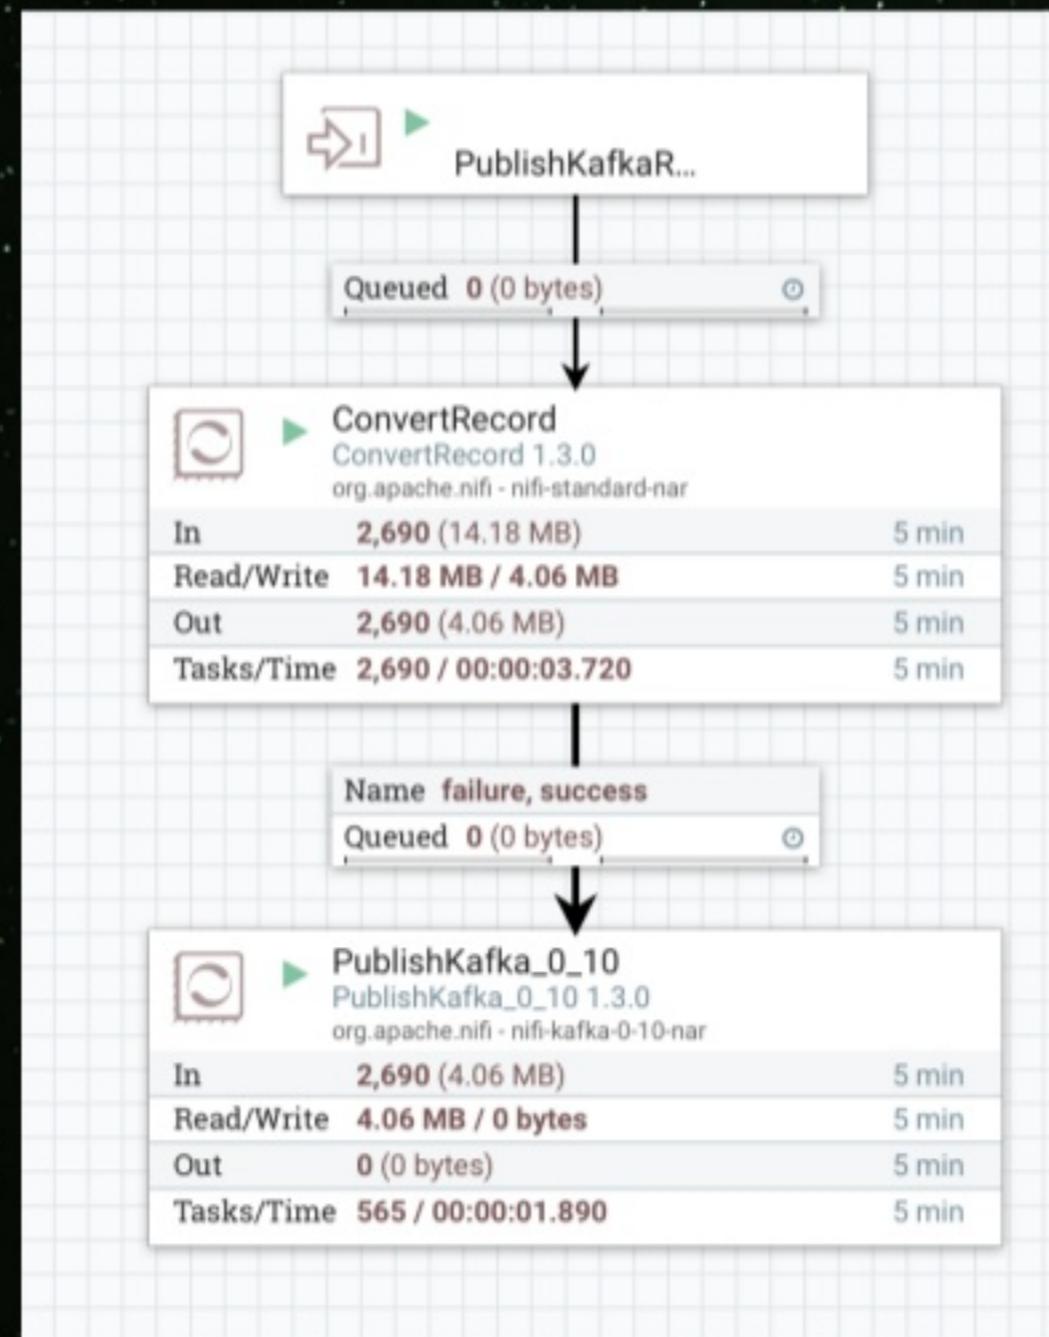
Serialization with Hortonworks Schema Reference

0000000	001	\0	\0	\0	\0	\0	\0	\0	001	\0	\0	\0	001	200	200	3
0000020	?	214	204	?	227	031	\$	9	0	7	3	1	2	6	6	7
0000040	5	8	8	6	8	1	7	2	8	j	R	T	@	B	T	
0000060	S	-	t	w	t	:		T	h	a	n	k	y	o	u	
0000100		?	230	201	?	?	217		h	t	t	p	s	:	/	/
0000120	t	.	c	o	/	8	g	w	a	z	v	b	U	J	C	004
0000140	e	n	\0	<	M	o	n		S	e	p		1	1		1
0000160	8	:	3	9	:	3	1		+	0	0	0	0		2	0
0000200	1	7	032	1	5	0	5	1	5	5	1	7	1	6	6	4
0000220	226	?	225	221	\b	024	1	0	9	1	7	4	6	6	9	9
0000240	006	P	M	Y	022	A	n	a	t	i	A	m	i	r	002	032
0000260	M	a	t	o	k	i		P	l	a	n	e	t	\0	002	
0000300	D	o	n	t		b	l	a	m	e		m	e	,	I	
0000320	m	w	e	i	r	d	?	?	\n	?	\f	?	031	\0	001	
0000340	N	214	?	?	?	002	022	3	3	5	1	4	1	6	3	8
0000360	016	B	T	S		t	w	t	036	?	?	?	?	203	204	?
0000400	206	214	?	205	204	?	213	?	\0							

Serialization with Confluent Schema Reference

0000000	\0	\0	\0	\0	Q	230	?	?	?	201	?	?	227	031	\$	9
0000020	0	7	3	1	8	0	9	5	0	1	7	9	9	6	2	8
0000040	8	v	@	-	-	k	i	l	e	y	@	o	n	l	y	t
0000060	s	i	n	w	o	r	l	d	Y	e	a	h	w	t		
0000100	f	?	?	?	T	h	a	t	i	s	s	u	p			
0000120	e	r		w	e	i	r	d	?	237	230	?	004	e	r	
0000140	\0	<	M	o	n	S	e	p	1	1	1	1	1	9	:	
0000160	0	1	:	0	5	+	0	0	0	0	0	2	0	1	7	
0000200	032	1	5	0	5	1	5	6	4	6	5	6	6	4	?	?
0000220	?	211	216	?	204	?	023	\$	7	0	2	0	0	8	8	6
0000240	7	3	6	6	9	3	2	4	8	1	\n	c	y	n	d	i
0000260	030	c	y	n	d	a	q	u	i	l	1	1	1	022	s	a
0000300	n	d	y	,	U	T	\0	>	a	y	o	u	n	c	o	
0000320	u	m	o	m	l	i	v	i	n	g	i	n	s			
0000340	u	b	u	r	b	i	a	.	\0	?	\a	204	\a			

Convert Record Processor Group



ConvertRecord Properties

Processor Details

SETTINGS

SCHEDULING

PROPERTIES

COMMENTS

Required field

Property	Value	
Record Reader	Tweet JsonTreeReader with Schema Text	→
Record Writer	Tweet AvroRecordSetWriter with Schema Text	→

OK

AvroRecordSetWriter Properties without Schema Registry

Configure Controller Service

SETTINGS PROPERTIES COMMENTS

Required field

Property	Value
Schema Write Strategy	⑦ Embed Avro Schema
Schema Access Strategy	⑦ Use 'Schema Text' Property
Schema Registry	⑦ No value set
Schema Name	⑦ \${schema.name}
Schema Text	⑦ {"namespace": "twitter", "type": "record", "name": "T...}

OK

AvroRecordSetWriter Properties with Schema Registry

Configure Controller Service

SETTINGS PROPERTIES COMMENTS

Required field

Property	Value
Schema Write Strategy	HWX Content-Encoded Schema Reference
Schema Access Strategy	Use 'Schema Name' Property
Schema Registry	HortonworksSchemaRegistry
Schema Name	\$(schema.name)
Schema Text	\$(avro.schema)

OK

PublishKafkaRecord Properties

Processor Details

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field

Property	Value
Kafka Brokers	localhost:9092
Topic Name	tweets.avro.encoded
Record Reader	Tweet JsonTreeReader with HWX Schema Registry →
Record Writer	Tweet AvroRecordSetWriter with HWX Schema Registry →
Security Protocol	PLAINTEXT
Kerberos Service Name	No value set
Kerberos Principal	No value set
Kerberos Keytab	No value set
SSL Context Service	No value set
Delivery Guarantee	Best Effort
Message Key Field	No value set
Max Request Size	1 MB
Acknowledgment Wait Time	5 secs
Max Metadata Wait Time	5 sec

OK

Implementation with Flink Deserialization Schemas

Hortonworks Deserialization Schema

Confluent Deserialization Schema

```
40 @SerialVersionUID(1L)
41 class ConfluentRegistryDeserializationSchema[T](avroType: Class[T], url: String = "http://localhost:8081/") extends DeserializationSchema[T] {
42
43
44
45
46
47
48
49
50
51
52
53
54     private[this] def getSchemaId(message: Array[Byte]): Int = {
55         if (message.length < 5)
56             throw new IllegalArgumentException(s"Message is too short for schema
57             encoding reference: ${message.mkString(" ")}")
58
59         val buffer = ByteBuffer.wrap(message)
60
61         val magicByte = buffer.get
62         if (magicByte > 0)
63             throw new IllegalArgumentException(s"Unrecognized magic byte:
64             ${magicByte}")
65
66         val schemaId = buffer.getInt
67
68         schemaId
69     }
70
71     private[this] def getContentsWithSchemaId(message: Array[Byte]): (Array[Byte], Int) = {
72         val schemaId = getSchemaId(message)
73         val contents = message.drop(5)
74
75         if (contents.isEmpty)
76             throw new IllegalArgumentException(s"Message is 0 bytes")
77
78         (contents, schemaId)
79     }
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111 }
```

Next Steps with Apache Flink

Higher level SerDes for:

- Source/Sink
- TableSource/TableSink

References

- Apache NiFi – Records and Schema Registries -
<https://bryanbende.com/development/2017/06/20/apache-nifi-records-and-schema-registries>
- Confluent Schema Registry – <https://github.com/confluentinc/schema-registry>
- Github – <https://github.com/jfrazee/schema-registry-examples>
- HortonWorks Schema Registry – <http://github.com/hortonworks/registry>
- Record-Oriented Data with NiFi – <https://blogs.apache.org/nifi/entry/record-oriented-data-with-nifi>

Credits

- Bryan Bende – Staff Software Engineer, Hortonworks and PMC on Apache NiFi
- Bruno P. Kinoshita – PMC on Apache OpenNLP and Apache Commons

Questions ???